

Classical Synchronization Problems in C++ and Go

CS3211 Parallel and Concurrent Programming

Outline

- Why study classical synchronization problems?
 - Recap
- Problem solutions in C++ and Go
 - Barrier
 - Dining philosophers
 - Barber shop

Motivation

- Classical synchronization problems model problems that we have nowadays in our computer systems


Problem	CS problem
Barrier	Wait until threads/processes reach a specific point in the execution
Producer-consumer	Model interactions between a processor and devices that interact through FIFO channels.
Readers-writers	Model access to shared memory
Dining philosophers	Allocation of limited resources to a group of processes in a deadlock-free and starvation-free manner.
Barbershop	Coordinating the execution of a processor.
FIFO Semaphore	Needed to avoid starvation and increase fairness in the system.
H2O	Allocation of specific resource to a process.
Cigarette smokers	The agent represents an operating system that allocates resources, and the smokers represent applications that need resources.

Recap from your OS class

- Producer consumer problem
- Readers writers problem

Producer consumer: specification

- Processes share a buffer (bounded or unbounded)
 - Producers produce items to insert in buffer
 - Only when the buffer is not full ($< K$ items)
 - Consumers remove items from buffer
 - Only when the buffer is not empty (> 0 items)



Producer consumer in tutorial 1

- <https://godbolt.org/z/66zsjWrhd>

Producer consumer: buffered version



Example Buffer (**K = 4**)

```
while (TRUE) {  
    Produce Item;  
  
    wait(notFull) ;  
    wait(mutex) ;  
    buffer[in] = item;  
    in = (in+1) % K;  
    count++;  
    signal(mutex) ;  
    signal(notEmpty) ;  
}
```

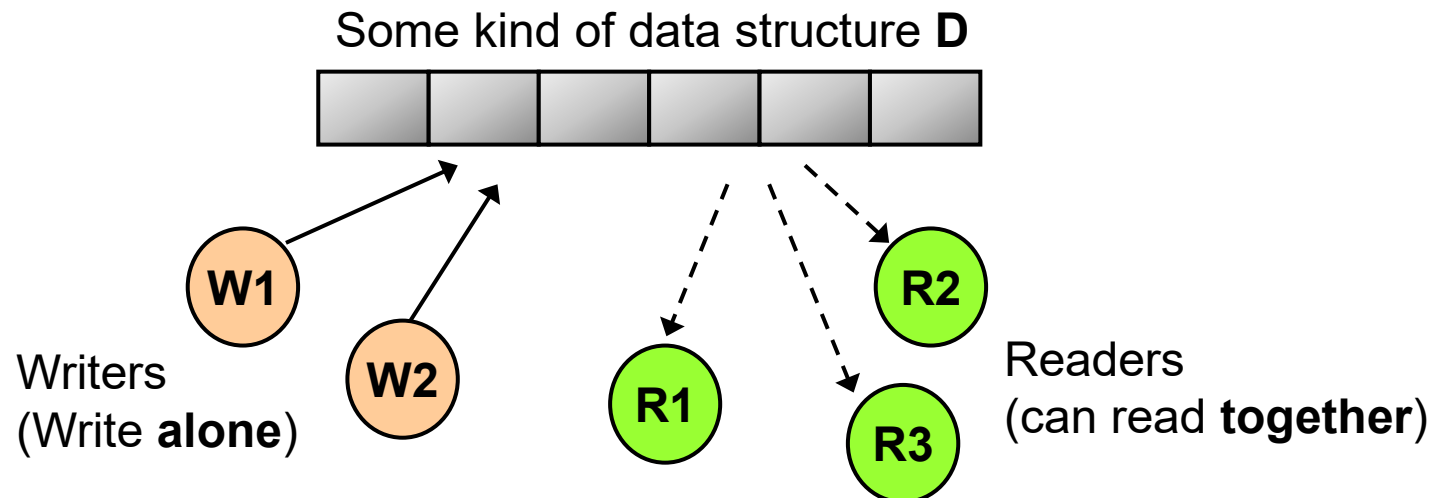
Producer

```
while (TRUE) {  
  
    wait(notEmpty) ;  
    wait(mutex) ;  
    item = buffer[out];  
    out = (out+1) % K;  
    count--;  
    signal(mutex) ;  
    signal(notFull) ;  
  
    Consume Item;  
}
```

Consumer

Readers writers: specification

- Processes share a data structure D:
 - Reader: retrieves information from D
 - Writer: modifies information in D
- Writer must have exclusive access to D
- Reader can access with other readers



Readers writers

Writers

```
roomEmpty.wait ()  
    #critical section for writers  
roomEmpty.signal ()
```

Starvation of writers is possible

Readers

```
mutex.wait ()  
readers += 1  
if readers == 1:  
    roomEmpty.wait () # first in locks  
mutex.signal ()  
# critical section for readers  
mutex.wait ()  
readers -= 1  
if readers == 0:  
    roomEmpty.signal () # last out unlocks  
mutex.signal ()
```

Lightswitch definition

Lightswitch :

counter = 0

mutex = Semaphore (1)

lock (semaphore):

mutex.wait ()

counter += 1

if counter == 1:

semaphore.wait ()

mutex.signal ()

unlock (semaphore):

mutex.wait ()

counter -= 1

if counter == 0:

semaphore.signal ()

mutex.signal ()

Readers writers with light switch

Writers

```
roomEmpty.wait ()  
    #critical section for writers  
roomEmpty.signal ()
```

Readers

```
readLightswitch.lock(roomEmpty)  
    # critical section  
readLightswitch.unlock(roomEmpty)
```

#starving writers

Use a

```
turnstile = Semaphore (1)
```

No-starve readers writers

Writers

turnstile.wait ()

roomEmpty.wait ()

critical section for writers

turnstile.signal ()

roomEmpty.signal ()

Readers

turnstile.wait ()

turnstile.signal ()

readSwitch.lock (roomEmpty)

critical section for readers

readSwitch.unlock (roomEmpty)

Readers writers locks

- Programming languages have rwlocks
 - C++17 onwards: use `shared_mutex`
 - Go has `RWLock`


```
6 class ThreadSafeCounter {
7     public:
8         ThreadSafeCounter() = default;
9
10        // Multiple threads/readers can read
11        // the counter's value at the same time.
12        unsigned int get() const {
13            std::shared_lock lock(mutex_);
14            return value_;
15        }
16
17        // Only one thread/writer can
18        // increment/write the counter's value.
19        unsigned int increment() {
20            std::unique_lock lock(mutex_);
21            return ++value_;
22        }
23
24        // Only one thread/writer can reset/write
25        // the counter's value.
26        void reset() {
27            std::unique_lock lock(mutex_);
28            value_ = 0;
29        }
30
31    private:
32        mutable std::shared_mutex mutex_;
33        unsigned int value_ = 0;
34    };
```

Outline


- Why study classical synchronization problems?
- Problem solutions in C++ and Go
 - Barrier
 - Dining philosophers
 - Barber shop

Barrier: specification


- Any thread/process must stop at this point and cannot proceed until all other threads/processes reach this barrier



(a)



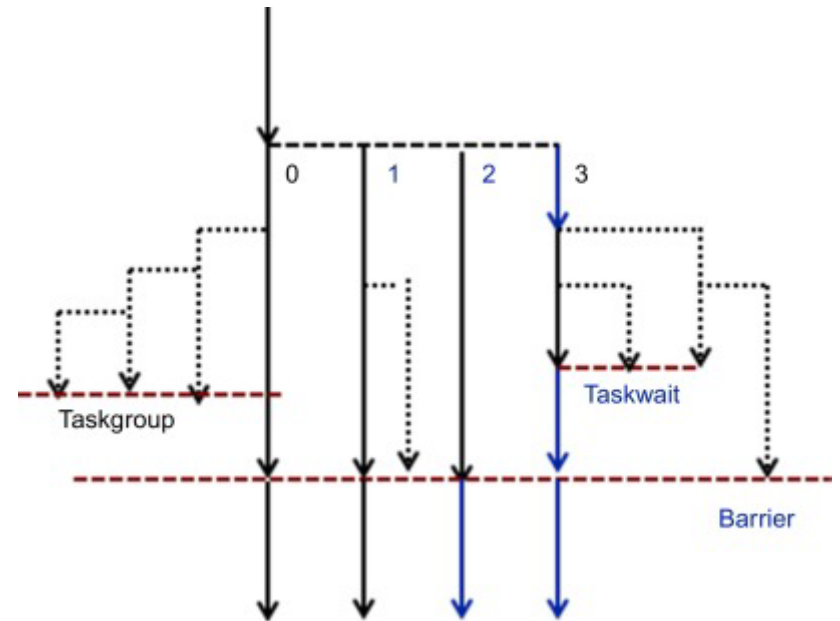
(b)



(c)

Barrier usage

- Appears in many collective routines as part of directive-based parallel languages
 - Parallel for loop in OpenMP
 - Collective communication in MPI
- Part of the programming language
 - `std::barrier` in C++20
 - `WaitGroup wait` in Go



Types of barriers

- Single use barrier or latch (`std::latch` in C++20)
 - Starts in the raised state and cannot be re-raised once it is in the lowered state
- Reusable barriers (`std::barrier` in C++20)
 - Once the arriving threads are unblocked from a barrier phase's synchronization point, the same barrier can be reused
 - Combining tree barrier - a hierarchical way of implementing barrier
 - Resolve the scalability by avoiding the case that all threads are spinning at the same location

C++: std::barrier

- In C++, create a synchronization point
 - Lines 21 & 25: wait for all threads to arrive at the barrier

```
7 int main() {
8     const auto tas = { "hw", "g", "zr", "c" };
9
10    auto on_completion = []() noexcept {
11        // locking not needed here
12        static auto phase = "... done\n" "Cleaning up...\n";
13        std::cout << phase;
14        phase = "... done\n";
15    };
16    std::barrier sync_point(std::ssize(tas), on_completion);
17
18    auto work = [&](std::string name) {
19        std::string product = " " + name + " worked\n";
20        std::cout << product;
21        sync_point.arrive_and_wait();
22
23        product = " " + name + " cleaned\n";
24        std::cout << product;
25        sync_point.arrive_and_wait();
26    };
27
28    std::cout << "Starting...\n";
29    std::vector<std::thread> threads;
30    for (auto const& worker : tas) {
31        threads.emplace_back(work, worker);
32    }
33    for (auto& thread : threads) {
34        thread.join();
35    }
36 }
```

C++: Barrier implementation

- Attempt 1:
 - Threads can go ahead other threads by one lap

```
7 struct BarrierAttempt1 {
8     std::ptrdiff_t expected;
9     std::ptrdiff_t count;
10    std::mutex mut;
11    std::counting_semaphore<> turnstile;

12
13    BarrierAttempt1(std::ptrdiff_t expected)
14        : expected{expected}, count{0}, mut{}, turnstile{0} {}
15
16    void arrive_and_wait() {
17    {
18        std::scoped_lock lock{mut};
19        count++;
20        if (count == expected) {
21            // Open turnstile
22            turnstile.release();
23        }
24    }

25
26    turnstile.acquire();
27    turnstile.release();
28
29    {
30        std::scoped_lock lock{mut};
31        count--;
32        if (count == 0) {
33            // Close turnstile to reset barrier
34            turnstile.acquire();
35        }
36    }
37 }
38 };
```

C++: Barrier implementation (2)

```
40 struct Barrier2 {
41     std::ptrdiff_t expected;
42     std::ptrdiff_t count;
43     std::mutex mut;
44     std::counting_semaphore<> turnstile;
45     std::counting_semaphore<> turnstile2;
46
47     Barrier2(std::ptrdiff_t expected)
48         : expected{expected}, count{0}, mut{},
49         |turnstile{0}, turnstile2{1} {}
```

```
51 void arrive_and_wait() {
52     {
53         std::scoped_lock lock{mut};
54         count++;
55         if (count == expected) {
56             // Close waiter turnstile
57             turnstile2.acquire();
58             // Open turnstile into the critical section
59             turnstile.release();
60         }
61     }
62 }
```

```
63 turnstile.acquire();
64 turnstile.release();
```

- Use 2
turnstiles: raise
and lower the
barrier

```
40 struct Barrier2 {
41     std::ptrdiff_t expected;
42     std::ptrdiff_t count;
43     std::mutex mut;
44     std::counting_semaphore<> turnstile;
45     std::counting_semaphore<> turnstile2;
46
47     Barrier2(std::ptrdiff_t expected)
48         : expected{expected}, count{0}, mut{},
49         |turnstile{0}, turnstile2{1} {}
```

```
51 void arrive_and_wait() {
52     {
53         std::scoped_lock lock{mut};
54         count++;
55         if (count == expected) {
56             // Close waiter turnstile
57             turnstile2.acquire();
58             // Open turnstile into the critical section
59             turnstile.release();
60         }
61     }
62 }
```

```
63 turnstile.acquire();
64 turnstile.release();
```

```
66 {
67     std::scoped_lock lock{mut};
68     count--;
69     if (count == 0) {
70         // Close turnstile to reset barrier
71         turnstile.acquire();
72         // Open second turnstile to let waiters through
73         turnstile2.release();
74     }
75 }
```

```
77 turnstile2.acquire();
78 turnstile2.release();
79 }
80 };
```

C++: Barrier implementation (3)

```
82 // Use a preloaded turnstile to let threads through faster
83 struct Barrier3 {
84     std::ptrdiff_t expected;
85     std::ptrdiff_t count;
86     std::mutex mut;
87     std::counting_semaphore<> turnstile;
88     std::counting_semaphore<> turnstile2;
89
90     Barrier3(std::ptrdiff_t expected)
91         : expected{expected}, count{0}, mut{},
92         |turnstile{0}, turnstile2{1} {}
93
94     void arrive_and_wait() {
95     {
96         std::scoped_lock lock{mut};
97         count++;
98         if (count == expected) {
99             // Close waiter turnstile
100             turnstile2.acquire();
101             // Open turnstile into the critical section
102             turnstile.release(expected);
103         }
104     }
105
106     turnstile.acquire();
107 }
```


- Lines 102 and 115: counting semaphore can be increased by expected to allow threads to pass

```
108 {
109     std::scoped_lock lock{mut};
110     count--;
111     if (count == 0) {
112         // Close turnstile to reset barrier
113         turnstile.acquire();
114         // Open second turnstile to let waiters through
115         turnstile2.release(expected);
116     }
117 }
118
119     turnstile2.acquire();
120 }
121 };
```

Go: Barrier implementation

- Use 2 WaitGroups for a reusable barrier for a group of goroutines

```
5 type Barrier1 struct {
6     wg  sync.WaitGroup
7     wg2 sync.WaitGroup
8 }
9
10 func (b *Barrier1) Init(expected int) {
11     b.wg.Add(expected)
12     b.wg2.Add(expected)
13 }
14
15 func (b *Barrier1) Wait() {
16     b.wg.Done()
17     b.wg.Wait()
18     // This line only reached when expected
19     // threads have called Wait
20     // Reset the barrier now
21     b.wg.Add(1)
22     b.wg2.Done()
23     // Wait because the barrier might not
24     // be fully reset yet
25     b.wg2.Wait()
26     // Now barrier is fully reset
27     b.wg2.Add(1)
28 }
```




Dining Philosophers

Dining philosophers: specification

- Appeared when computers were competing for access to tape drive peripherals
- Models the problem of allocating limited resources to a group of processes in a deadlock-free and starvation-free manner
- An algorithm that solves this problem without deadlock
 - Low contention: performs wonderfully when the philosophers spend any appreciable amount of time thinking, compared to eating
 - High contention: philosophers are hungry

Dining philosophers: specification



Dining philosophers: attempt 1

```
#define N 5
#define LEFT i
#define RIGHT ((i+1) % N)


//For philosopher i
while (TRUE){

    Think();
    //hungry, need food!
    takeChpStick(LEFT);
    takeChpStick(RIGHT);

    Eat();

    putChpStick(LEFT);
    putChpStick(RIGHT);

}
```



Dining philosophers: attempt 1

- Deadlock:
 - All philosopher simultaneously takes up the left chopstick, and none can proceed
- Fix the attempt:
 - Make the philosopher to put down the left chopstick if right chopstick cannot be acquired
 - Try again later
 - No deadlock:
 - Livelock: All philosopher take up left chopstick, put it down, take it up, put it down,...

Dining philosophers: attempt 2


```
#define N 5
#define LEFT i
#define RIGHT ((i+1) % N)

//For philosopher i
while (TRUE){
    Think( );

    wait(mutex);

    takeChpStick(LEFT);
    takeChpStick(RIGHT);
    Eat( );
    putChpStick(LEFT);
    putChpStick(RIGHT);

    signal(mutex);
}
```



C++: Dining philosophers' implementation (1)

- Uses **some** deadlock avoidance algorithm
 - The objects are locked by an unspecified series of calls to lock, try_lock, and unlock. If a call to lock or unlock results in an exception, unlock is called for any locked objects before rethrowing
- https://howardhinnant.github.io/dining_philosophers.html

```
7 // Basic solution to the dining philosopher's problem
8 template <size_t NumP>
9 struct DiningTable1 {
10     using ChpStick = std::mutex;
11
12     ChpStick chpSticks[NumP];
13
14     // pid = philosopher id
15     ChpStick& get_left_chpStick(size_t pid) { return chpSticks[pid]; }
16     ChpStick& get_right_chpStick(size_t pid) { return chpSticks[(pid + 1) % NumP]; }
17
18     void eat(size_t pid, void (*eat_callback)(size_t pid)) {
19         std::scoped_lock lock{get_left_chpStick(pid), get_right_chpStick(pid)};
20         eat_callback(pid);
21     }
22 };
23
```

Go: Dining philosophers' implementation

- Use odd-even ring communication to avoid the deadlock

```
3 type ChpStick struct{}
4
5 type DiningTable1 struct {
6     numPhilosophers int
7     chpStickChs      []chan ChpStick
8 }
9
10 func (t *DiningTable1) Init(numPhilosophers int) {
11     t.numPhilosophers = numPhilosophers
12
13     t.chpStickChs = make([]chan ChpStick, 0, numPhilosophers)
14     for i := 0; i < numPhilosophers; i++ {
15         chpStick := make(chan ChpStick, 1)
16         chpStick <- ChpStick{}
17         t.chpStickChs = append(t.chpStickChs, chpStick)
18     }
19 }
42 func (t *DiningTable1) Eat(pid int, eat_callback func(pid int)) {
43     evenChpStickCh := t.evenChpStickCh(pid)
44     oddChpStickCh := t.oddChpStickCh(pid)
45
46     // Use even / odd chpSticks so the resulting
47     // chpStick locking order is acyclic
48     <-evenChpStickCh
49     <-oddChpStickCh
50
51     eat_callback(pid)
52
53     evenChpStickCh <- ChpStick{}
54     oddChpStickCh <- ChpStick{}
}
```

Go: Dining philosophers' implementation

- Use odd-even ring communication to avoid the deadlock

```
21 func (t *DiningTable1) leftChpStickCh(pid int) chan ChpStick {
22     return t.chpStickChs[pid]
23 }
24 func (t *DiningTable1) rightChpStickCh(pid int) chan ChpStick {
25     return t.chpStickChs[(pid+1)%t.numPhilosophers]
26 }
27 func (t *DiningTable1) evenChpStickCh(pid int) chan ChpStick {
28     if pid%2 == 0 {
29         return t.leftChpStickCh(pid)
30     } else {
31         return t.rightChpStickCh(pid)
32     }
33 }
34 func (t *DiningTable1) oddChpStickCh(pid int) chan ChpStick {
35     if pid%2 == 1 {
36         return t.leftChpStickCh(pid)
37     } else {
38         return t.rightChpStickCh(pid)
39     }
40 }
```

```
42 func (t *DiningTable1) Eat(pid int, eat_callback func(pid int
43     evenChpStickCh := t.evenChpStickCh(pid)
44     oddChpStickCh := t.oddChpStickCh(pid)
45
46     // Use even / odd chpSticks so the resulting
47     // chpStick locking order is acyclic
48     <-evenChpStickCh
49     <-oddChpStickCh
50
51     eat_callback(pid)
52
53     evenChpStickCh <- ChpStick{}
54     oddChpStickCh <- ChpStick{}
```

Go: Dining philosophers' implementation (2)

- Swap order of chopsticks to avoid deadlock

```
62 func (t *DiningTable2) Init(numPhilosophers int) {
63     t.numPhilosophers = numPhilosophers
64
65     t.chpStickChs = make([]chan ChpStick, 0, numPhilosophers)
66     for i := 0; i < numPhilosophers; i++ {
67         chpStick := make(chan ChpStick, 1)
68         chpStick <- ChpStick{}
69         t.chpStickChs = append(t.chpStickChs, chpStick)
70     }
71 }
```

```
94 func (t *DiningTable2) Eat(pid int, eat_callback func(pid int)) {
95     evenChpStickCh := t.evenChpStickCh(pid)
96     oddChpStickCh := t.oddChpStickCh(pid)
97
98     |
99     breakLock:
100     for {
101         <-evenChpStickCh
102         select {
103             case <-oddChpStickCh:
104                 break breakLock
105             default:
106                 // Couldn't get oddChpStickCh, swap order of chpSticks
107             }
108             evenChpStickCh <- ChpStick{}
109
110             <-oddChpStickCh
111             select {
112                 case <-evenChpStickCh:
113                     break breakLock
114                 default:
115                     // Couldn't get evenChpStickCh, swap order of chpSticks
116                 }
117             oddChpStickCh <- ChpStick{}
118         }
119
120         eat_callback(pid)
121
122         evenChpStickCh <- ChpStick{}
123         oddChpStickCh <- ChpStick{}
124     }
```



Dining philosophers: Tanenbaum solution

```
#define N 5
#define LEFT ((i+N-1)% N)
#define RIGHT ((i+1) % N)

#define THINKING 0
#define HUNGRY 1
#define EATING 2

int state[N];
Semaphore mutex = 1;
Semaphore s[N];


void philosopher(int i){
    while (TRUE){
        Think();
        takeChpSticks(i);
        Eat();
        putChpSticks(i);
    }
}
```



Dining philosophers: Tanenbaum solution

```
void takeChpSticks( i )  
{  
    wait(mutex);  
    state[i] = HUNGRY;  
    safeToEat(i);  
    signal(mutex);  
    wait(s[i]);  
}
```

```
void safeToEat( i )  
{  
    if( (state[i] == HUNGRY) &&  
        (state[LEFT] != EATING) &&  
        (state[RIGHT] != EATING) ) {  
        state[i] = EATING;  
        signal(s[i]);  
    }  
}
```



Dining philosophers: Tanenbaum solution


```
void safeToEat(i)
{
    if( (state[i] == HUNGRY) &&
        (state[LEFT] != EATING) &&
        (state[RIGHT] != EATING) ) {

        state[i] = EATING;
        signal(s[i]);
    }
}
```

```
void putChpSticks(i)
{
    wait(mutex);


    state[i] = THINKING;
    safeToEat(LEFT);
    safeToEat(RIGHT);

    signal(mutex);
}
```



Dining philosophers: limited eater

```
void philosopher(int i){  
  
    while (TRUE){  
        Think( );  
        wait(seats);  
        wait(chpStick[LEFT]);  
        wait(chpStick[RIGHT]);  
        Eat( );  
        signal(chpStick[LEFT]);  
        signal(chpStick[RIGHT]);  
        signal(seats);  
    }  
}
```



C++: Dining philosophers' implementation (2)

- footman_sem used to limit the number of eaters
- Starvation is possible because mutexes and semaphores are not fair in C++

```
30     template <size_t NumP>
31     struct DiningTable2 {
32         using ChpStick = std::mutex;
33
34         ChpStick chpSticks[NumP];
35         std::counting_semaphore<> footman_sem{NumP - 1};
36
37         // pid = philosopher id
38         ChpStick& get_left_chpStick(size_t pid) { return chpSticks[pid];
39         ChpStick& get_right_chpStick(size_t pid) { return chpSticks[(pid
40
41     void eat(size_t pid, void (*eat_callback)(size_t pid)) {
42         footman_sem.acquire();
43         std::scoped_lock left_lock{get_left_chpStick(pid)};
44         std::scoped_lock right_lock{get_right_chpStick(pid)};
45         eat_callback(pid);
46         footman_sem.release();
47     }
48 };
```

C++: Dining philosophers' implementation (3)

- Use a fair semaphore
 - FIFO semaphore will be discussed in tutorial 7

```
69     template <size_t NumP>
70     struct DiningTable3 {
71         using ChpStick = FairMutex;
72
73         ChpStick chpSticks[NumP];
74         FairSemaphore footman_sem{NumP - 1};
75
76         // pid = philosopher id
77         ChpStick& get_left_chpStick(size_t pid) { return chpSticks[p
78         ChpStick& get_right_chpStick(size_t pid) { return chpSticks[
79
80         void eat(size_t pid, void (*eat_callback)(size_t pid)) {
81             footman_sem.acquire();
82             std::scoped_lock left_lock{get_left_chpStick(pid)};
83             std::scoped_lock right_lock{get_right_chpStick(pid)};
84             eat_callback(pid);
85             footman_sem.release();
86         }
87     };
```


Go: Dining philosophers' implementation (3)

- This time it doesn't matter what order each philosopher tries to obtain each chopsticks

```
126 type DiningTable3 struct {
127     numPhilosophers int
128     chpStickChs      []chan ChpStick
129     footman          chan struct{}
130 }
131
132 func (t *DiningTable3) Init(numPhilosophers int) {
133     t.numPhilosophers = numPhilosophers
134
135     t.chpStickChs = make([]chan ChpStick, 0, numPhilosophers)
136     for i := 0; i < numPhilosophers; i++ {
137         chpStick := make(chan ChpStick, 1)
138         chpStick <- ChpStick{}
139         t.chpStickChs = append(t.chpStickChs, chpStick)
140     }
141
142     t.footman = make(chan struct{}, numPhilosophers-1)
143     for i := 0; i < numPhilosophers-1; i++ {
144         t.footman <- struct{}{}
145     }
146 }
```

sical Synchron


```
155 func (t *DiningTable3) Eat(pid int,
156     ..... eat_callback func(pid int)) {
157     leftChpStickCh := t.leftChpStickCh(pid)
158     rightChpStickCh := t.rightChpStickCh(pid)
159
160     <-t.footman
161
162     select {
163     case <-leftChpStickCh:
164         ..... <-rightChpStickCh
165     case <-rightChpStickCh:
166         ..... <-leftChpStickCh
167     }
168
169     eat_callback(pid)
170
171     leftChpStickCh <- ChpStick{}
172     rightChpStickCh <- ChpStick{}
173     t.footman <- struct{}{}
174 }
```



Barbershop

(Sleeping Barber)


Problem description



- A barbershop consists of a waiting room with n chairs and the barber chair
- If there are no customers to be served, the barber goes to sleep
- If the barber is busy, but chairs are available, then the customer sits in one of the free chairs. If the barber is asleep, the customer wakes up the barber.
- If a customer enters the barbershop and all chairs are occupied, then the customer leaves the shop.

Solution

Customer is leaving



	Customer Pseudo-code		Barber Pseudo-code
1	<code>wait(mutex);</code>	21	
2	<code>if (customers == n) {</code>	22	
3	<code>signal(mutex);</code>	23	
4	<code>exit();</code>	24	
5	<code>}</code>	25	
6	<code>customers += 1;</code>	26	
7	<code>signal(mutex);</code>	27	<code>while (TRUE) {</code>
8	<code>signal(customer);</code>	28	<code>wait(customer);</code>
9	<code>wait(barber);</code>	29	<code>signal(barber);</code>
10	<code>getHairCut ();</code>	30	<code>cutHair();</code>
11	<code>signal(customerDone);</code>	31	<code>wait(customerDone);</code>
12	<code>wait (barberDone);</code>	32	<code>signal(barberDone);</code>
13	<code>wait(mutex);</code>	33	<code>}</code>
14	<code>customers -= 1;</code>	34	
15	<code>signal(mutex);</code>	35	

Line#	Initialization
1	<code>customers = 0</code>
2	<code>mutex = Semaphore (1)</code>
3	<code>customer = Semaphore (0)</code>
4	<code>barber = Semaphore (0)</code>
5	<code>customerDone = Semaphore (0)</code>
6	<code>barberDone = Semaphore (0)</code>

C++: Barbershop implementation

```
23 void customer(void (*balk)(), void (*getHairCut)()) {
24     {
25         std::scoped_lock lock{mut};
26         if (customers == MaxCustomers + 1) {
27             balk();
28             return;
29         }
30         customers++;
31     }
32
33     customer_sem.release();
34     barber_sem.acquire();
35     getHairCut();
36     done_customer_sem.release();
37     done_barber_sem.acquire();
38
39     {
40         std::scoped_lock lock{mut};
41         customers--;
42     }
43 }
44 };
```

```
4 template <size_t MaxCustomers>
5 struct Barbershop1 {
6     size_t customers;
7     std::mutex mut;
8     std::counting_semaphore<> customer_sem;
9     std::counting_semaphore<> barber_sem;
10    std::counting_semaphore<> done_customer_sem;
11    std::counting_semaphore<> done_barber_sem;
12
13    void barber(void (*cutHair)()) {
14        while (true) {
15            customer_sem.acquire();
16            barber_sem.release();
17            cutHair();
18            done_customer_sem.acquire();
19            done_barber_sem.release();
20        }
21    }
```

Go: Barbershop implementation

- Using channels
 - Customers might not get served in the FIFO order

```
23 func (bs *Barbershop1) Customer(balk func(), getHairCut func()) {
24     select {
25     case bs.chairs <- 1:
26         // Sat in chair, ask for barber
27         <-bs.barberAvailable
28         // Get hair cut
29         getHairCut()
30     default:
31         // Chairs are full
32         balk()
33     }
34 }
```

```
3 type Barbershop1 struct {
4     chairs          chan int
5     barberAvailable chan struct{}
6 }
7
8 func (bs *Barbershop1) Init(numChairs int) {
9     bs.chairs = make(chan int, numChairs)
10    bs.barberAvailable = make(chan struct{})
11 }
12
13 func (bs *Barbershop1) Barber(cutHair func()) {
14     for {
15         bs.barberAvailable <- struct{}{}
16         // Someone is sitting in the chair, free them
17         <-bs.chairs
18         // Cut their hair
19         cutHair()
20     }
21 }
```

Variations

- FIFO barbershop
- Multiple barbers

Summary

- Take advantage of the specification of the programming language when possible
- No fair semaphore
- Deadlock avoidance in
 - Acquiring multiple locks in C++
 - Odd-even ring communication
- References
- The little book of semaphores:
<https://greenteapress.com/semaphores/LittleBookOfSemaphores.pdf>
- Dining philosophers:
https://howardhinnant.github.io/dining_philosophers.html