

# Safety in Concurrent Programming with Rust

CS3211 Parallel and Concurrent Programming

# Outline

- Challenges in concurrent programming
- Introducing Rust
  - Safety, ownership and borrowing
- Concurrency in Rust
  - Threads, Mutex, Atomic
  - Libraries: Crossbeam, Rayon

# Challenges in concurrent programming

- Parallelizing ... anything is a daunting task
  - The goal is to make things faster
  - Many times, parallelizing is done by just adding another instance that does the same work
- Race conditions, data races, deadlocks, starvation
- Unsafe usage of memory in C/C++
  - Use after free (UAF)
  - Double free


**Closed** Bug 631527 Opened 11 years ago Closed 4 years ago

### Parallelize selector matching

**Categories**

Product: Core ▼

Component: CSS Parsing and Computation ▼

Type:  defect

Priority: Not set

Severity: normal

# Fearless Concurrency in Rust

- Rust was initiated with two goals in mind:
  - Safety in system programming
  - Painless concurrency



# Rust nowadays

- Strong safety guarantees
  - No seg-faults, no data races, expressive type system
- Without compromising on performance
  - No garbage collector, no runtime
  - Same level of performance as C/C++
- Goal
  - Confident, productive systems programming



# Rust

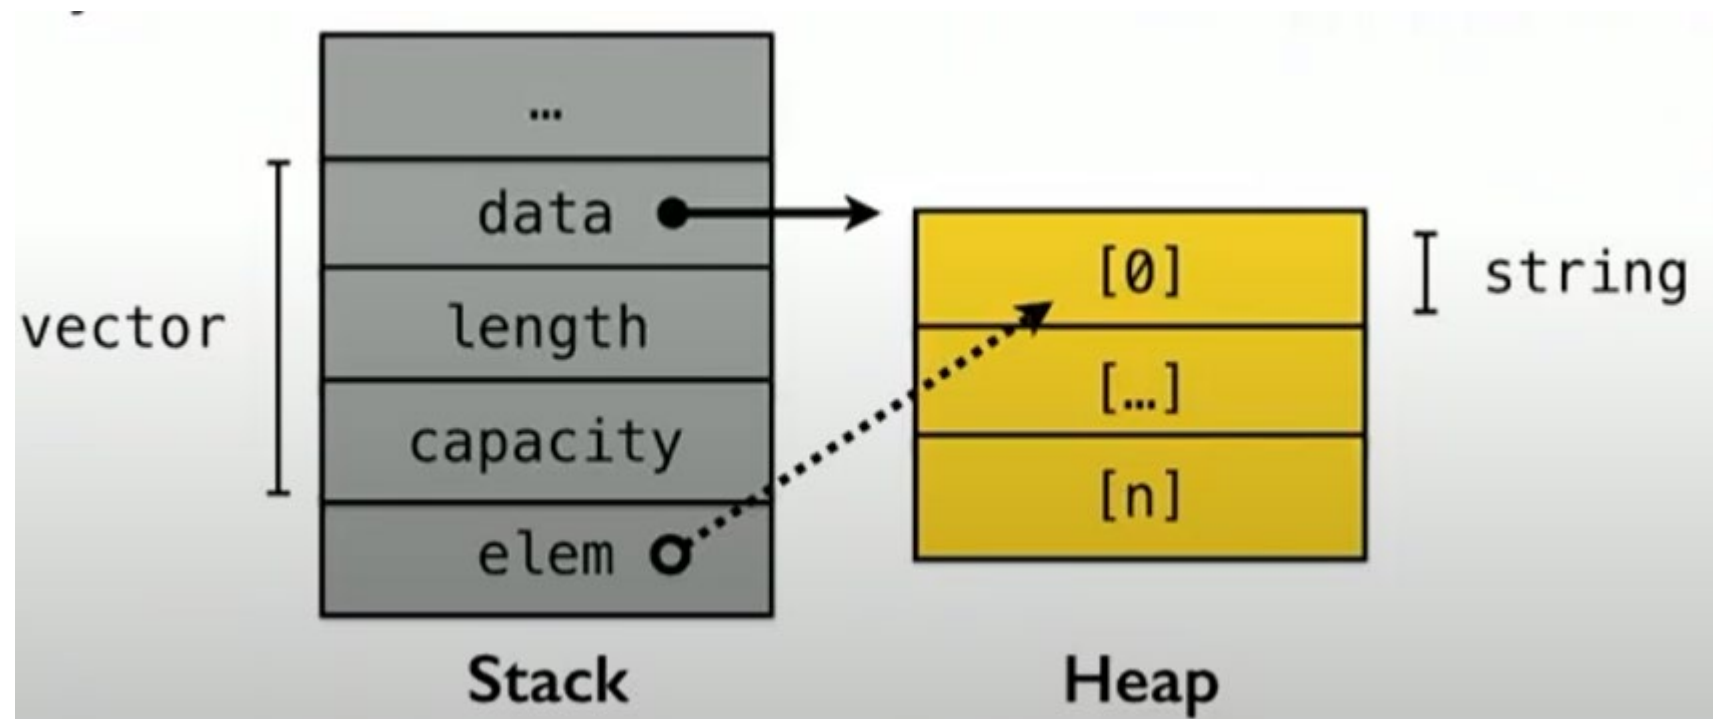
- Rustup – to install your rust tools
- Rustc – the Rust compiler
- Cargo
  - Calls the compiler – rustc
  - TOML (Tom's Obvious, Minimal Language) format for the configuration file
- Packages, crates, modules
  - A package is one or more crates that provide a set of functionality
  - A crate is a binary or library
  - Modules are used to organize code within a crate into groups
    - Privacy control



# C++ is unsafe

- Vector is freed when we exit the scope


```
1 void example() {  
2     vector<string> vector;  
3     ...  
4     auto& elem = vector[0];  
5     ...  
6 }
```



# C++ is unsafe

- Dangling pointers issues

```
1 void example() {  
2     vector<string> vector;  
3     ...  
4     auto& elem = vector[0];  
5     vector.push_back(some_string);  
6     cout << elem;  
7 }
```

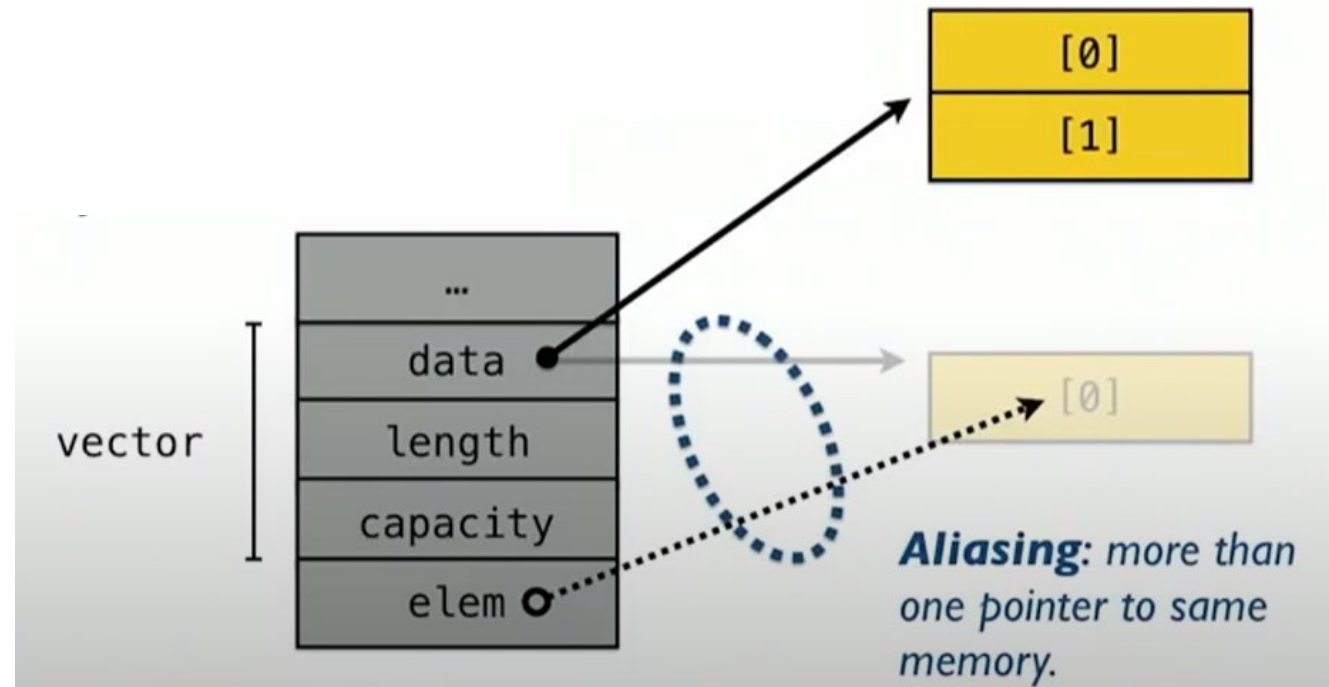




# C++ is unsafe

- Aliased pointers – pointers that point to the same chunk of memory
  - `elem` and `vector[0]`
- Mutation – changing a pointer
- **Aliasing + mutation** – changing (modifying) pointers that point same chunk of memory

```
1 void example() {  
2     vector<string> vector;  
3     ...  
4     auto& elem = vector[0];  
5     vector.push_back(some_string);  
6     cout << elem;  
7 }
```



# Solution


- **Ownership and borrowing**
  - Prevent simultaneous mutation and aliasing
- No runtime like in C++
- Memory safety like in garbage collected languages
- No data races like in ...**Rust**



# Ownership (1)

- Lines 2-4:
  - Vector `book` is initialized
  - Owner: `main` function


```
1 fn main() {  
2     let mut book = Vec::new();  
3     book.push(...);  
4     book.push(...);  
5     publish(book);  
6     // a second call to publish would  
7     // generate a compilation error  
8     // publish(book);  
9 }  
10 fn publish(book: Vec<String>) {  
11     ...  
12 }
```



# Ownership (2)

- Line 5: give ownership to `publish`
  - Pass the without `&`
- Runtime
  - Copy over the fields from `main`'s stack to `publish`'s stack
  - Forget about the first `book` in `main`
- Line 11: runs the destructor for `book`


```
1 fn main() {  
2     let mut book = Vec::new();  
3     book.push(...);  
4     book.push(...);  
5     publish(book);  
6     // a second call to publish would  
7     // generate a compilation error  
8     // publish(book);  
9 }  
10 fn publish(book: Vec<String>) {  
11     ...  
12 }
```



# Ownership (3)

- Line 5: give ownership to `publish`
- Line 8: compilation error
  - Error: use of moved value `book`

```
1 fn main() {  
2     let mut book = Vec::new();  
3     book.push(...);  
4     book.push(...);  
5     publish(book);  
6     // a second call to publish would  
7     // generate a compilation error  
8     // publish(book);  
9 }  
10 fn publish(book: Vec<String>) {  
11     ...  
12 }
```



**Ownership does not allow aliasing!**



# Rust ownership compared to C++

- Rust: giving ownership is the default
  - Not like the copy constructor in C++
  - A bit like a move in C++, but **enforced at compilation time** and **no ownership is retained**
- Rust: deep copy of data is explicit using `clone()`
  - In C++, the copy constructor does a deep copy

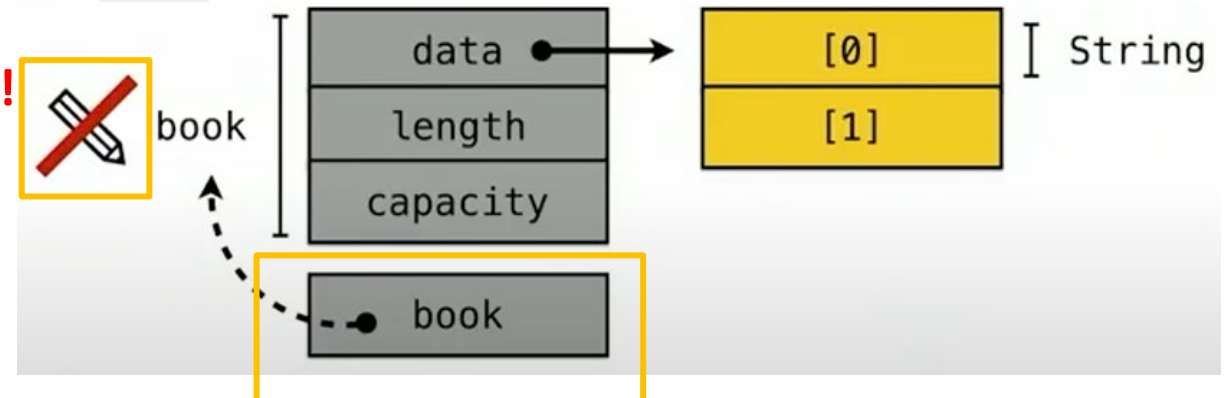


# Shared borrow

- Line 12: type is a **reference** to a vector → use &
- Line 5, 10: **borrow** the vector, creating a **shared reference**

```
1 fn main() {  
2     let mut book = Vec::new();  
3     book.push(...);  
4     book.push(...);  
5     publish(&book);  
6     // a second call to publish  
7     // borrows again the reference  
8     // to book.  
9     // compilation is successful  
10    publish(&book);  
11 }  
12 fn publish(book: &Vec<String>) {  
13     ...  
14 }
```

A shared borrow allows aliasing, but no mutation!



# Consequences of shared borrow

- **Line 4:** vector is (shared) borrowed here
  - Freezes the whole container vector
- **Line 5:** cannot mutate (compilation error)

```
1 fn example() {  
2     let mut vector = Vec::new();  
3     ...  
4     let elem = &vector[0];  
5  
6     // mutation is not allowed while  
7     // a shared borrow exists for book.  
8     // compilation error  
9     vector.push(some_string);  
10    ...  
11 }
```

**A shared borrow allows aliasing, but no mutation!**






# Mutable borrow

- Line 9: mutable reference to a vector
- Line 5, 6: mutable borrow

```
1 fn main() {  
2     let mut book = Vec::new();  
3     book.push(...);  
4     book.push(...);  
5     publish(&mut book);  
6     publish(&mut book);  
7 }  
8  
9 fn publish(book: &mut Vec<String>) {  
10     book.push(...);  
11 }
```



# Shared borrow: “Don’t break your friend’s toys”

```
1 let mut book = Vec::new();
2 book.push(...); //success: book is mutable
3 {
4     let r = &book; // shared borrow of book
5     book.push();   // compilation error: cannot mutate
6     |             | // while shared
7     r.push(...);   // compilation error: cannot mutate
8     |             | // while shared
9 } // shared borrow ends
10 book.push(...);   // success: book can be mutated
```

**A shared borrow allows aliasing, but no mutation!**



# Mutable borrow: “No, it’s my turn now!”


```
12 let mut book = Vec::new();
13 book.push(...);           //success: book is mutable
14 {
15     let r = &mut book; // mutable borrow of book
16     book.len();         // compilation error: cannot access
17     |                   | // while mutable borrow exists
18     r.push(...);         // success: reference can be mutated
19 } // mutable borrow ends
20 book.push(...);         // success: book can be mutated
```

**A mutable borrow allows mutation, but no aliasing!**



# Memory safety in Rust

- The borrow checker *statically* prevents **aliasing + mutation**
  - Compile time
  - Fighting the borrow checker!
    - Don't give up! Don't use `unsafe`!
- Ownership prevents **double-free**
  - The owner frees
- Borrowing prevents **use-after-free**
- No segfaults!



Type	Ownership	Alias?	Mutate?
T	Owned		Yes
&T	Shared reference	Yes	
&mut T	Mutable reference		Yes



# No data races in Rust

- Data race = sharing + mutation + no ordering
- Sharing + mutation are prevented in Rust



# Library-based concurrency

- Not build into the language
  - Rust had message passing build into the language – removed
- Library-based – in std or other libraries
  - Multi-paradigm
  - Leverage on ownership/borrowing



# Create a thread

- Line 1: create a thread
  - `loc` is a `JoinHandle`
  - If `loc` is dropped, the spawned thread is detached
- Line 5: join a thread

```
1 let loc = thread::spawn(|| {  
2     "world"  
3 });  
4 println!("Hello, {}!",  
5     loc.join().unwrap());
```



# Transfer the vector to a thread

- move converts any variables captured by reference or mutable reference to variables captured by *value*
  - move keyword: the **closure will take ownership of the values it uses** from the environment, thus transferring ownership of those values from one thread to another
- Line 5: **error**: use after move

```
1 let mut dst = Vec::new();  
2 thread::spawn(move || {  
3     dst.push(3);  
4 });  
5 dst.push(3);
```





# Remove the move

- Line 3: **error**: value doesn't live long enough
  - Possible memory issues: UAF
- Spawn a thread
  - Capture everything as a borrow
  - The close captures `dst` (mutable borrow)
    - Rust *infers* how to capture `dst`

```
1 let mut dst = Vec::new();
2 thread::spawn(|| {
3     dst.push(3);
4 });
5 dst.push(3);
```



# Reference counting (RC)

- Line 4: **error**: `Rc<T>` can't be sent across threads
  - RC type is not atomically managed
  - **No Send trait**

```
1 let v = Rc::new(vec![1, 2, 3]);
2 let v2 = v.clone();
3 thread::spawn(move || {
4     println!("{}", v.len());
5 });
6 another_fn(v2.clone());
```



# Traits

- Like an interface that you can implement for a given type
- It might have methods
- Example
  - Clone
- Marker traits:
  - Send – transferred across thread boundaries
    - String, u32, Arc<String>
  - Sync – safe to share references between threads
    - Type T is Sync if and only if &T is Send
  - Copy – safe to memcpy (for built-in types)
    - u32, f32
    - Not for Strings


```
trait Clone {  
    fn clone(&self) -> Self;  
}  
  
impl<T: Clone> Clone for Vec<T> {  
    fn clone(&self) -> Vec<T> {  
        let mut v = Vec::new();  
        for elem in self {  
            v.push(elem.clone());  
        }  
        return v;  
    }  
}
```



# Atomically Reference Counted - Arc

- Arc: allows only shared references
  - References cannot be mutated
- Line 3: **move** reference
- Line 4: it's safe to access v

```
1 let v = Arc::new(vec![1, 2]);  
2 let v2 = v.clone();  
3 thread::spawn(move || {  
4     println!("{}", v.len());  
5 });  
6 another_fn(&v2);
```



# Synchronization in Rust


- Shared memory
  - Mutex
  - Atomics
- Message-passing
  - Channels: MPSC (multi-producer, single-consumer FIFO queue)



# Mutex

- Based on ownership
- Data protected by the mutex
- Lock returns a guard - a proxy through which we can access the data

```
1 fn sync_inc(counter: &Mutex<i32>) {  
2     let mut data Guard<i32> = counter.lock();  
3     *data += 1;  
4 }  
    }
```



# Atomics

- Similar to modern C++
  - Same memory model
- Ordering of memory operations - SeqCst

```
1 let number = AtomicUsize::new(10);  
2 let prev = number.fetch_add(1, SeqCst);  
3 assert_eq!(prev, 10);  
4 let prev = number.swap(2, SeqCst);  
5 assert_eq!(prev, 11);  
6 assert_eq!(number.load(SeqCst), 2);
```



# Multi-Producer, Single-Consumer FIFO queue

- Channel with a reading and writing reference
  - Accepts one reader and multiple writers

```
1 let (tx, rx) = mpsc::channel ();
2 let tx2 = tx.clone();
3 thread::spawn(move || tx.send(5));
4 thread::spawn(move || tx2.send(4));
5
6 //prints 4 and 5 in an unspecified order
7 println!("{:?}", rx.recv());
8 println!("{:?}", rx.recv());
```





# Outline

- Challenges in concurrent programming
- Introducing Rust
  - Safety, ownership and borrowing
- Concurrency in Rust
  - Threads, Mutex, Atomic
  - Libraries: Crossbeam, Rayon



# Crossbeam

- Ability to create scoped threads – now in std
  - Scope is like a little container we are going to put our threads in
  - You cannot borrow variables mutably into two threads in the same scope
- Message passing using multiple-producer-multiple-consumer channel
  - With exponential backoff



# Scoped threads

- Line 5: create the scope
- Line 6: spawn threads

```
1 fn main() {  
2     let v = vec![1, 2, 3];  
3     println!("main thread has id {}", thread_id::get());  
4  
5     std::thread::scope(|scope| {  
6         scope.spawn(|inner_scope| {  
7             println!("Here's a vector: {:?}", v);  
8             println!("Now in thread with id {}", thread_id::get());  
9         });  
10    }).unwrap();  
11  
12    println!("Vector v is back: {:?}", v);  
13 }
```



# Producer consumer

- Line 12: use a bounded channel

```
11 fn main() {
12     let (send_end, receive_end) = bounded(CHANNEL_CAPACITY);
13
14     let mut threads = vec![];
15     for _i in 0 .. NUM_THREADS {
16         let send_end = send_end.clone();
17         threads.push(
18             thread::spawn(move || {
19                 for _k in 0 .. ITEMS_PER_THREAD {
20                     let produced_value = produce_item();
21                     send_end.send(produced_value).unwrap();
22                 }
23             })
24         );
25     }
26
27     for _j in 0 .. NUM_THREADS {
28         // create consumers
29         let receive_end = receive_end.clone();
30         threads.push(
31             thread::spawn(move || {
32                 for _k in 0 .. ITEMS_PER_THREAD {
33                     let to_consume = receive_end.recv().unwrap();
34                     consume_item(to_consume);
35                 }
36             })
37         );
38     }
39
40     for t in threads {
41         let _ = t.join();
42     }
43     println!("Done!");
44 }
```



# Exponential backoff

- Resources might not be available right now?
  - Retry later
- It's unhelpful to have a tight loop that simply retries as fast as possible
- Wait a little bit and try again
  - If the error occurs, next time wait a little longer
- Rationale: if the resource is overloaded right now, the reaction of requesting it more will make it even more overloaded and makes the problem worse!

# Backoff with scoped threads

- Line 12: backoff in lock-free loop
- Line 20: wait for another thread to take its turn first

```
17 fn spin_wait(ready: &AtomicBool) {  
18     let backoff = Backoff::new();  
19     while !ready.load(SeqCst) {  
20         backoff.snooze();  
21     }  
22 }
```

```
1 use crossbeam_utils::Backoff;  
2 use std::sync::atomic::AtomicUsize;  
3 use std::sync::atomic::Ordering::SeqCst;  
4  
5 fn fetch_mul(a: &AtomicUsize, b: usize) -> usize {  
6     let backoff = Backoff::new();  
7     loop {  
8         let val = a.load(SeqCst);  
9         if a.compare_and_swap(val, val.wrapping_mul(b), SeqCst) == val {  
10             return val;  
11         }  
12         backoff.spin();  
13     }  
14 }
```



# Rayon

- A data parallelism library
  - Parallelize some spots without full/major rewrite
- Similar in functionality with OpenMP
  - But OpenMP uses compiler directives
- Rationale: reasonably common that computationally-intensive parts of the program happen in a loop, so parallelizing loops is likely to be quite profitable



# Example: Sequential maximum of a vector

```
1 fn main() {
2     let vec = init_vector();
3     let max = MIN;
4     vec.iter().for_each(|n| {
5         if *n > max {
6             max = *n;
7         }
8     });
9
10    println!("Max value in the array is{}",max);
11    if max == MAX {
12        println!("This is the max value for an i64.")
13    }
14 }
```





# Example: Maximum of a vector with Rayon

```
 9 fn main() {
10     let vec = init_vector();
11     let max = AtomicI64::new(MIN);
12     vec.par_iter().for_each(|n| {
13         loop {
14             let old = max.load(Ordering::SeqCst);
15             if *n <= old {
16                 break;
17             }
18             let returned = max.compare_and_swap(old, *n, Ordering::SeqCst);
19             if returned == old {
20                 println!("Swapped {} for {}.", n, old);
21                 break;
22             }
23         }
24     });
25     println!("Max value in the array is {}", max.load(Ordering::SeqCst));
26     if max.load(Ordering::SeqCst) == MAX {
27         println!("This is the max value for an i64.")
28     }
29 }
```



# Rust in a nutshell

- Zero-cost abstraction (like C/C++)
- Memory safety & data-race freedom
- Results in
  - Confident, productive systems programming
- References:
  - <https://www.youtube.com/watch?v=SiUBdUE7xnA>
  - <https://www.youtube.com/watch?v=L0dEE2IqbD8>
  - <https://www.youtube.com/watch?v=6BYKw0Y758Q>

