

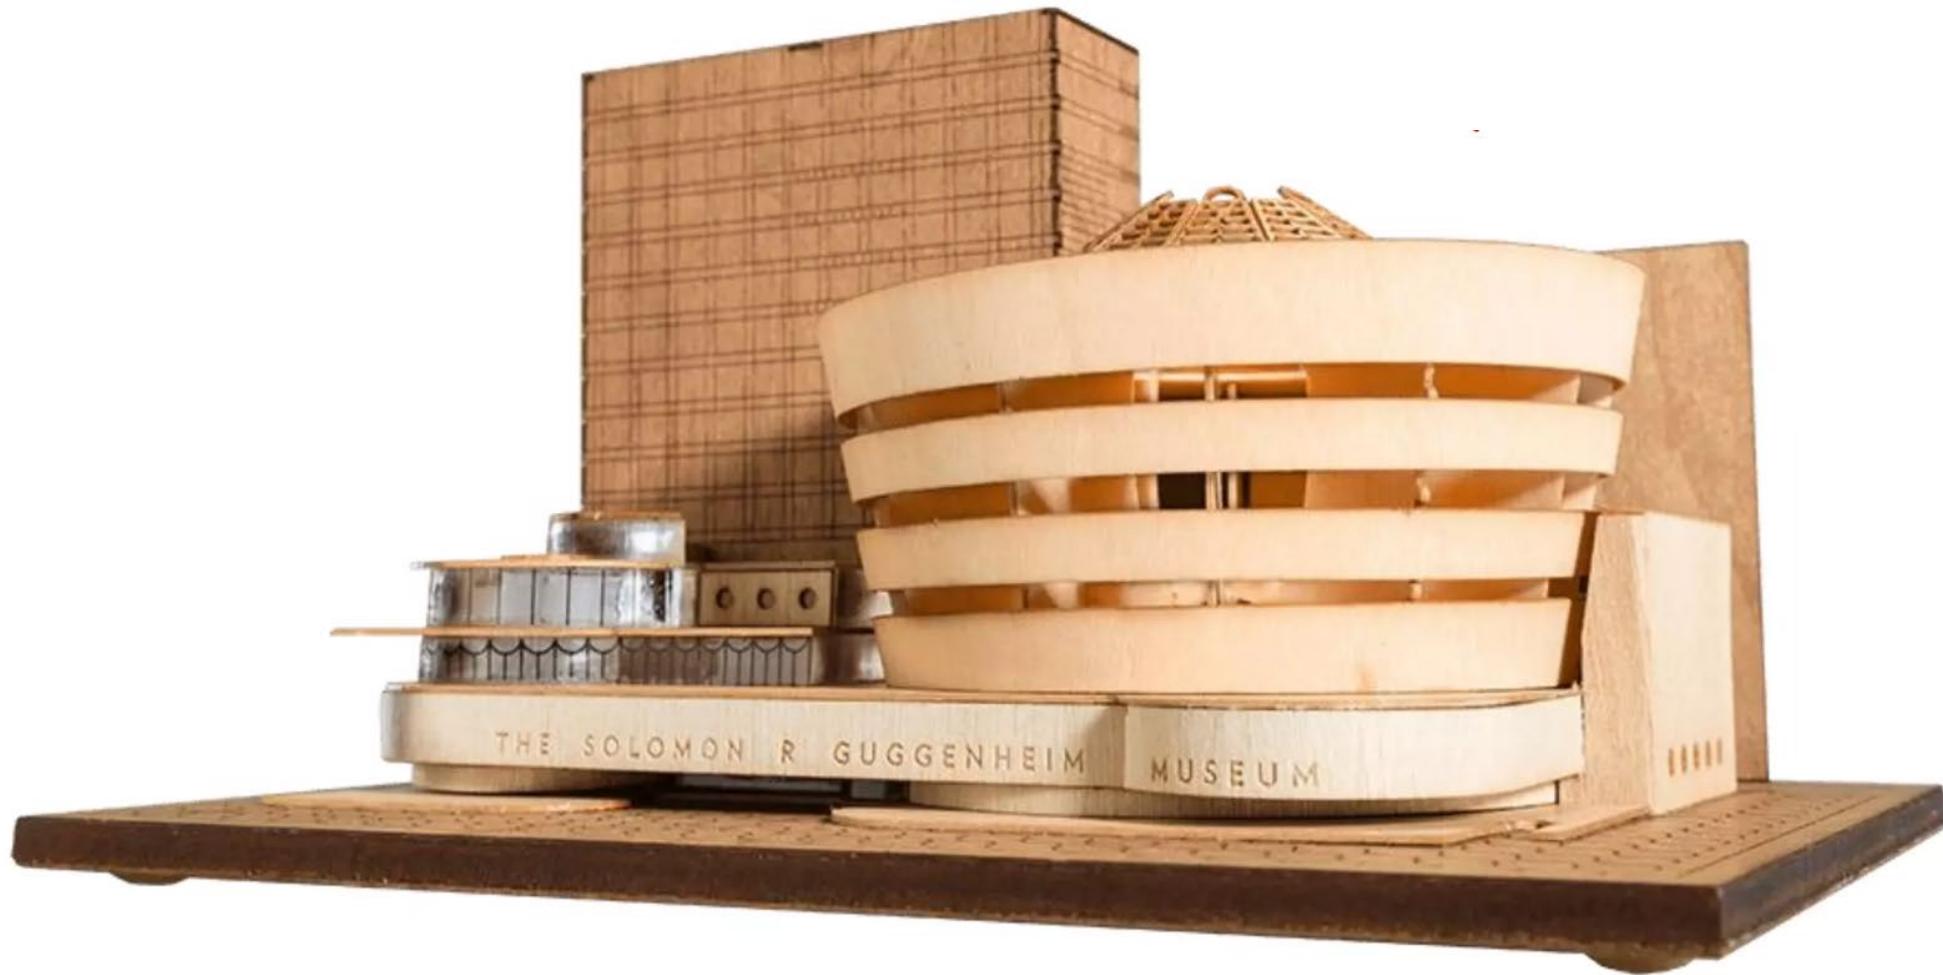
# Checking Concurrent Programs

CS3211 Parallel and Concurrent Programming

# Increase the confidence in your concurrent programs

- My code works and I don't know why
  - No confidence that the code will always work
  - Will making a change break my code?
- How to increase the confidence in your code
  - (Classic testing methods)
  - Sanitizers (in CS3211 lecture 4)
  - **Model checking** – mathematically proving that the code is correct
    - Many approaches out there
    - Unfortunately, these approaches are not really used in the industry (much)

# Model and checking the model



# Model checking of a formal specification

- Build the **model** using a special Domain Specific Language (DSL)
  - Write a formal specification, usually done in a new language
- Then **check** the model
  - Are all the constraints met?
  - Does anything unexpected happen?
  - Does it deadlock?
- Why spend time with this?
  - Check things make sense before starting the (costly) implementation
  - Prove certain properties for existing code
  - Allow for aggressive optimizations without compromising correctness

aka “**formal methods**” approach

# Formal Methods

## Advantages

- Rigorous
- Verify all traces exhaustively
- Produce a system run that violates the requirement

## Disadvantages

- The specification used is faulty
- Tedious in coming up with a complete specification
- Time consuming

# Approaches to Model Checking

- Write a formal specification of the system and check it
  - Use a model checker or proof assistant
    - Various degrees of automation
    - The checker usually checks the states exhaustively
  - Use the specification to write the code
    - Manually write code
    - Automatically generate code from specification
- Add formal specification (invariants that should hold) in the code, as comments
  - Use a model checker to check the invariants
  - Difficult to make the model checker understand the code
    - Use some symbolic execution
  - Limited in functionalities for concurrent code

# Model checkers for concurrent programs

- TLA<sup>+</sup> (TLC) - Temporal Logic of Actions+
  - Focusses on temporal properties
  - Good for modeling concurrent systems (and distributed systems)
- Coq Proof Assistant
  - Generates oCaml, Haskell and Scheme
  - Good for interactive proof methods
- Alloy (alloy analyzer)
  - Focusses on relational logic
  - Good for modeling structures

# TLA+

- Proposed by Leslie Lamport in 1999
  - Defines TLA+ as a "quixotic attempt to overcome engineers' antipathy towards mathematics"
- High-level language for modeling programs and systems – especially concurrent and distributed ones
- Based on the idea that the best way to describe things precisely is with simple mathematics
- Approach
  - A specification in TLA+ is written
  - The specification is proven (verified) using a checker by exhaustively testing the states
  - Manually write the code based on the TLA+ spec

# How it works?

- The model checker finds all possible system behaviours (states) up to some number of execution steps
- Examines the states for violations of desired invariance properties such as **safety** and **liveness**
- TLA+ specifications use basic set theory to define **safety** (bad things won't happen) and temporal logic to define **liveness** (good things eventually happen)

# TLA+

- Temporal (time)
- Logic (Boolean logic) of
- Actions (state machines)
- Plus (some stuff)

# Boolean Logic

Boolean	Mathematics	TLA+	Programming
AND	$a \wedge b$	$a \wedge b$	$a \&\& b$
OR	$a \vee b$	$a \vee b$	$a \mid\mid b$
NOT	$\neg a$	$\sim a$	$!a; \text{not } a$

LaTeX friendly!



# Boolean Logic

- A **predicate** is an expression that returns a Boolean

```
// programming language definition
```

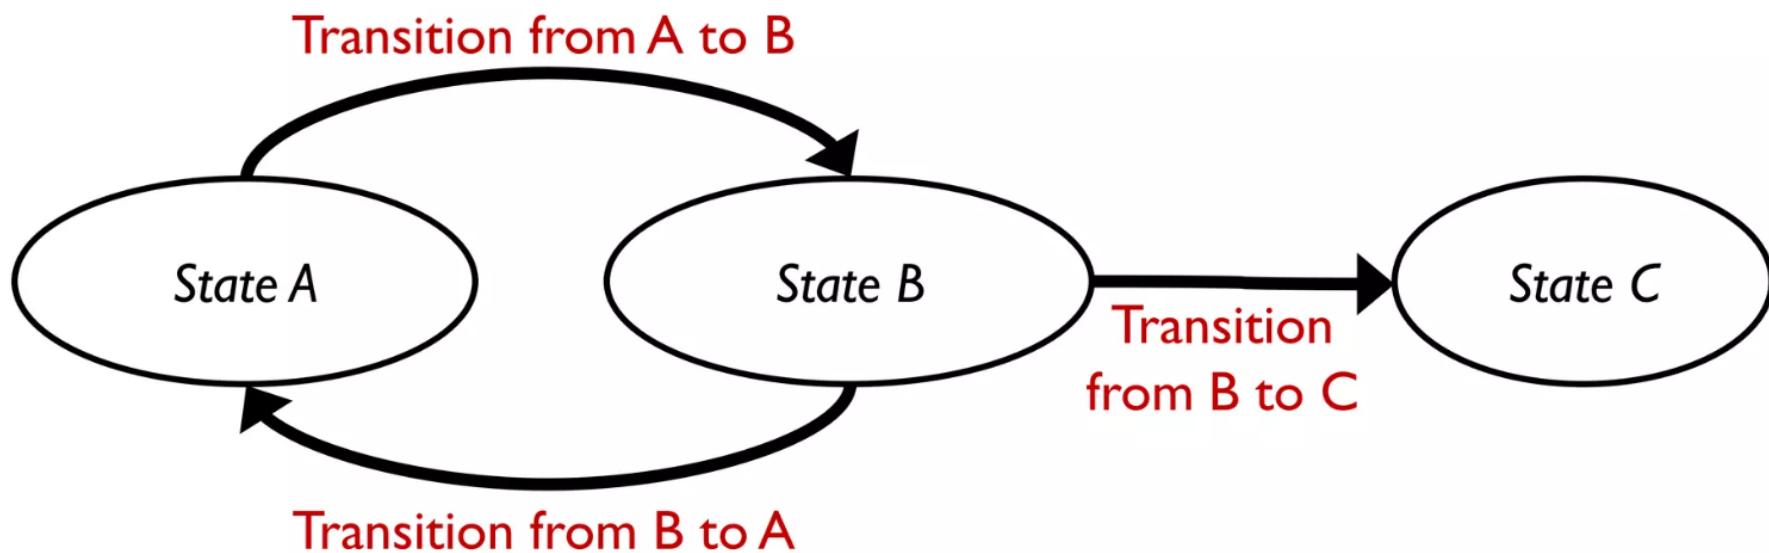
```
function(a,b,c) {  
    (a && b) || (a && !c)  
}
```

```
\* TLA-style definition
```

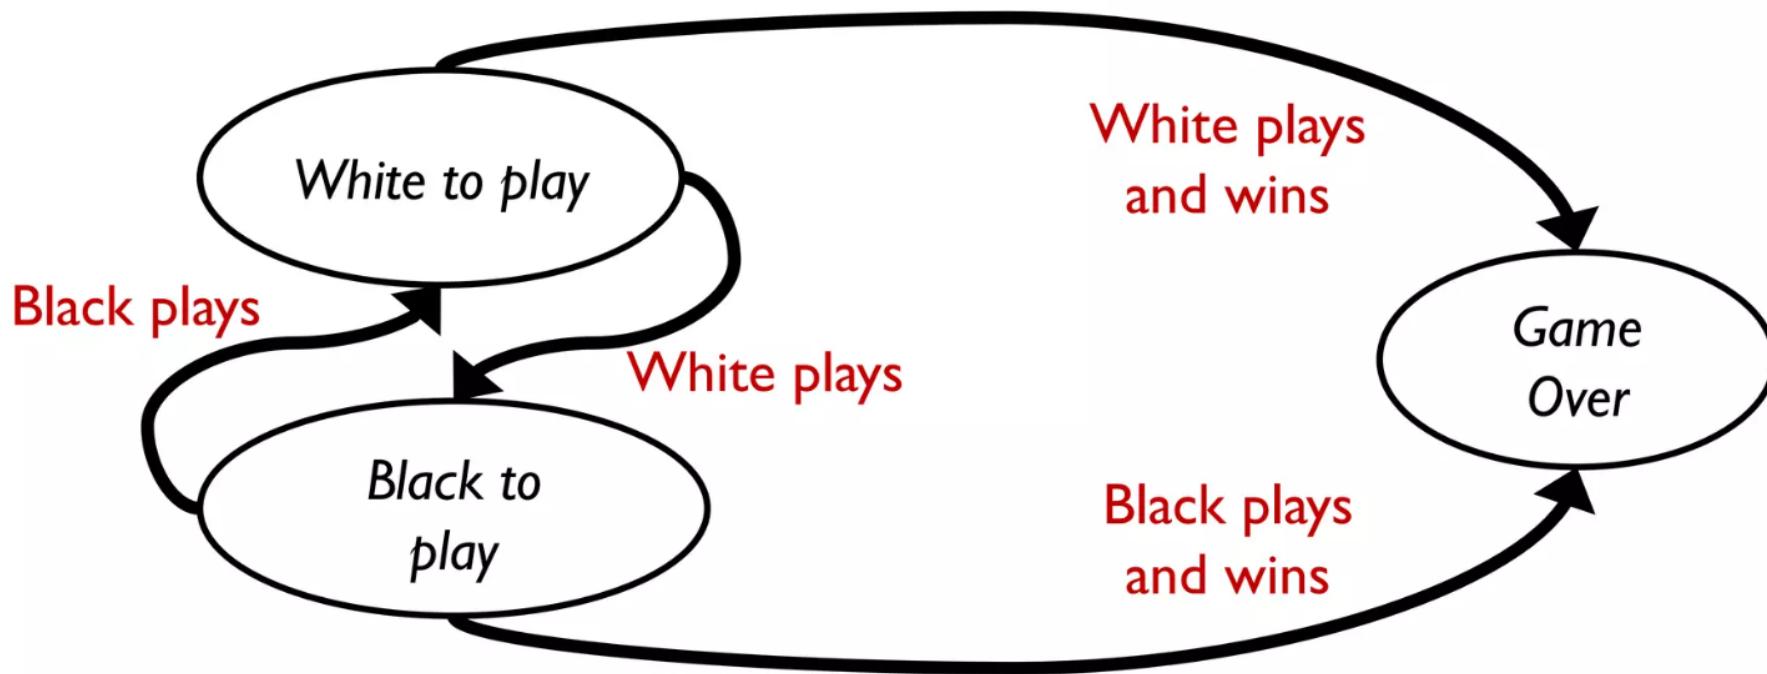
```
operator(a,b,c) ==  
    (a /\ b) \vee (a /\ ~c)
```

# Actions – state machines

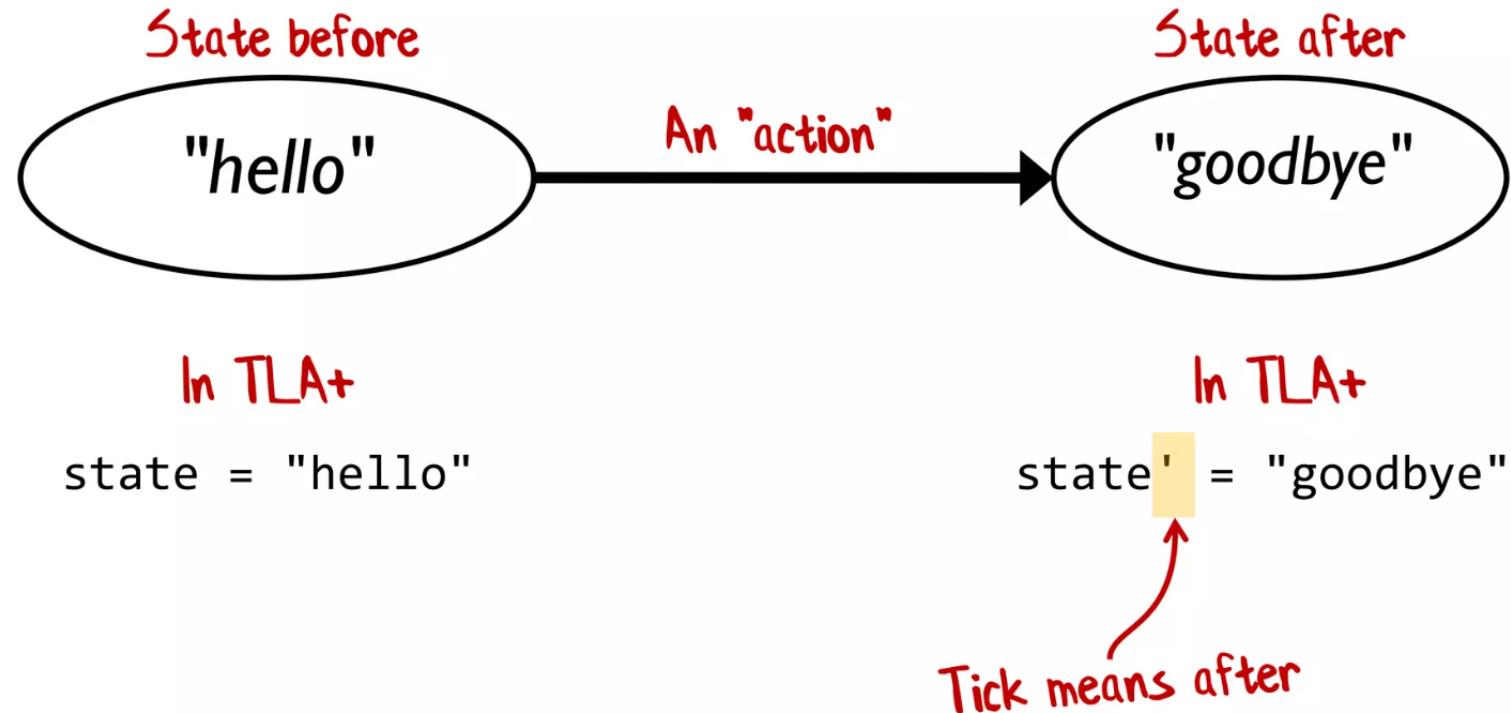
- State machines
  - States
  - Transitions



# State machine for playing chess



# Formalizing the actions



# Formalizing the actions



In TLA+, define the action "Next" like this

```
Next ==  
    state = "hello"  
    /\ state' = "goodbye"
```

# The action is the transition

- This is test, not an assignment!



```
Next ==  
state = "hello"  
/\ state' = "goodbye"
```

This is not an  
assignment!

# Actions are tests

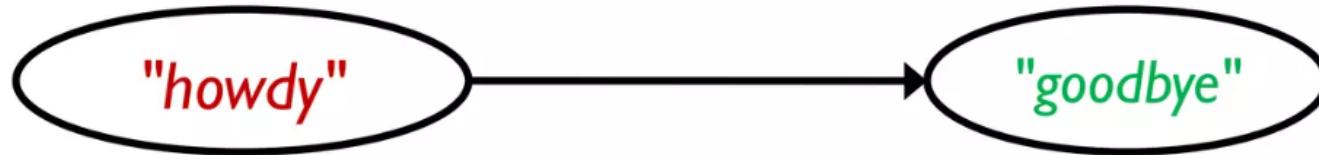
`state = "hello"  $\wedge$  state' = "goodbye"`



✓ Does match



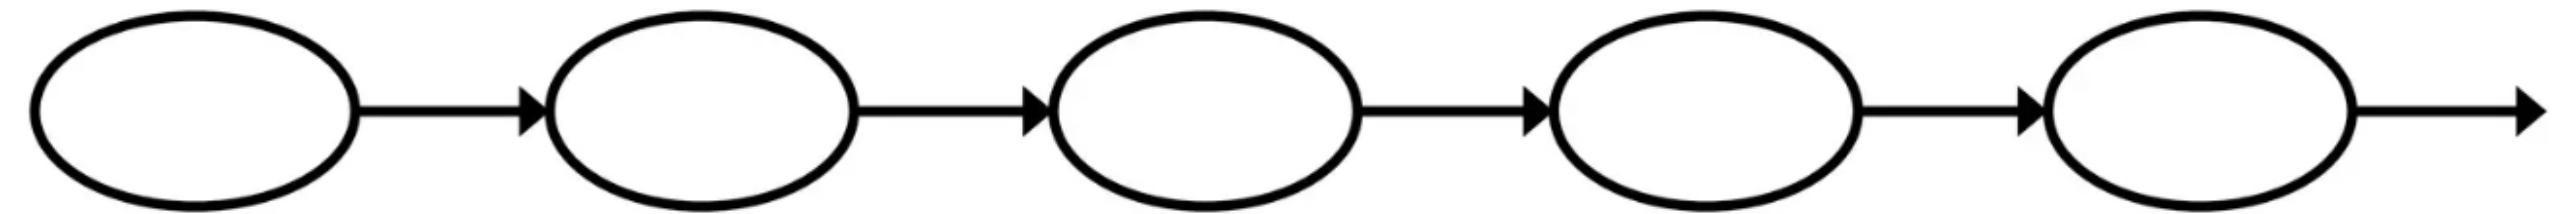
✗ Doesn't match



✗ Doesn't match

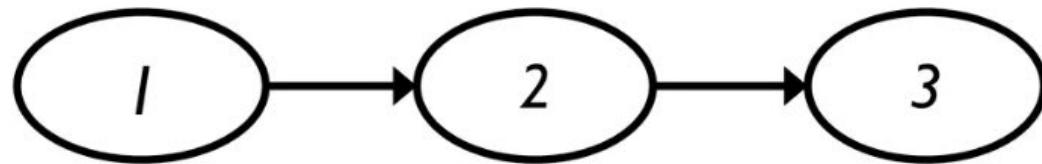
# Temporal – state transitions over time

- Infinite amount of time



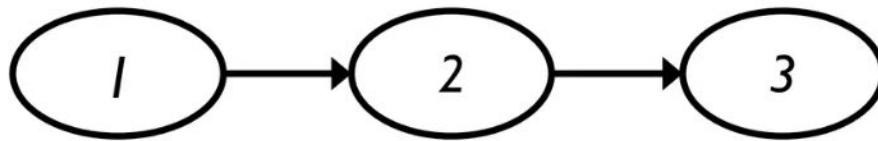
- TLA+ can ask questions like:
  - Is something always true?
  - Is something ever true?
  - If X happens, must Y happen afterwards?over time

# Count to three



```
// programming language version
var x = 1
x = 2
x = 3
```

# Count to three



\\* TLA version

**Init** ==

x=1

\\* initial state

**Next** ==

(x=1 /\ x'=2)

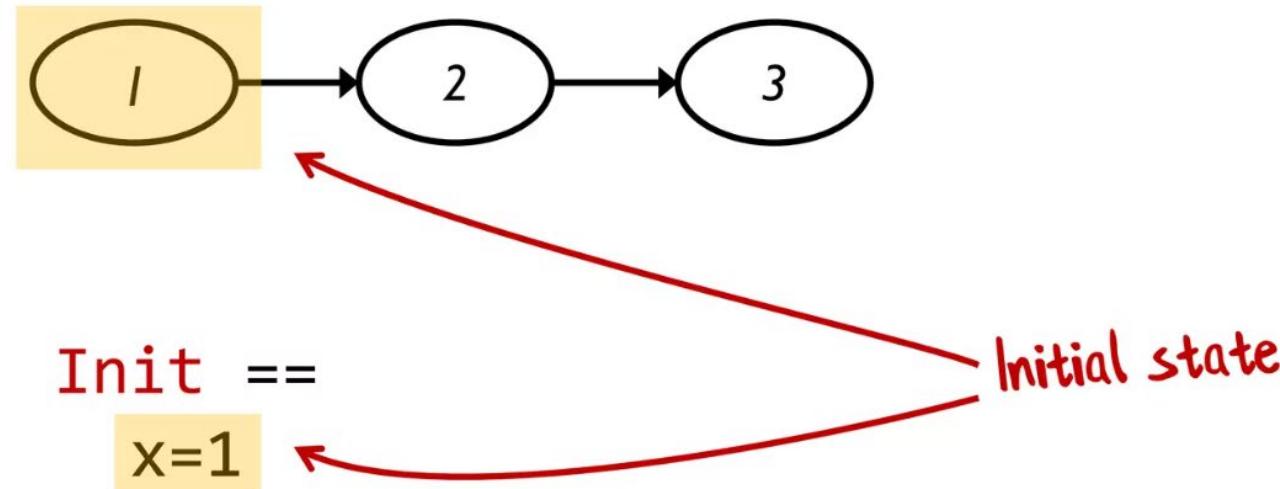
\/ (x=2 /\ x'=3)

\\* transition

\\* match step 1

\\* or match step 2

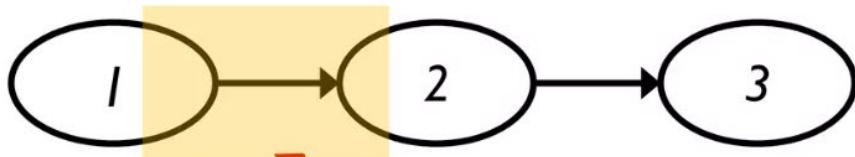
# Count to three



Next ==

$$(x=1 \wedge x'=2) \quad \text{\scriptsize/* match step 1}$$
$$\vee (x=2 \wedge x'=3) \quad \text{\scriptsize/* or match step 2}$$

# Count to three



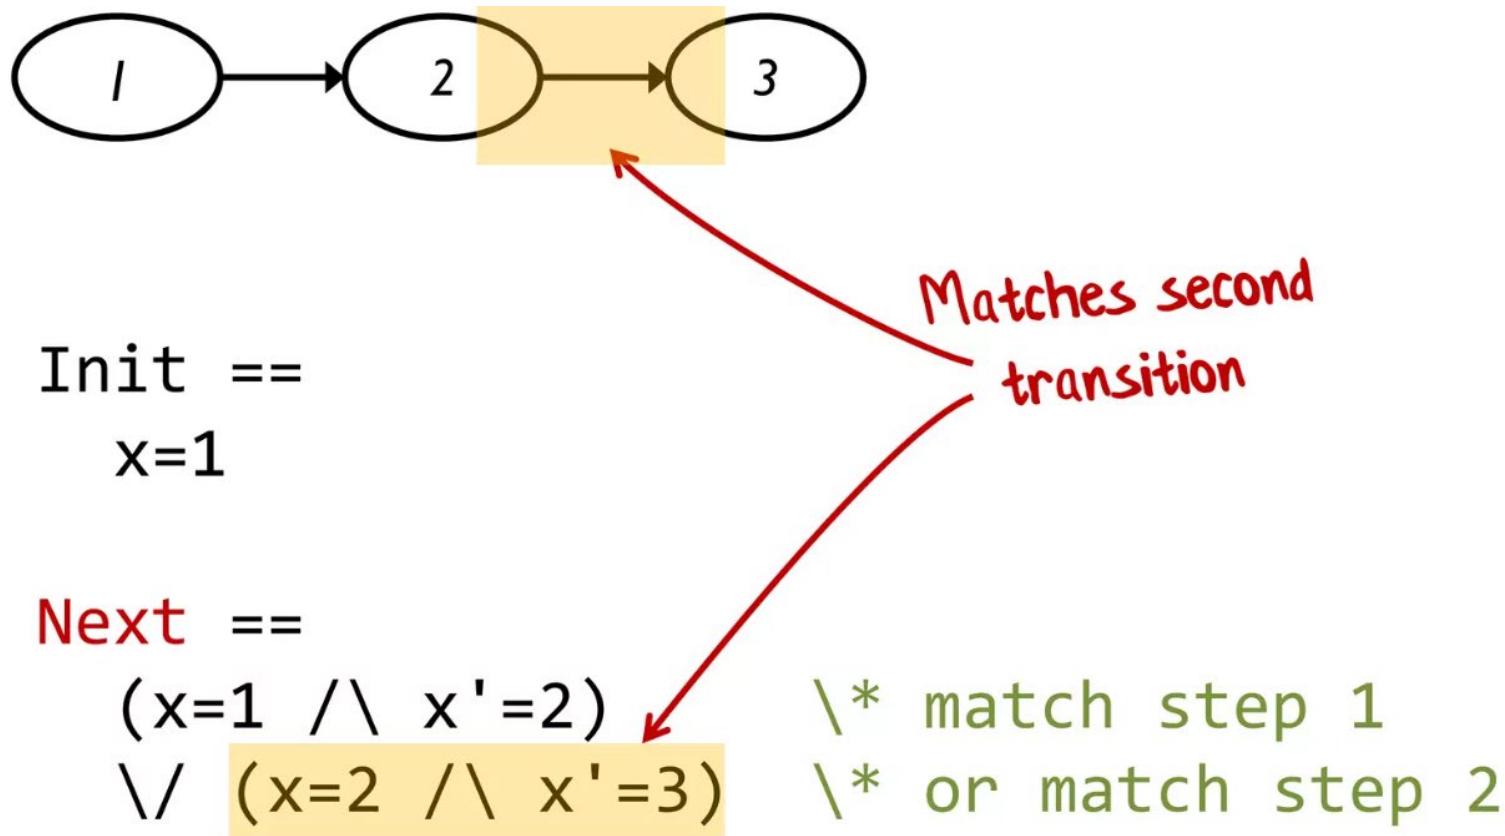
Init ==  
x=1

Matches first  
transition

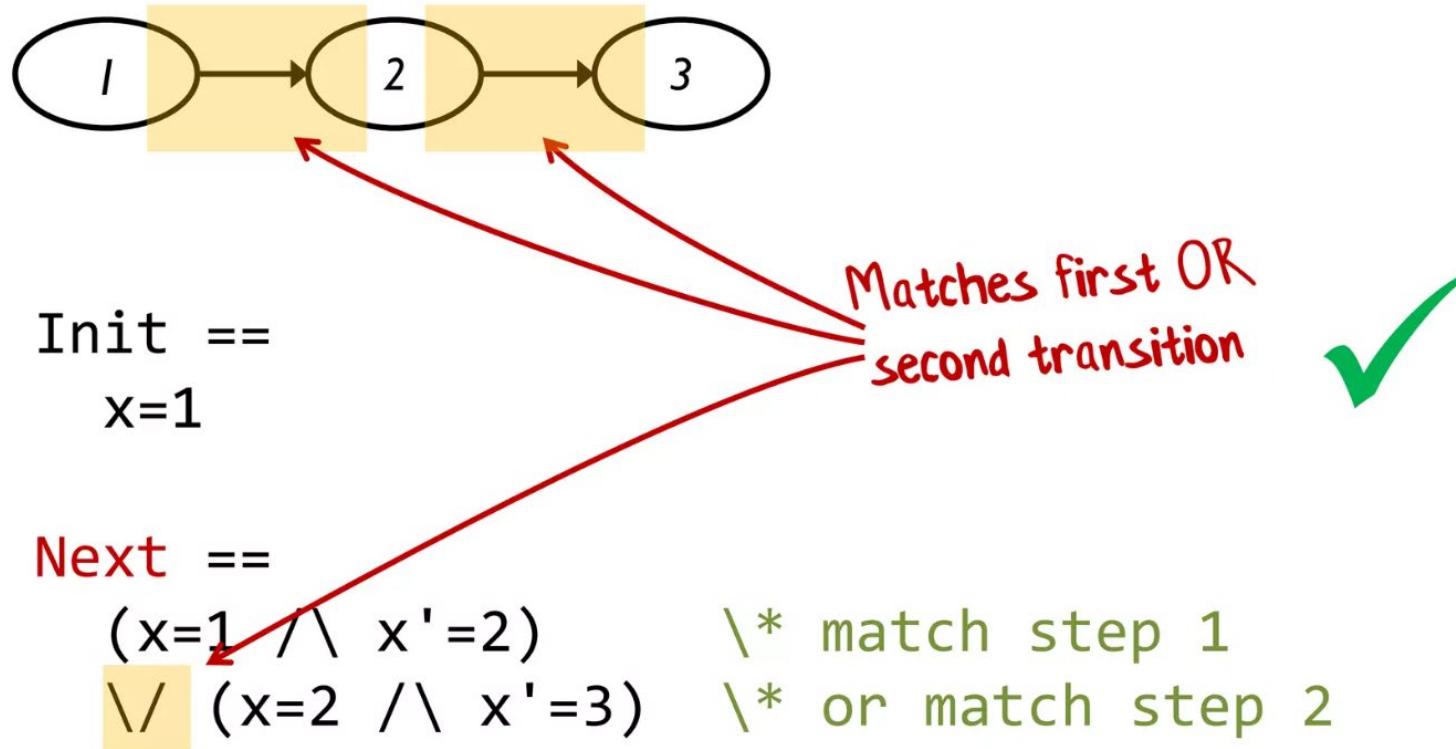
Next ==  
 $(x=1 \wedge x'=2) \vee (x=2 \wedge x'=3)$

\\* match step 1      \\* or match step 2

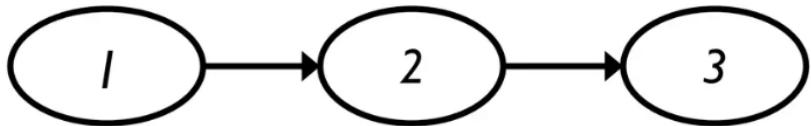
# Count to three



# Count to three



# Count to three, refactored



```
Init == x=1  
Step1 == x=1 /\ x'=2  
Step2 == x=2 /\ x'=3  
Next == Step1 \/\ Step2
```

These are called "operators".  
They are like macros,  
not procedures/functions

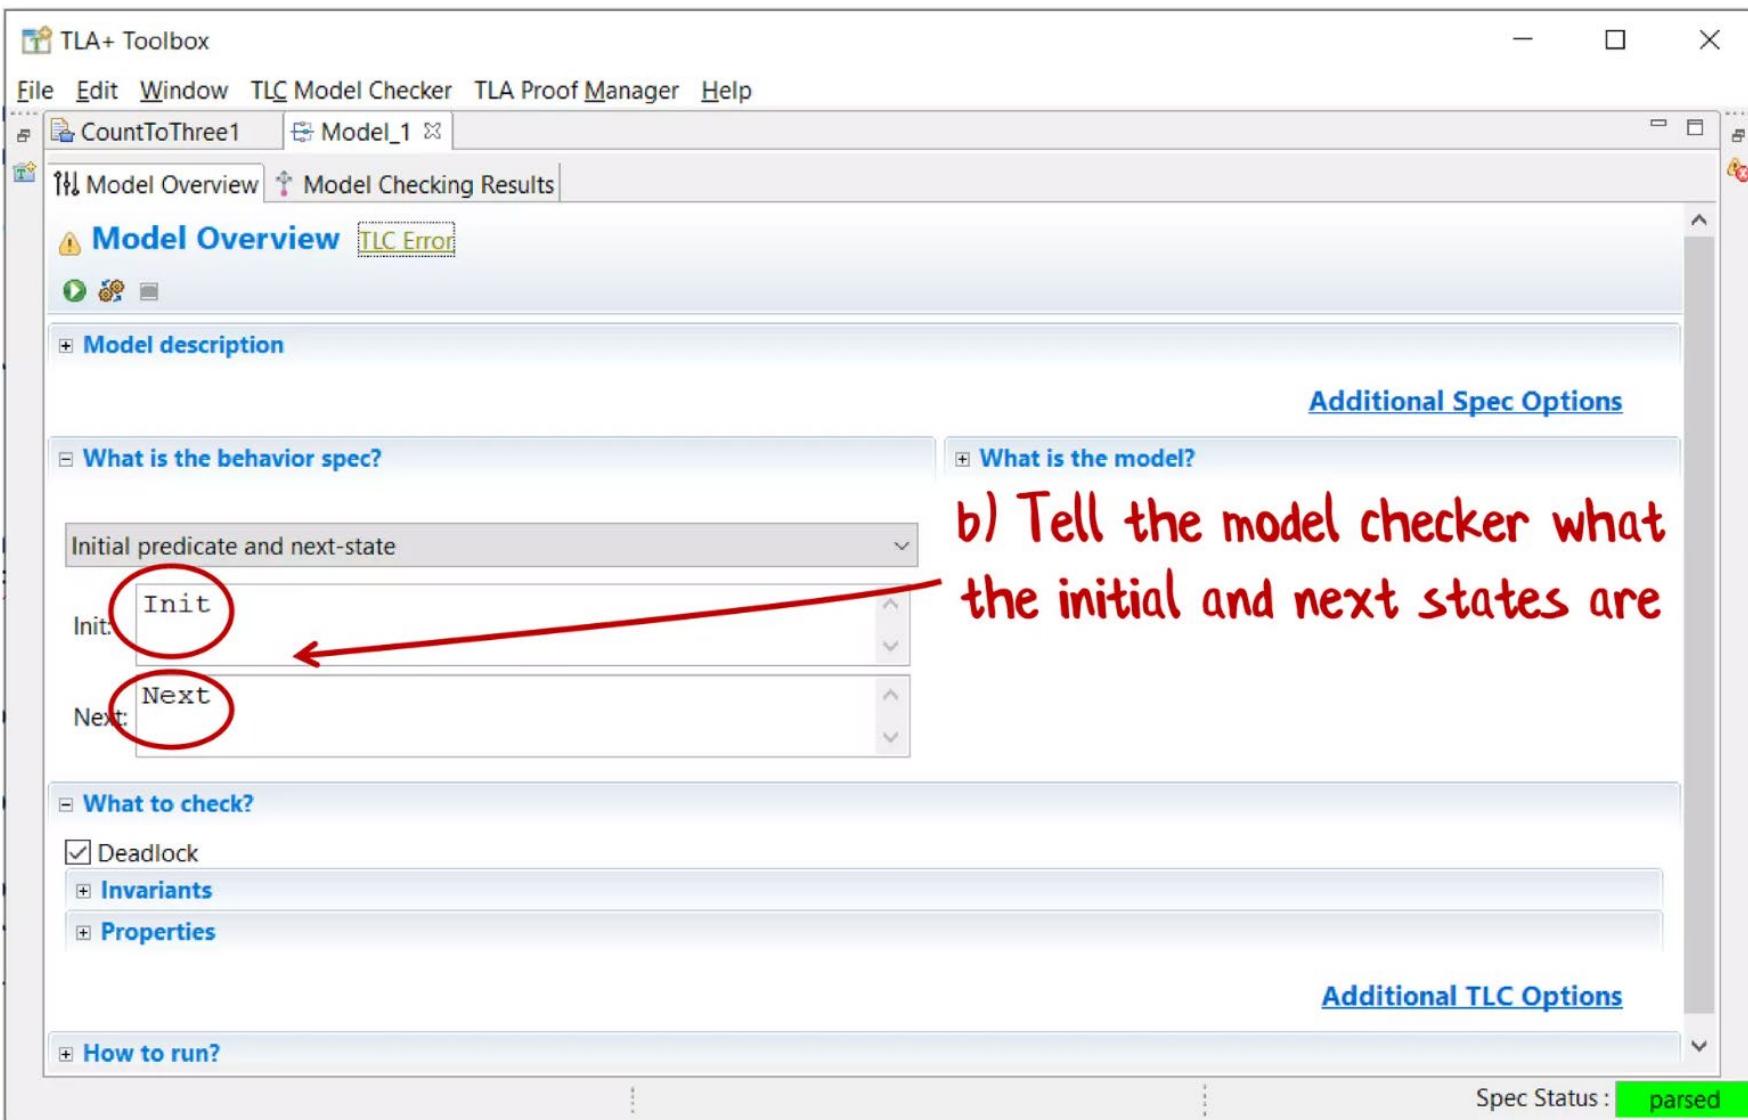
Refactored version.  
Steps are now explicitly named

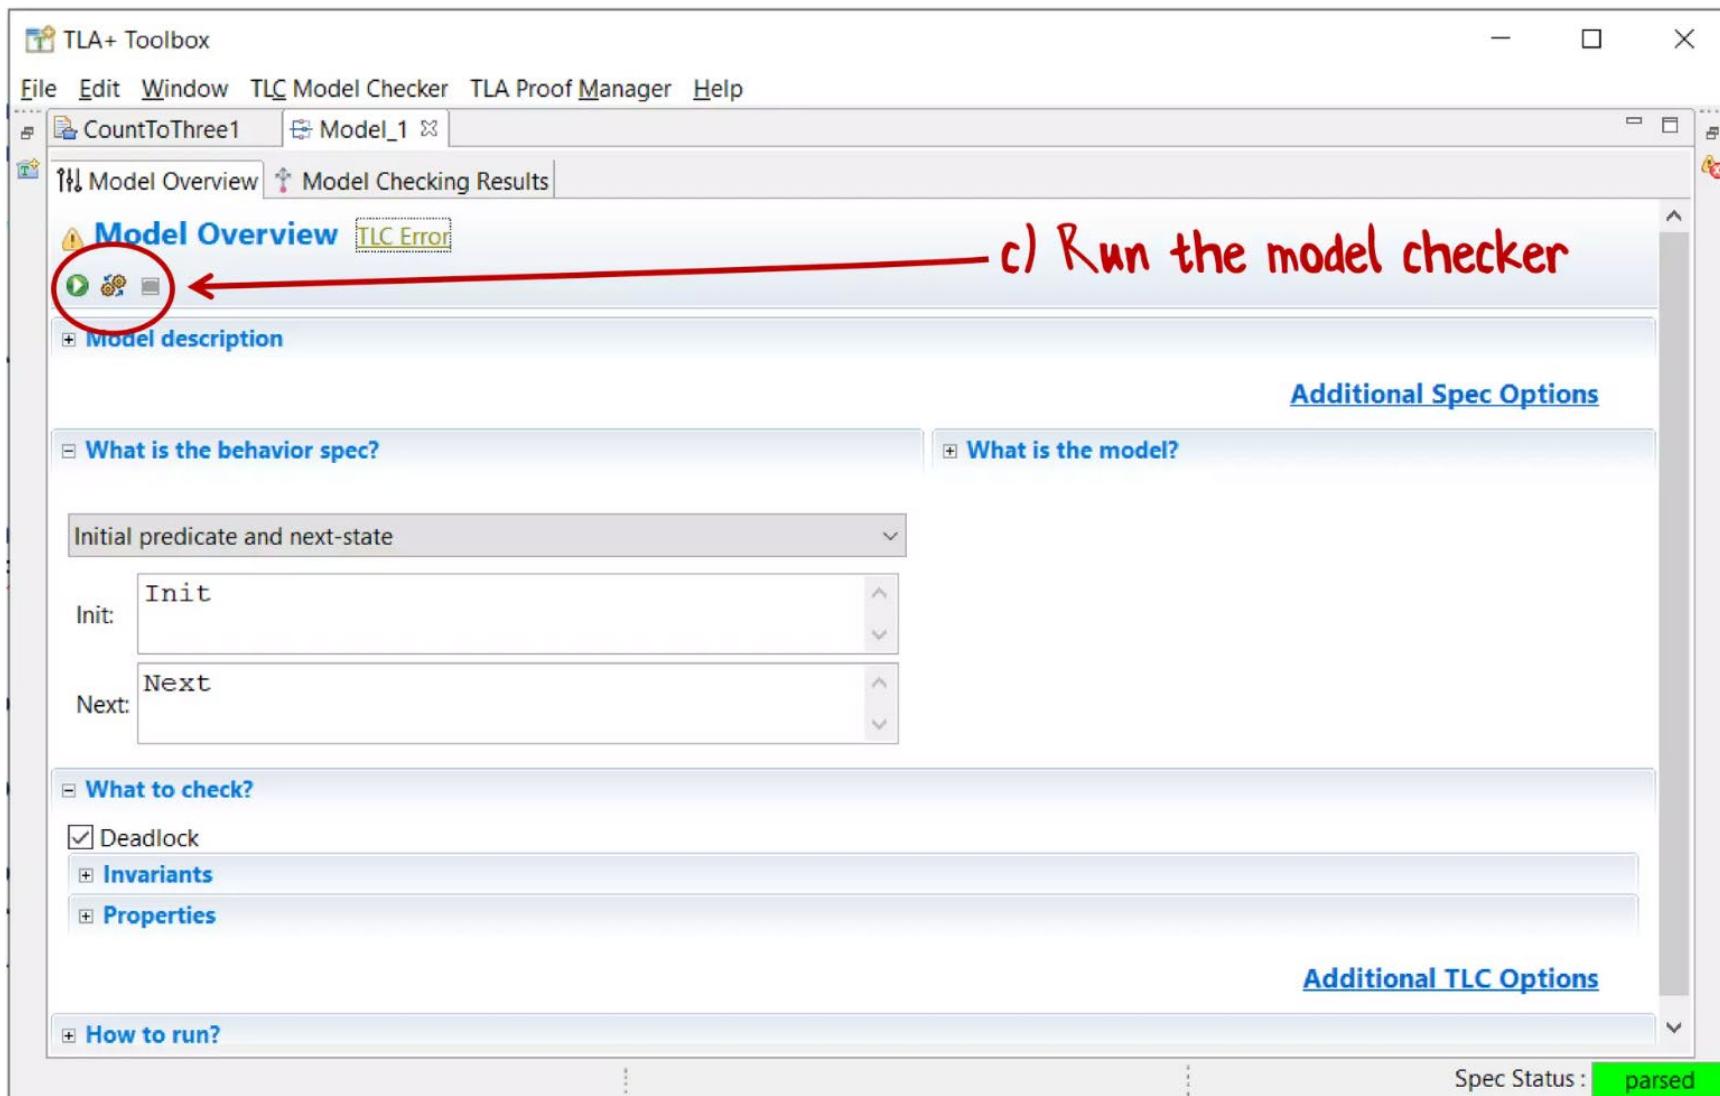
# TLA+ toolbox: IDE and checker

a) Type in the  
script and check  
that there are no  
syntax errors

```
1 MODULE CountToThree1
2
3 EXTENDS Integers
4
5 VARIABLE x
6
7
8 Init == x=1
9 Step1 == x=1 /\ x'=2
10 Step2 == x=2 /\ x'=3
11 Next == Step1 \wedge Step2
12
13
14 =====
15 /* Modification History
16 /* Last modified Fri Jun 05 09:24:53 BST 2020 by swlas
17
```

Spec Status : parsed





**Look at the error**

The screenshot shows the TLA+ Toolbox interface with two main windows:

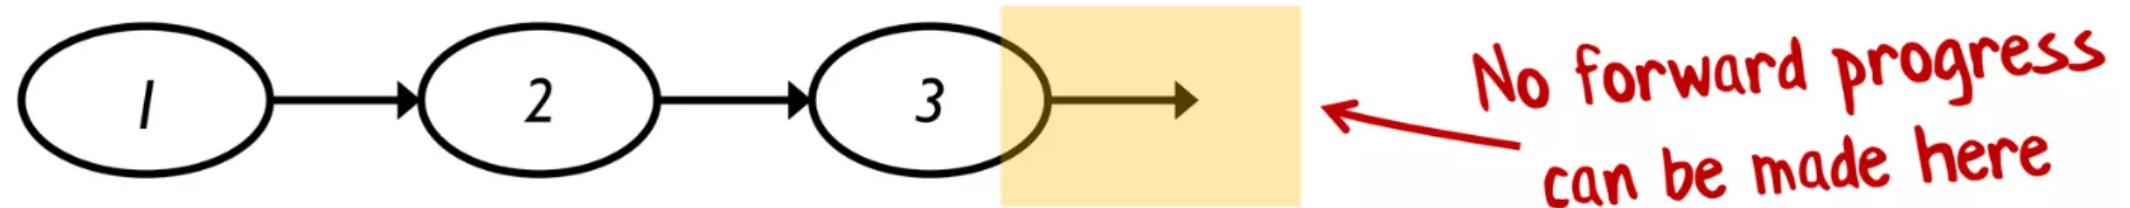
- TLA+ Toolbox** window (left):
  - File Edit Window TLC Model Checker TLA Proof Manager Help
  - CountToThree1 Model\_1
  - Model Overview Model Checking Results
  - Model Overview** (selected) TLC Error
  - Initial predicate and next-state:
    - Init: `Init`
    - Next: `Next`
  - What to check?:
    - Deadlock
    - + Invariants
    - + Properties
- TLC Errors** window (right):
  - Model 1
  - Deadlock reached.
  - Error-Trace Exploration
  - Error-Trace:
 

Name	Value
<Initial predicate>	State (num = 1)
x	1
<Next line 12, col 4 to	State (num = 2)
x	2
<Next line 13, col 7 to	State (num = 3)
x	3
  - Select line in Error Trace to show its value here.
  - Spec Status: parsed

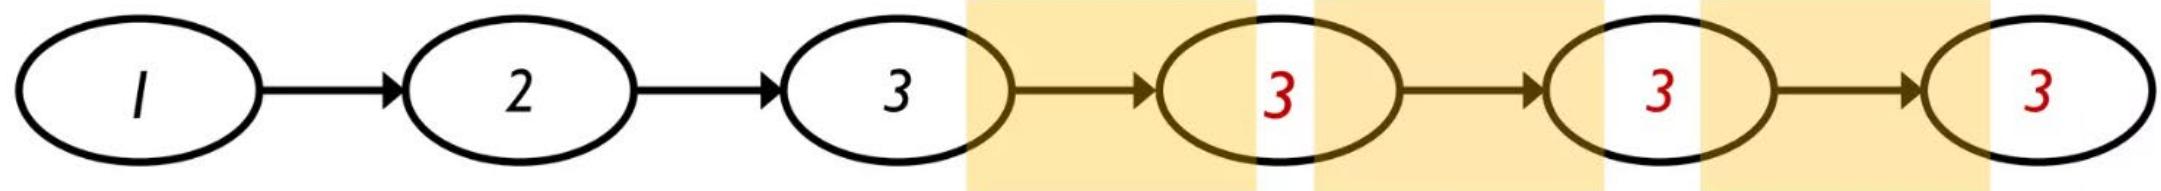
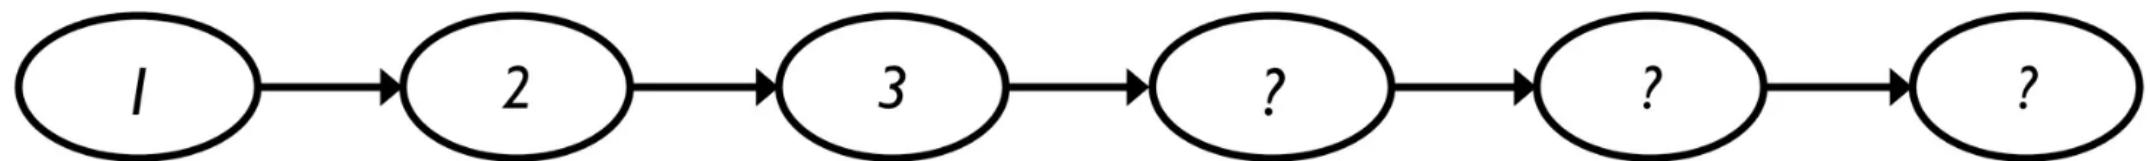
Annotations in red:

- A red arrow points from the text "Look at the error" to the "Deadlock reached." message in the TLC Errors window.
- A red bracket on the right side of the Error-Trace table groups the three rows under each column header.
- Two red circles highlight the values "2" and "3" in the "Value" column of the Error-Trace table.
- A red bracket on the left side of the Error-Trace table groups the first three rows under each column header.
- A red arrow points from the text "Look at the trace" to the "x" entry in the third row of the Error-Trace table.

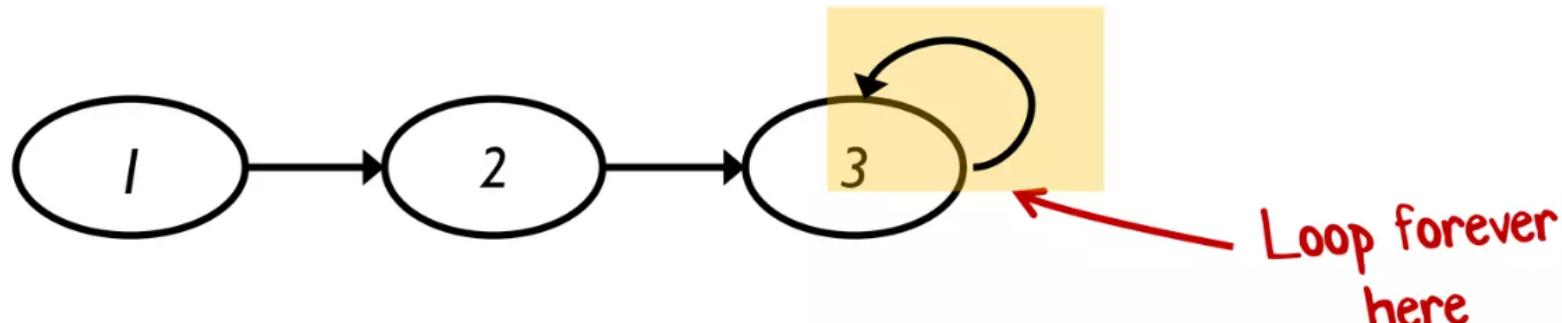
# Deadlock in TLA+



- Infinite time in TLA+



# Count to three, updated



Init ==  $x=1$

Step1 ==  $x=1 \wedge x'=2$

Step2 ==  $x=2 \wedge x'=3$

Done ==  $x=3 \wedge \text{UNCHANGED } x$

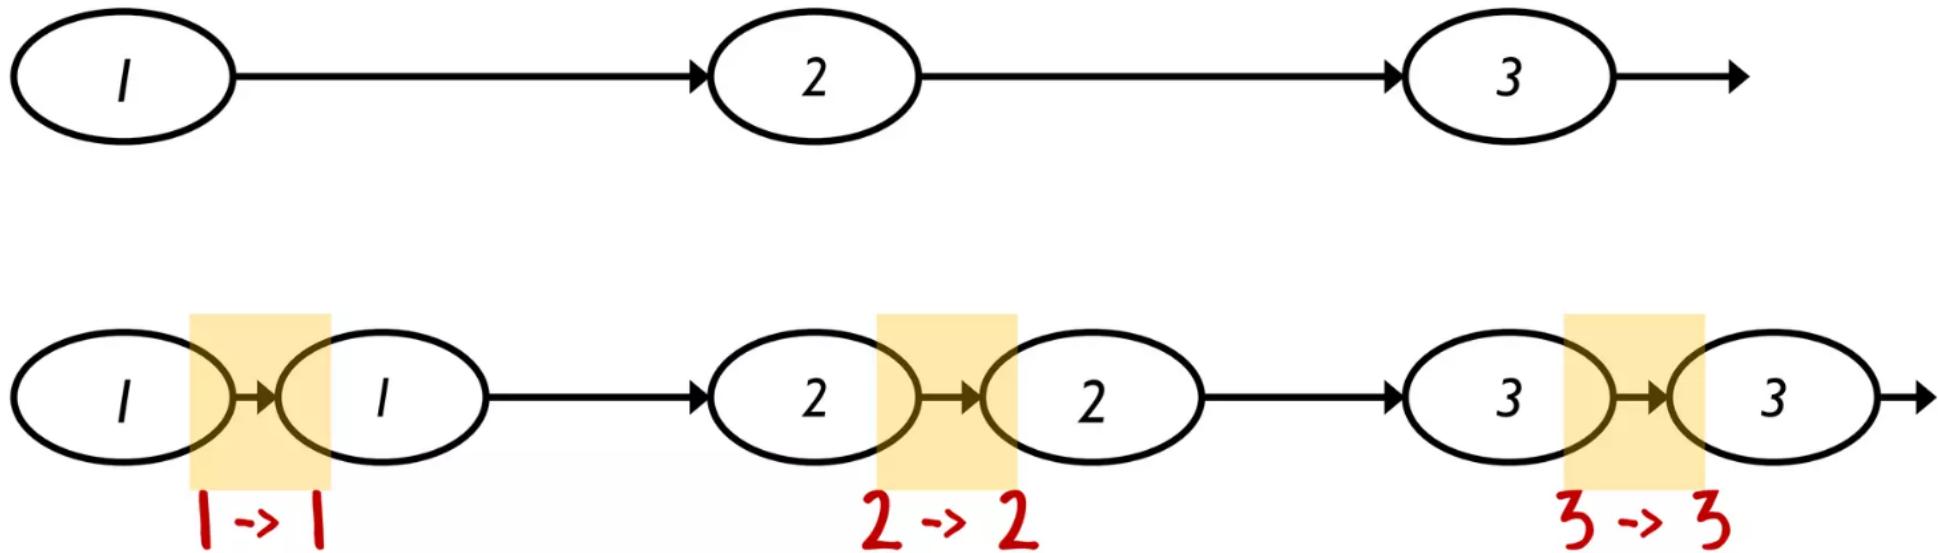
Next == Step1  $\vee$  Step2  $\vee$  Done

If we run this now, we no longer get a deadlock



Add a final step to the model

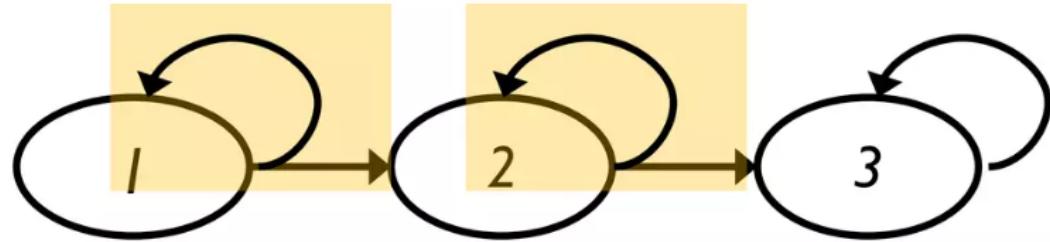
# Doing nothing is always an option!



Staying in the same state is  
almost always a valid state transition!

The TLA+ term for  
this is "stuttering"

# Count to three, with stuttering



Init ==  $x=1$

Step1 ==  $x=1 \wedge x'=2$

Step2 ==  $x=2 \wedge x'=3$

Done ==  $x=3 \wedge \text{UNCHANGED } x$

Next == Step1  $\vee$  Step2  $\vee$  Done  $\vee$  **UNCHANGED  $x$**

With added  
stuttering

# The power of temporal properties

- A **property** applies to the whole system over time
  - Not just to individual states
- Checking these properties is important
  - Humans are bad at this
  - Programming languages are bad at this too
  - TLA+ can help with this!

# Properties in TLA+

- Always true
  - For all tests,  $x > 0$
- Eventually true
  - At some point in time,  $x = 2$
- Eventually always
  - Eventually becomes true (done) and stays there (done)
  - $x$  eventually becomes 3 and then stays there
- Leads to
  - If  $x$  ever becomes 2, then it will become 3 later

# Properties for “count to three”

In English	Formally	In TLA+
$x$ is always $> 2$	Always ( $x > 0$ )	$[] (x > 0)$
At some point $x$ is 2	Eventually ( $x = 2$ )	$\langle\rangle (x = 2)$
$x$ eventually becomes 3 and then stays there.	Eventually (Always ( $x = 3$ ))	$\langle\rangle[] (x = 3)$
if $x$ ever becomes 2 then it will become 3 later.	$(x=2)$ leads to $(x=3)$	$(x=2) \sim> (x=3)$

The “leads to” operator

# Adding properties to the script

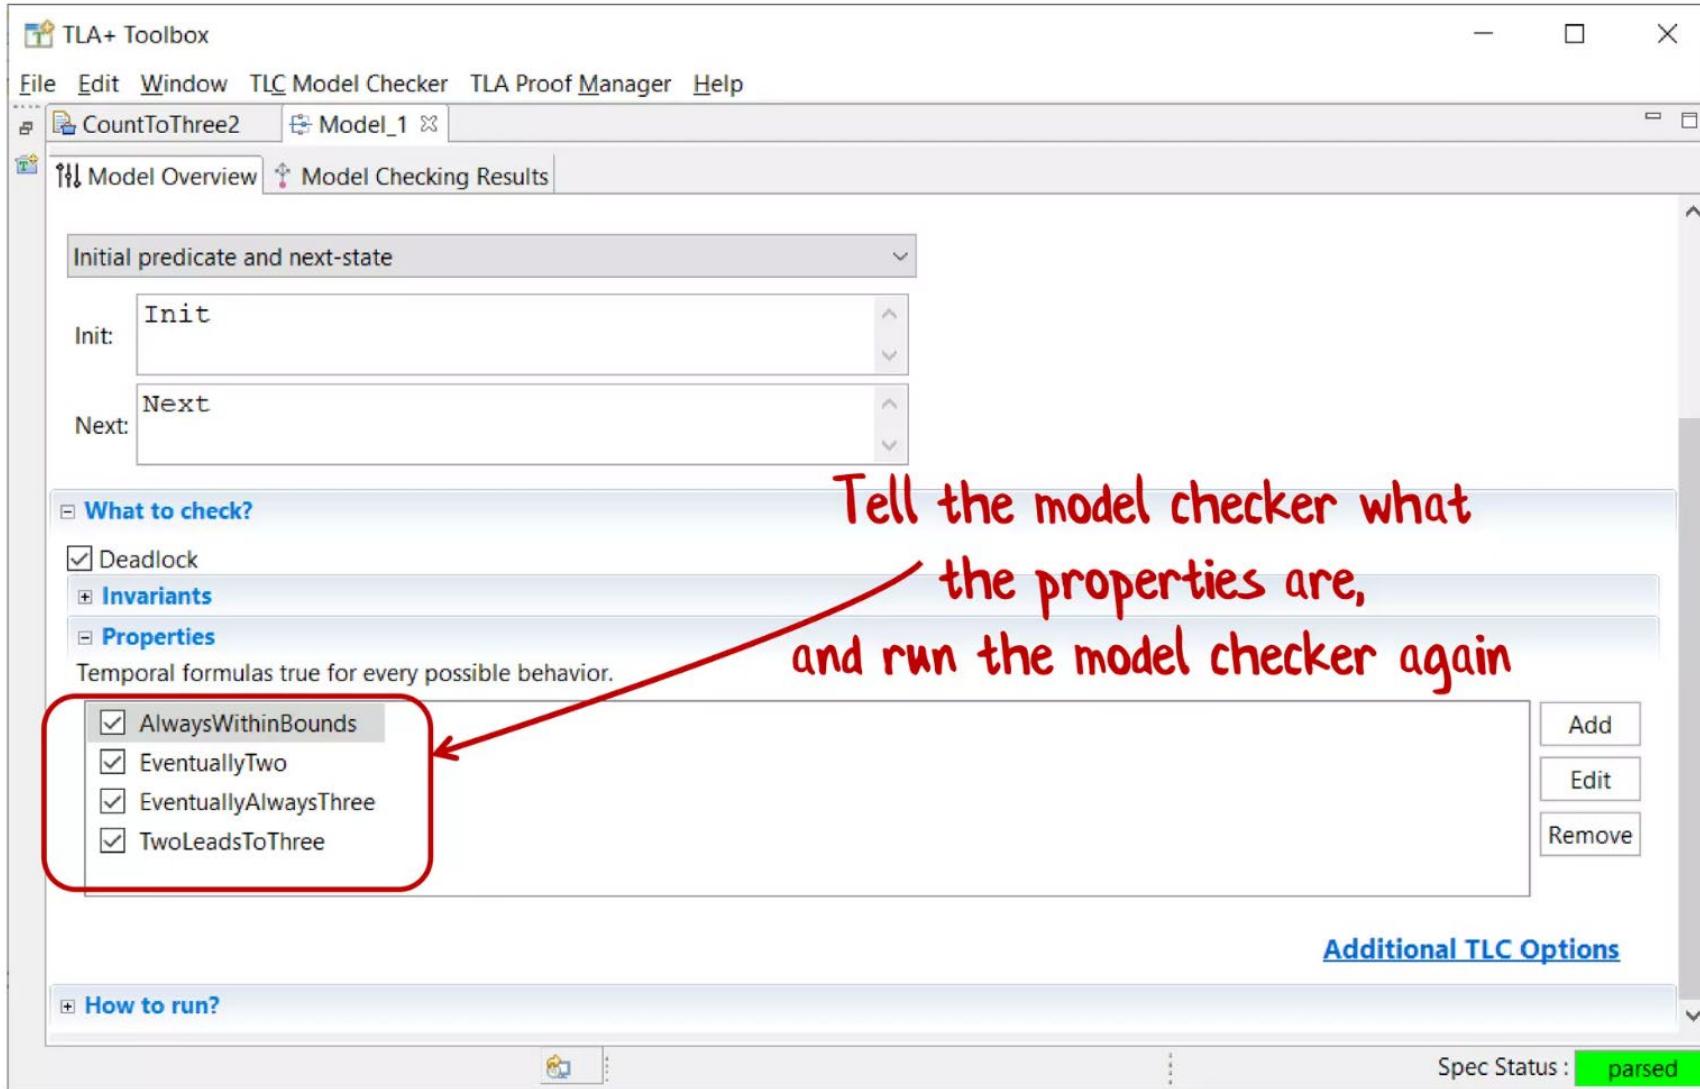
```
\* Always, x >= 1 && x <= 3
AlwaysWithinBounds == [](x >= 1 /\ x <= 3)
```

```
\* At some point, x = 2
EventuallyTwo == <>(x = 2)
```

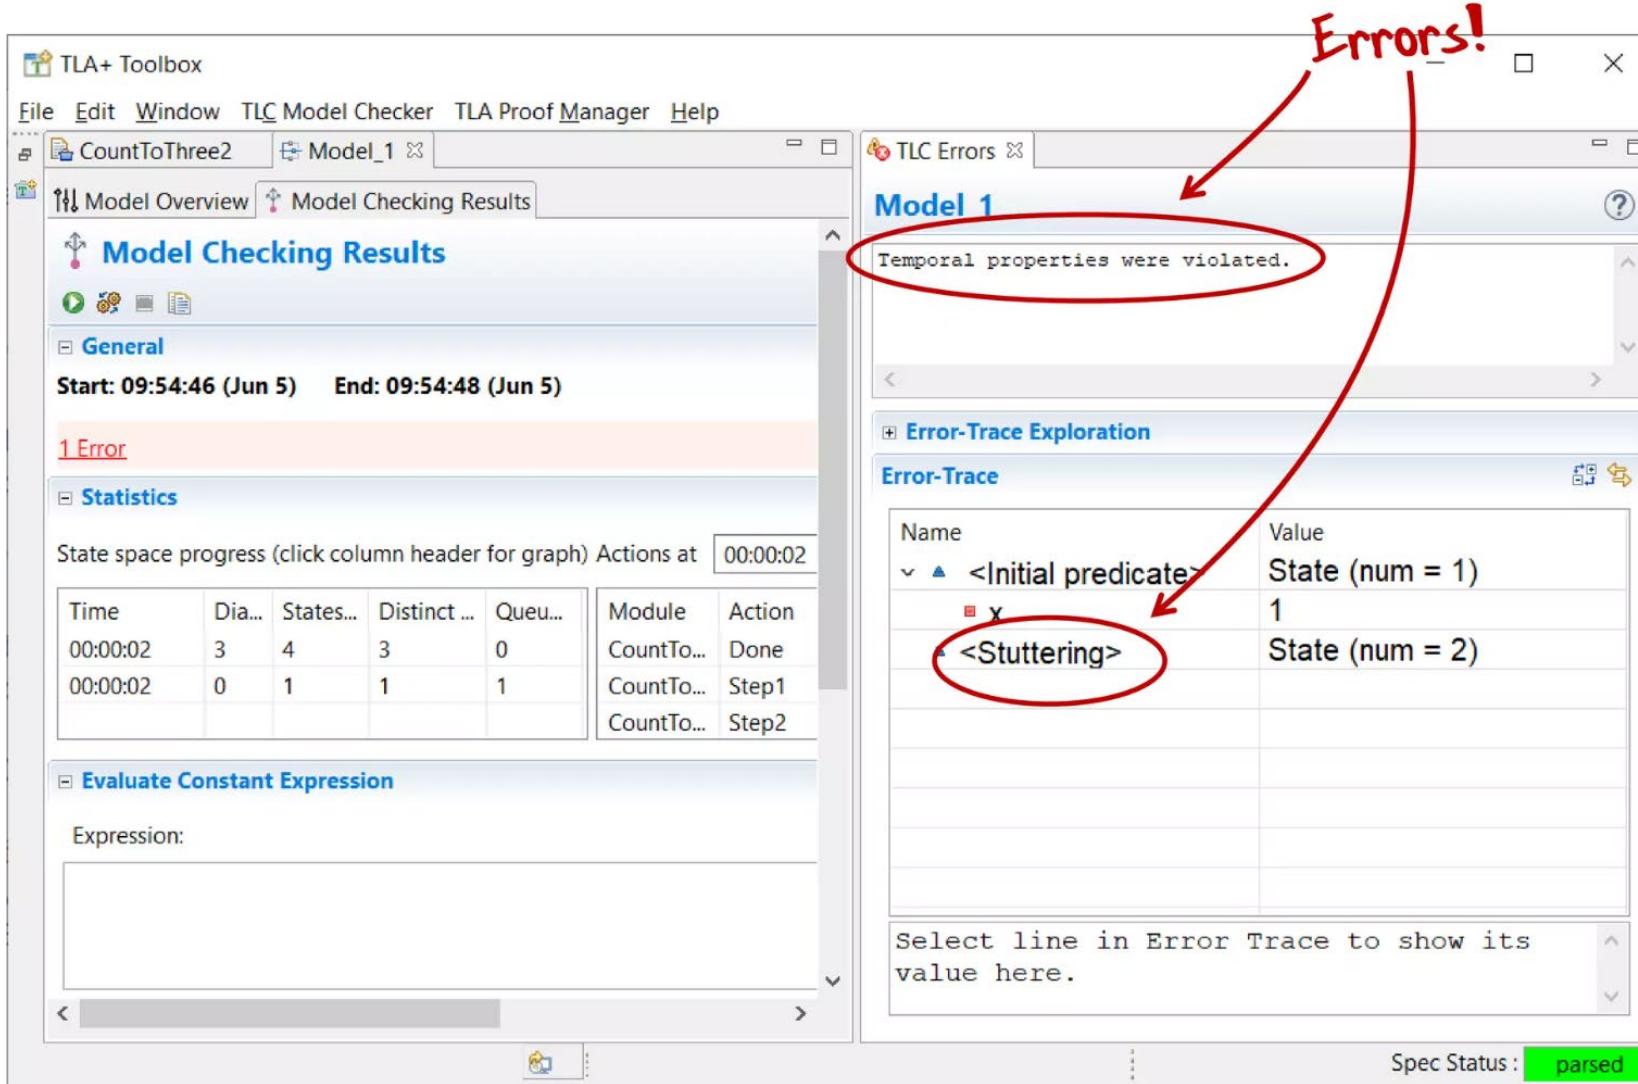
```
\* At some point, x = 3 and stays there
EventuallyAlwaysThree == <>[](x = 3)
```

```
\* Whenever x=2, then x=3 later
TwoLeadsToThree == (x = 2) ~> (x = 3)
```

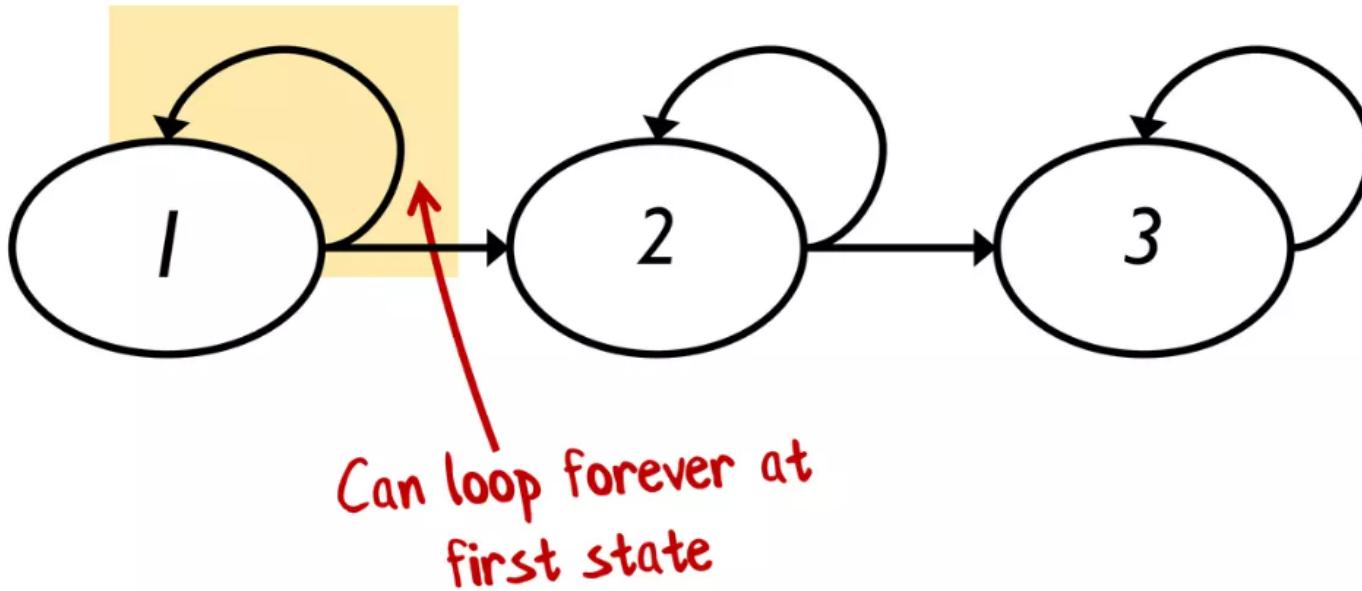
# Adding properties to the script



# Oh no! Model checker says we have errors!



# Stuttering caused a loop!



So only the  
"AlwaysWithinBounds"  
property is true

# Fixing the error

- Make sure every possible transition is followed
- Don't get stuck in an infinite loop

Add fairness! TLA+ can model this

## Refactor #1: change the spec to merge init/next

Init ==  $x=1$

Step1 ==  $x=1 \wedge x'=2$

Step2 ==  $x=2 \wedge x'=3$

Done ==  $x=3 \wedge \text{UNCHANGED } x$

Next == Step1  $\vee$  Step2  $\vee$  Done

Spec = Init  $\wedge$  [] (Next  $\vee$  UNCHANGED  $x$ )

An overall specification

## Refactor #1: change the spec to merge init/next

Init ==  $x=1$

Step1 ==  $x=1 \wedge x'=2$

Step2 ==  $x=2 \wedge x'=3$

Done ==  $x=3 \wedge \text{UNCHANGED } x$

Next == Step1  $\vee$  Step2  $\vee$  Done

Spec = Init  $\wedge$  [] (Next  $\vee$  UNCHANGED  $x$ )



Init AND then always  
some state transition

## Refactor #1: change the spec to merge init/next

```
Init == x=1
```

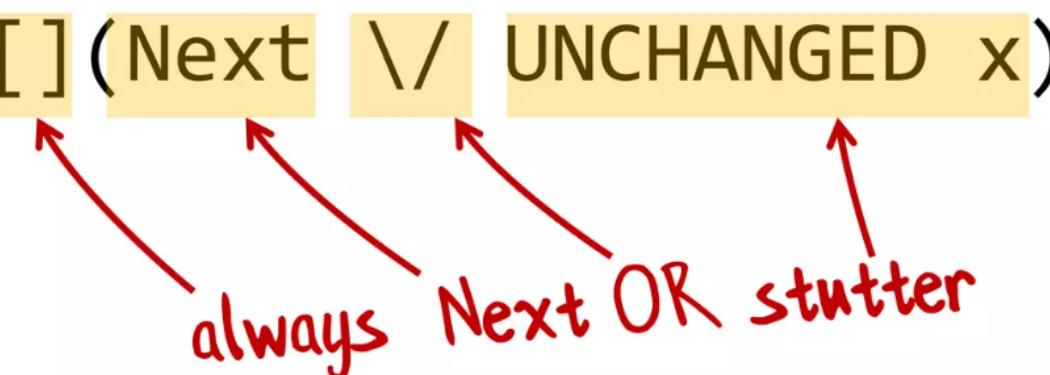
```
Step1 == x=1 /\ x'=2
```

```
Step2 == x=2 /\ x'=3
```

```
Done == x=3 /\ UNCHANGED x
```

```
Next == Step1 \vee Step2 \vee Done
```

```
Spec = Init /\ [](Next \vee UNCHANGED x)
```



## Refactor #2: Use a special syntax for stuttering

Before

```
Spec = Init /\ [](Next \/ UNCHANGED x)
```

After

```
Spec = Init /\ [] [Next]_x
```

Special syntax for  
"Next or stutter on x"

## Refactor #3: Now we can add fairness!

Spec = Init  $\wedge$  [] [Next]\_x

With fairness

Spec = Init  $\wedge$  [] [Next]\_x  $\wedge$  WF\_x(Next)

If "Next" step is available,  
eventually try it, ignoring  
stutter

The new "weak"  
fairness condition



# The complete spec with fairness

```
Init == x=1
Step1 == x=1 /\ x'=2
Step2 == x=2 /\ x'=3
Done == x=3 /\ UNCHANGED x
Next == Step1 \vee Step2 \vee Done
Spec == Init /\ [] [Next]_x /\ WF_x(Next)
```

\\* properties to check

```
AlwaysWithinBounds == [](x >= 1 /\ x <= 3)
EventuallyTwo == <>(x = 2)
EventuallyAlwaysThree == <>[](x = 3)
TwoLeadsToThree == (x = 2) ~> (x = 3)
```

If we run this new spec, all four temporal properties are now true!



# The complete spec with fairness

```
Init == x=1
Step1 == x=1 /\ x'=2
Step2 == x=2 /\ x'=3
Done == x=3 /\ UNCHANGED x
Next == Step1 \vee Step2 \vee Done
Spec == Init /\ [] [Next]_x /\ WF_x(Next)
```

\* properties to check

```
AlwaysWithinBounds == [] (x >= 1 /\ x <= 3)
EventuallyTwo == <> (x = 2)
EventuallyAlwaysThree == <> [] (x = 3)
TwoLeadsToThree == (x = 2) ~> (x = 3)
```

This is where TLA+ differs  
from a programming language

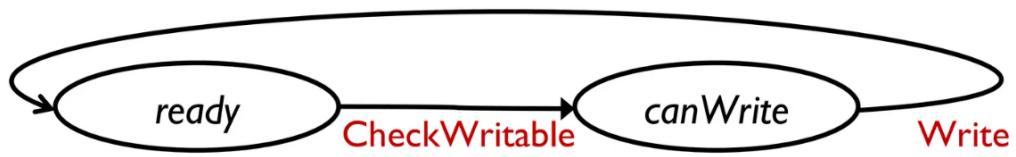
# A more complicated example

- Very exciting: we can count to three!
- What about a more complicated problem?
- What about concurrency?
  - Property checking is where TLA+ is powerful and it can help

# Producer/consumer problem

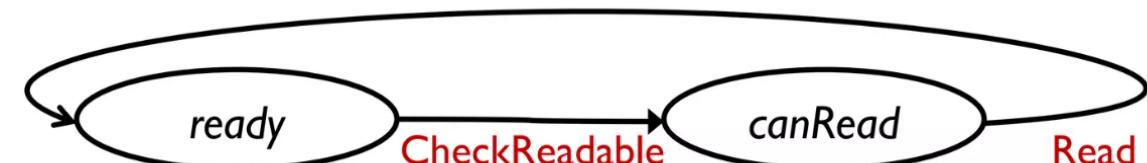
Producer:

- Check if queue is not full
- If true, then write item to queue



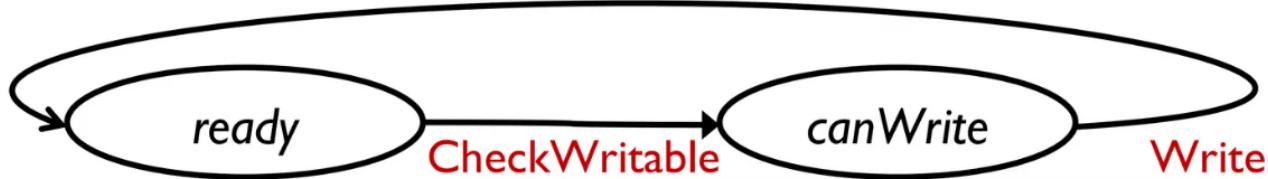
Consumer:

- Check if queue is not empty
- If true, then read item from queue



We're choosing to model this  
as two distinct state transitions,  
not one atomic step

# States for a Producer



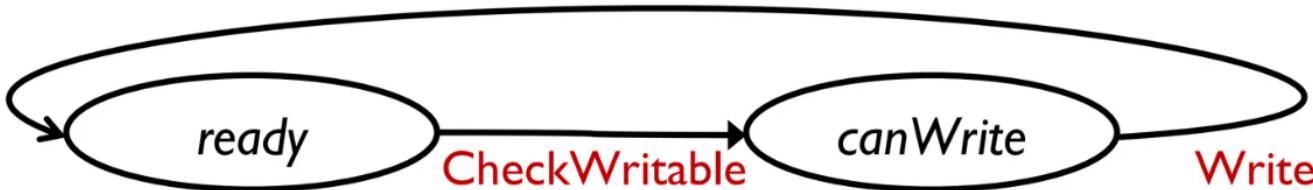
```
def CheckWritable():
    if (queueSize < MaxQueueSize)
        && (producerState = "ready")
    then
        producerState = "canWrite";
```

```
def Write():
    if producerState = "canWrite"
    then
        producerState = "ready";
        queueSize = queueSize + 1;
```

Programming  
language version

Don't care about the  
contents of the queue

# States for a Producer



CheckWritable ==

producerState = "ready"  
/\ queueSize < MaxQueueSize

/\ producerState' = "canWrite" /\* transition  
/\ UNCHANGED queueSize

*TLA+ version*

*Important: All variables must  
be accounted for in an action*

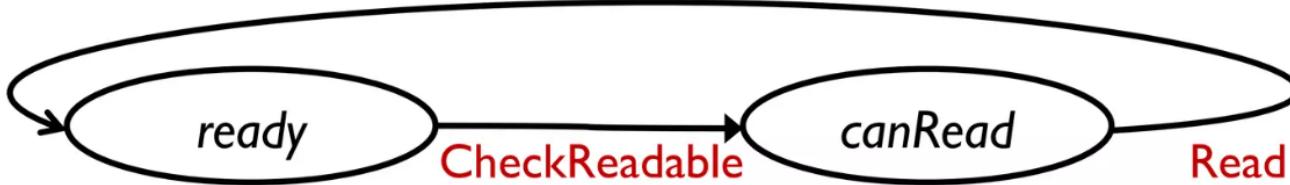
Write ==

producerState = "canWrite"  
/\ producerState' = "ready" /\* transition  
/\ queueSize' = queueSize + 1 /\* push to queue

ProducerAction == CheckWritable \/ Write

*All the valid actions  
for a producer*

# States for a Consumer



CheckReadable ==

```
consumerState = "ready"  
/\ queueSize > 0  
/\ consumerState' = "canRead" /* transition  
/\ UNCHANGED queueSize
```

Read ==

```
consumerState = "canRead"  
/\ consumerState' = "ready" /* transition  
/\ queueSize' = queueSize - 1 /* pop from queue
```

ConsumerAction == CheckReadable ∨ Read

All the valid actions  
for a consumer

# Complete TLA<sup>+</sup> script (1/2)

```
VARIABLES
  queueSize,
  producerState,
  consumerState

  MaxQueueSize == 2 /* can be small

Init ==
  queueSize = 0
  /\ producerState = "ready"
  /\ consumerState = "ready"

CheckWritable ==
  producerState = "ready"
  /\ queueSize < MaxQueueSize
  /\ producerState' = "canWrite"
  /\ UNCHANGED queueSize
  /\ UNCHANGED consumerState

Write ==
  producerState = "canWrite"
  /\ producerState' = "ready"
  /\ queueSize' = queueSize + 1
  /\ UNCHANGED consumerState

ProducerAction ==
  CheckWritable \/ Write
```

# Complete TLA<sup>+</sup> script (2/2)

```
CheckReadable ==  
  consumerState = "ready"  
  /\ queueSize > 0  
  /\ consumerState' = "canRead"  
  /\ UNCHANGED queueSize  
  /\ UNCHANGED producerState
```

```
Read ==  
  consumerState = "canRead"  
  /\ consumerState' = "ready"  
  /\ queueSize' = queueSize - 1  
  /\ UNCHANGED producerState
```

```
ConsumerAction ==  
  CheckReadable \/ Read
```

```
Next ==  
  ProducerAction  
  \/  
  ConsumerAction
```



Non-determinism for free!  
No need to create threads,  
actors, etc

Embedded concurrency!

# Complete TLA<sup>+</sup> script (2/2)

```
CheckReadable ==  
  consumerState = "ready"  
  /\ queueSize > 0  
  /\ consumerState' = "canRead"  
  /\ UNCHANGED queueSize  
  /\ UNCHANGED producerState
```

```
Read ==  
  consumerState = "canRead"  
  /\ consumerState' = "ready"  
  /\ queueSize' = queueSize - 1  
  /\ UNCHANGED producerState
```

```
ConsumerAction ==  
  CheckReadable \/ Read
```

```
Next ==  
  ProducerAction  
  \/ ConsumerAction
```

Question: Should we add  
stuttering here?

# Complete TLA<sup>+</sup> script (2/2)

```
CheckReadable ==  
    consumerState = "ready"  
    /\ queueSize > 0  
    /\ consumerState' = "canRead"  
    /\ UNCHANGED queueSize  
    /\ UNCHANGED producerState
```

```
Read ==  
    consumerState = "canRead"  
    /\ consumerState' = "ready"  
    /\ queueSize' = queueSize - 1  
    /\ UNCHANGED producerState
```

```
ConsumerAction ==  
    CheckReadable \vee Read
```

```
Next ==  
    ProducerAction  
    \vee ConsumerAction  
    \vee (UNCHANGED producerState  
        /\ UNCHANGED consumerState  
        /\ UNCHANGED queueSize)
```

*Doing nothing could  
represent producer or  
consumer crashing!*

# Temporal properties for the producer/consumer

```
AlwaysWithinBounds ==  
  [] (queueSize >= 0  
    /\ queueSize <= MaxQueueSize)
```

- 8 states, no errors
- BUT only for 1 producer and 1 consumer!

# Concurrent version with multiple producers/consumers

- Use the **Plus** in TLA+
- We need
  - A **set** of producers
  - A **set** of consumers
- Use the set-description part of TLA+

```
producers={"p1","p2"}  
consumers={"c1","c2"}
```

# Plus... set theory!

Set theory	Mathematics	TLA+	Programming
e is an element of set S	$e \in S$	$e \in S$	
<b>Define a set</b> by enumeration	{1,2,3}	{1,2,3}	[1,2,3]
<b>Define a set</b> by predicate "p"	$\{ e \in S \mid p \}$	$\{e \in S : p\}$	Set.filter(p)
<b>For all</b> e in Set, some predicate "p" is true	$\forall e \in S : p$	$\forall e \in S : p$	Set.all(p)
<b>There exists</b> e in Set such that some predicate "p" is true	$\exists e \in S : p$	$\exists x \in S : p$	Set.any(p)

# Producer/Consumer Spec, part I

CONSTANT producers, consumers

\\* e.g

\\* 2 producers={"p1","p2"}

\\* 2 consumers={"c1","c2"}

VARIABLES queueSize, producerState, consumerState

MaxQueueSize == 2

Init ==

queueSize = 0

/\ producerState = [p \in producers |-> "ready"]  
  \\* same as {"p1":"ready","p2":"ready"}

/\ consumerState = [c \in consumers |-> "ready"]

# Producer/Consumer Spec, part I

```
CONSTANT producers, consumers  
/* e.g  
/* 2 producers={"p1","p2"}  
/* 2 consumers={"c1","c2"}
```

These are now  
maps/dictionaries

```
VARIABLES queueSize, producerState, consumerState
```

```
MaxQueueSize == 2
```

```
Init ==
```

```
queueSize = 0
```

```
/\ producerState = [p \in producers |-> "ready"]  
/* same as {"p1":"ready","p2":"ready"}  
\ consumerState = [c \in consumers |-> "ready"]
```

For each producer, set  
the state to be "ready"

## Producer/Consumer Spec, part 2

```
CheckWritable(p) ==
  producerState[p] = "ready"
  /\ queueSize < MaxQueueSize
  /\ producerState' =
    [producerState EXCEPT ![p] = "canWrite"]
  /\ UNCHANGED queueSize
  /\ UNCHANGED consumerState
```

Parameterized by a producer

```
CheckWritable(p) ==  
    producerState[p] = "ready"  
    /\ queueSize < MaxQueueSize  
    /\ producerState' =  
        [producerState EXCEPT ![p] = "canWrite"]  
    /\ UNCHANGED queueSize  
    /\ UNCHANGED consumerState
```

Check the state

Update one element of the state map/dictionary only

## Producer/Consumer Spec, part 2

```
CheckWritable(p) ==
    producerState[p] = "ready"
    /\ queueSize < MaxQueueSize
    /\ producerState' =
        [producerState EXCEPT ![p] = "canWrite"]
    /\ UNCHANGED queueSize
    /\ UNCHANGED consumerState

Write(p) ==
    producerState[p] = "canWrite"
    /\ queueSize' = queueSize + 1
    /\ producerState' =
        [producerState EXCEPT ![p] = "ready"]
    /\ UNCHANGED consumerState

ProducerAction ==
    \E p \in producers : CheckWritable(p) \wedge Write(p)
```

## Producer/Consumer Spec, part 2

```
CheckWritable(p) ==
    producerState[p] = "ready"
    /\ queueSize < MaxQueueSize
    /\ producerState' =
        [producerState EXCEPT ![p] = "canWrite"]
    /\ UNCHANGED queueSize
    /\ UNCHANGED consumerState

Write(p) ==
    producerState[p] = "canWrite"
    /\ queueSize' = queueSize + 1
    /\ producerState' =
        [producerState EXCEPT ![p] = "ready"]
    /\ UNCHANGED consumerState

ProducerAction ==
    \E p \in producers : CheckWritable(p) \vee Write(p)
```

Find any producer which has a valid action

# Producer/Consumer Spec, part 3

```
CheckReadable(c) ==  
    consumerState[c] = "ready"  
    /\ queueSize > 0  
    /\ consumerState' =  
        [consumerState EXCEPT !(c) = "canRead"]  
    /\ UNCHANGED queueSize  
    /\ UNCHANGED producerState
```

Parameterized by a consumer

```
CheckReadable(c) ==  
    consumerState[c] = "ready"  
    /\ queueSize > 0  
    /\ consumerState' =  
        [consumerState EXCEPT ![c] = "canRead"]  
    /\ UNCHANGED queueSize  
    /\ UNCHANGED producerState
```

Check the state

Update one element of the state map/dictionary only

# Producer/Consumer Spec, part 3

```
CheckReadable(c) ==
    consumerState[c] = "ready"
    /\ queueSize > 0
    /\ consumerState' =
        [consumerState EXCEPT ![c] = "canRead"]
    /\ UNCHANGED queueSize
    /\ UNCHANGED producerState

Read(c) ==
    consumerState[c] = "canRead"
    /\ queueSize' = queueSize - 1
    /\ consumerState' =
        [consumerState EXCEPT ![c] = "ready"]
    /\ UNCHANGED producerState

ConsumerAction ==
    \E c \in consumers : CheckReadable(c) \vee Read(c)
```

```
CheckReadable(c) ==
    consumerState[c] = "ready"
    /\ queueSize > 0
    /\ consumerState' =
        [consumerState EXCEPT !(c) = "canRead"]
    /\ UNCHANGED queueSize
    /\ UNCHANGED producerState
```

```
Read(c) ==
    consumerState[c] = "canRead"
    /\ queueSize' = queueSize - 1
    /\ consumerState' =
        [consumerState EXCEPT !(c) = "ready"]
    /\ UNCHANGED producerState
```

```
ConsumerAction ==
    \E c \in consumers : CheckReadable(c) \vee Read(c)
```

Find any consumer which has a valid action

# Running the script

- Run the model checker with 2 producers and 2 consumers
  - Use the `AlwaysWithinBounds` property
- There are 38 states
- Error: `Invariant AlwaysWithinBounds is violated!`
  - The design does not work

# Fixing the error

- TLA+ won't tell you how to fix the error
  - You must fix the spec
- Easy to test the fixes
  - Update the spec to use atomic operations (or locks)
  - Re-run the model checker!
  - You gain confidence in your design

# The power of TLA+

- TLA+ can be used to model large concurrent systems
  - Such as distributed systems!
- Examples where TLA+ can help:
  - <https://hillelwayne.com/modeling-deployments/>
  - <https://hillelwayne.com/talks/distributed-systems-tlaplus/>
- Learn more: <https://learntla.com/index.html>

# Why concurrent programs are important?

- They are everywhere nowadays because we all use **distributed systems**
- Distributed systems use the most complex programs
  - systems that span the world
  - serve millions of users
  - and are always available
- Incredibly relevant today as everything is a distributed system!

# Definition of Distributed Systems

- **Distributed system** is a system where multiple processes located on networked computers communicate via messages to achieve a common goal
- "A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.", Leslie Lamport
- Examples: client-server applications, map-reduce, grid computing, peer-to-peer networks, skype, cloud computing, email clients, music streaming, ftp connection, hadoop, web service compositions, video streaming, etc.

# Definition of Distributed Systems

- Distributed system is a system where multiple processes located on networked computers communicate via messages to achieve a common goal
- "**A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.**", **Leslie Lamport**
- Examples: client-server applications, map-reduce, grid computing, peer-to-peer networks, skype, cloud computing, email clients, music streaming, ftp connection, hadoop, web service compositions, video streaming, etc.

# Definition of Distributed Systems

- Distributed system is a system where multiple processes located on networked computers communicate via messages to achieve a common goal
- "A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.", Leslie Lamport
- Examples: client-server applications, map-reduce, grid computing, peer-to-peer networks, skype, cloud computing, email clients, music streaming, ftp connection, hadoop, web service compositions, video streaming, etc.

# Challenges for distributed systems

- **Failure** – cope with partial failure - **reliability**
  - Hardware
  - Network
  - Software – distributed nature, but not only
  - Performance
- **No now**
  - No ordering, no centralized clock
- **Consistency**
- **Consensus**
- Security
- Testing

# Major Challenges (1)

- No global clock, no ordering of events
  - Events happen at different times
  - Different interleaving of events are possible
  - The participants see the events interleaving in different ways
  - Use a logical clock (happens-before)
    - ~ sounds like a memory model is needed

# Major Challenges (2)

## Resilience – Consistency - Consensus

Possible errors due to the concurrent nature of the system:

- Deadlocks
- Livelock/starvation
- Lack of consensus
  - ~ knowing classical synchronization problems might help you solve some particular situation

# 2015: Formal Methods at AWS \*

- Precise description of system in TLA+ (PlusCal language - like c)
  - In 6 large complex real world systems
  - 7 teams
- Found subtle bugs
- Confidence to make **aggressive optimizations** w/o sacrificing correctness
- Use formal specification **to teach** system to new engineers

\* **How Amazon Web Services Uses Formal Methods** by Chris Newcombe et al. (Communications of the ACM, 2015)

# 2015: Formal Methods at AWB

## Applying TLA+ to some of our more complex systems

System	Components	Line count (excl. comments)	Benefit
S3	Fault-tolerant low-level network algorithm	804 PlusCal	Found 2 bugs. Found further bugs in proposed optimizations.
	Background redistribution of data	645 PlusCal	Found 1 bug, and found a bug in the first proposed fix.
DynamoDB	Replication & group-membership system	939 TLA+	Found 3 bugs, some requiring traces of 35 steps
EBS	Volume management	102 PlusCal	Found 3 bugs.
Internal distributed lock manager	Lock-free data structure	223 PlusCal	Improved confidence. Failed to find a liveness bug as we did not check liveness.
	Fault tolerant replication and reconfiguration algorithm	318 TLA+	Found 1 bug. Verified an aggressive optimization.

# 2021: Using Lightweight Formal Methods to Validate a Key-Value Storage Node in Amazon S3

- S3's new ShardStore storage node
  - Built in Rust
  - Crash consistency, concurrency, IO,etc
- Specs alongside the code
  - Reference model spec
  - Decompose correctness checks
    - Sequential correctness
    - Crashes
    - Concurrency
  - Accept weaker correctness guarantees than full formal verification
- Adding continuous validation validation

\* Using Lightweight Formal Methods to Validate a Key-Value Storage Node in Amazon S3 by James Bornholt et al., SOSP2021. <https://www.youtube.com/watch?v=YdxvOPenjWI>

# Conclusion

- Formal verification (model checking) bring guarantees and allows us to check properties
- Formally checking the concurrent code is here to stay
- More engineers will need to write formal specs for their code
- Industry is adapting and using model checkers, especially for newly developed systems
- References:
  - <https://www.youtube.com/watch?v=tqwcz-Yt9gQ>