

Tutorial 03 - Debugging & shared_ptr

Warm up - Spot the bugs (Section 1)

```
void reader(int* foo) {  
    std::cout << *foo;  
    delete foo;  
}  
  
void writer(int* foo) {  
    (*foo)++;  
    delete foo;  
}  
  
void schedule_unsafe () {  
    int* foo = new int;  
  
    std::thread { reader, foo }.detach();  
    std::thread { writer, foo }.detach();  
}
```

1. Double Free (ASan)

```
void reader(int* foo) {  
    std::cout << *foo;  
    delete foo;  
}
```

```
void writer(int* foo) {  
    (*foo)++;  
    delete foo;  
}
```

```
void schedule_unsafe () {  
    int* foo = new int;  
  
    std::thread { reader, foo }.detach();  
    std::thread { writer, foo }.detach();  
}
```

```
=====
```

```
==39903==ERROR: AddressSanitizer: attempting double-free on 0x  
    #0 0x5600620a5d0a in operator delete(void*)  
    (truncated)
```

```
freed by thread T1 here:  
    #0 0x5600620a5d0a in operator delete(void*)  
    (truncated)
```

```
SUMMARY: AddressSanitizer: double-free in operator delete(void*)  
    (truncated)
```

2. Use After Free (ASan)

```
void reader(int* foo) {  
    std::cout << *foo;  
    delete foo;  
}
```

```
void writer(int* foo) {  
    (*foo)++;  
    delete foo;  
}
```

```
void schedule_unsafe () {  
    int* foo = new int;
```

```
    std::thread { reader, foo }.detach();  
    std::thread { writer, foo }.detach();  
}
```

```
=====
```

```
==39903==ERROR: AddressSanitizer: heap-use-after-free on address  
    at pc 0x5600620a8209 bp 0x7f69578fdd20 sp 0x7f69578fdd18  
READ of size 4 at 0x602000000010 thread T2  
    #0 0x5600620a8208 in writer(int*)  
    (truncated)
```

```
0x602000000010 is located 0 bytes inside of 4-byte region [0x602000000000, 0x602000000004)  
freed by thread T1 here:  
    #0 0x5600620a5d0a in operator delete(void*)  
    (truncated)
```

```
SUMMARY: AddressSanitizer: heap-use-after-free in writer(int*)  
    (truncated)
```

```
=====
```

3. Uninitialized variable (MSan)

```
void reader(int* foo) {  
    std::cout << *foo;  
    delete foo;  
}  
  
void writer(int* foo) {  
    (*foo)++;  
    delete foo;  
}  
  
void schedule_unsafe () {  
    int* foo = new int;  
  
    std::thread { reader, foo }.detach();  
    std::thread { writer, foo }.detach();  
}
```

<https://i.imgur.com/3wlxtl0.gifv>

std::shared_ptr

Why use it?

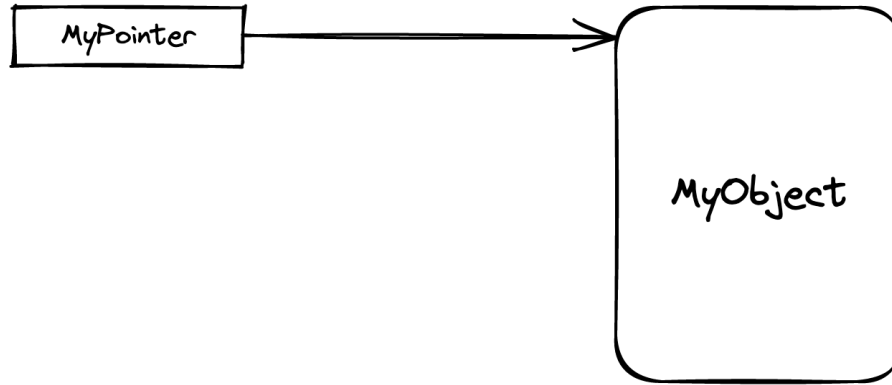
How does it solve the previous problems?

std::shared_ptr

Reference Counting
Section 2.1

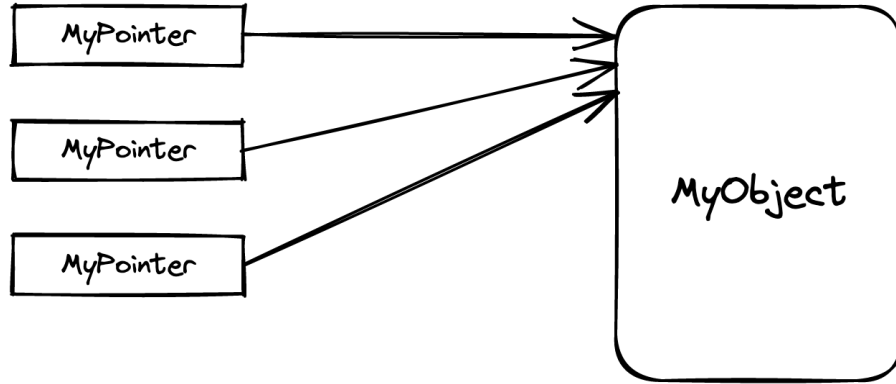
What's Reference Counting?

the object is stored somewhere on the heap



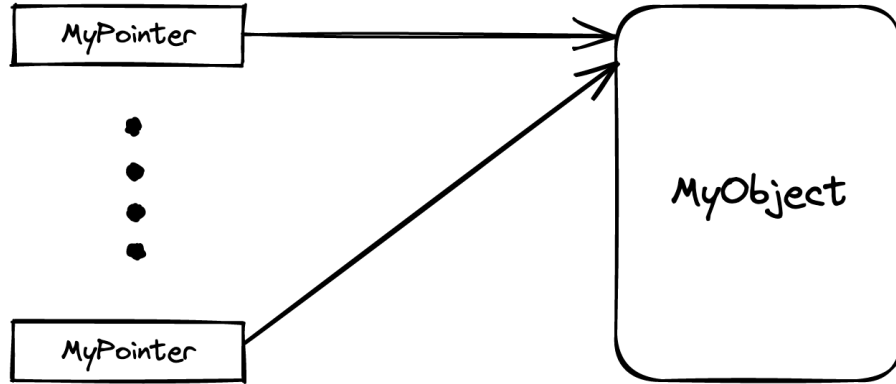
What's Reference Counting?

multiple pointers can point to the same object



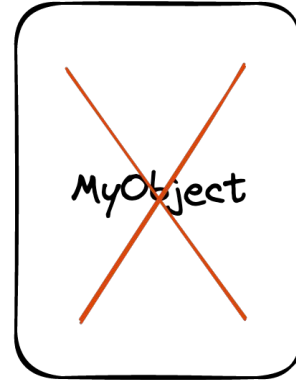
What's Reference Counting?

the pointers can get created and destroyed...



What's Reference Counting?

when the *last pointer* goes away,
the object destructor is called.



What's Reference Counting?

try it yourself

(example 2.1)

4. Is `std::shared_ptr` thread-safe?

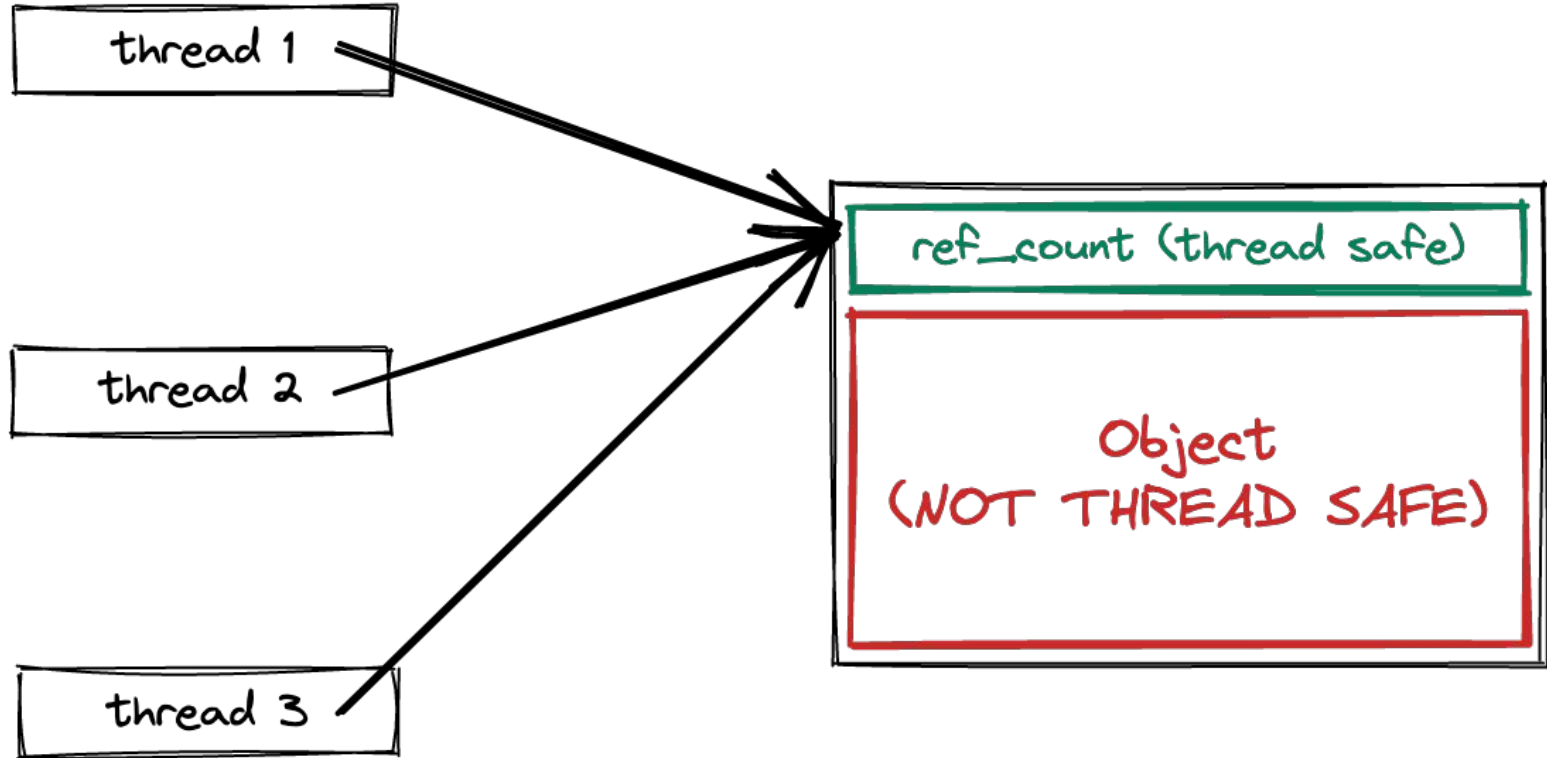
```
void reader2 (std::shared_ptr<int> foo) {
    std::cout << *foo;
} // delete foo

void writer2 (std::shared_ptr<int> foo) {
    (*foo)++;
} // delete foo

void schedule_safe () {
    std::shared_ptr<int> foo { std::make_shared<int>(0) };

    std::thread { reader2, foo }.detach();
    std::thread { writer2, foo }.detach();
} // delete foo
```

Is `std::shared_ptr` thread-safe?



Solution

```
void reader2 (std::shared_ptr<std::atomic<int>> foo) {
    std::cout << foo->load();
}

void writer2 (std::shared_ptr<std::atomic<int>> foo) {
    foo->fetch_add(1);
}

void schedule_safe () {
    auto foo { std::make_shared<std::atomic<int>>(0) };

    std::thread { reader2, foo }.detach();
    std::thread { writer2, foo }.detach();
}
```

std::shared_ptr

The antipattern
(Section 2.4)

Problem 1. Find the bug

```
int main() {
    std::shared_ptr<int> ptr = std::make_shared<int>(0);

    auto reader = std::jthread([](std::shared_ptr<int> ptr) {
        for(int i = 0; i < 100; i++)
            printf("%d\n", *ptr);
    }, ptr);

    auto writer = std::jthread([](std::shared_ptr<int> ptr) {
        for(int i = 0; i < 100; i++)
            *ptr = i;
    }, ptr);
}
```

Problem 1. Data Race on Value (TSan)

```
int main() {
    std::shared_ptr<int> ptr = std::make_shared<int>(0);

    auto reader = std::jthread([](std::shared_ptr<int> ptr) {
        for(int i = 0; i < 100; i++)
            printf("%d\n", *ptr);
    }, ptr);

    auto writer = std::jthread([](std::shared_ptr<int> ptr) {
        for(int i = 0; i < 100; i++)
            *ptr = i;
    }, ptr);
}
```

Recap: What's the fix?

Problem 1. Data Race on Value (TSan)

```
int main() {  
    std::shared_ptr<int> ptr = std:  
  
    auto reader = std::jthread([](s  
    for(int i = 0; i < 100; i++)  
        printf("%d\n", *ptr);  
    }, ptr);  
  
    auto writer = std::jthread([](s  
    for(int i = 0; i < 100; i++)  
        *ptr = i;  
    }, ptr);  
}
```

=====

WARNING: ThreadSanitizer: data race (pid=48593)

Write of size 4 at 0x7b0800000030 by thread T2:

#0 func2()::\$_3::operator()(std::shared_ptr<int>) const
(truncated)

Previous read of size 4 at 0x7b0800000030 by thread T1:

#0 func2()::\$_2::operator()(std::shared_ptr<int>) const
(truncated)

Recap: What's the fix?

Problem 2. Find the bug

```
int main() {
    std::shared_ptr<int> ptr;

    auto reader = std::jthread([](std::shared_ptr<int>& ptr) {
        while(ptr == nullptr)
            ;
        printf("%d\n", *ptr);
    }, std::ref(ptr));

    auto writer = std::jthread([](std::shared_ptr<int>& ptr) {
        for(int i = 0; i < 100; i++)
            ptr = std::make_shared<int>(i);
    }, std::ref(ptr));
}
```

Problem 2. Data Race on Pointer (TSan)

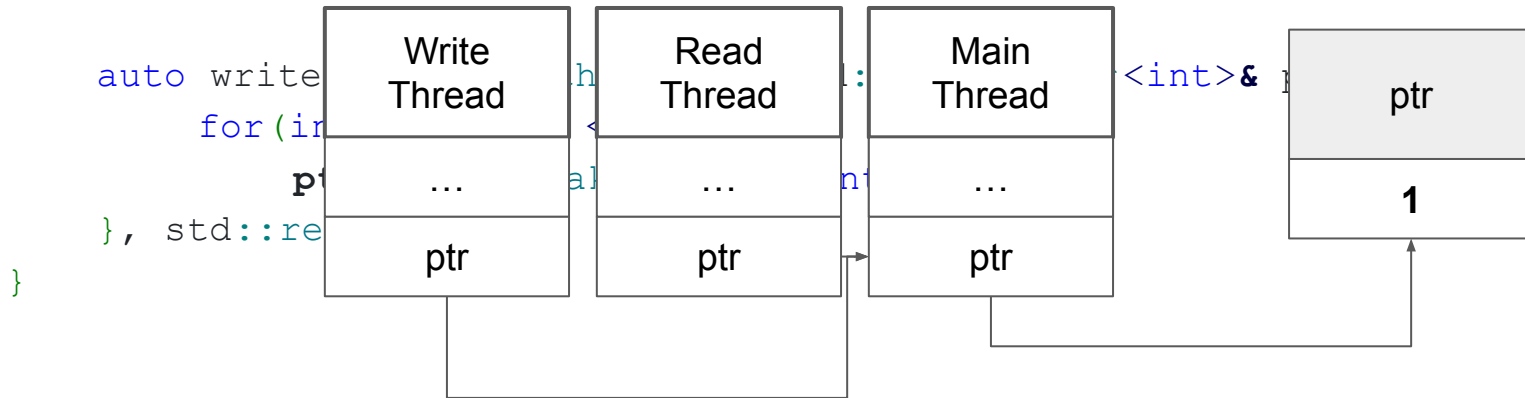
```
int main() {
    std::shared_ptr<int> ptr;

    auto reader = std::jthread([](std::shared_ptr<int>& ptr) {
        while (ptr == nullptr)
            ;
        printf("%d\n", *ptr);
    }, std::ref(ptr));

    auto writer = std::jthread([](std::shared_ptr<int>& ptr) {
        for(int i = 0; i < 100; i++)
            ptr = std::make_shared<int>(i);
    }, std::ref(ptr));
}
```

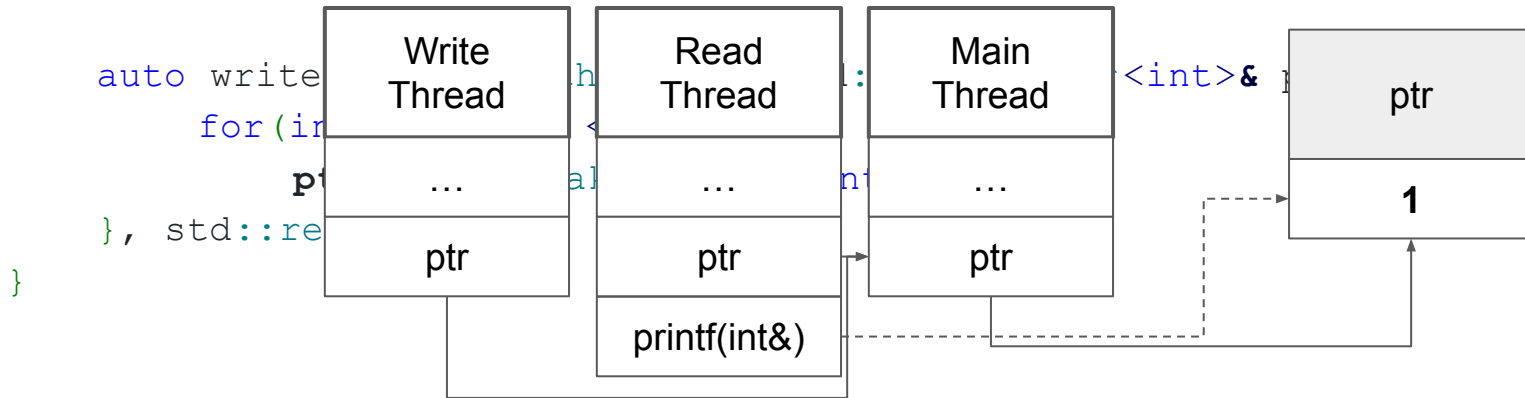
Problem 2. Data Race on Pointer (TSan)

```
int main() {  
    std::shared_ptr<int> ptr;  
  
    auto reader = std::jthread([](std::shared_ptr<int>& ptr) {  
        while(ptr == nullptr)  
            ;  
        printf("%d\n", *ptr);  
    }, std::ref(ptr));  
  
    auto writer = std::jthread([](std::shared_ptr<int>& ptr) {  
        for(int i = 0; i < 10; i++)  
            ptr = std::make_shared<int>(i);  
    }, std::ref(ptr));  
}
```



Problem 2. Data Race on Pointer (TSan)

```
int main() {  
    std::shared_ptr<int> ptr;  
  
    auto reader = std::jthread([](std::shared_ptr<int>& ptr) {  
        while(ptr == nullptr)  
            ;  
        printf("%d\n", *ptr);  
    }, std::ref(ptr));  
  
    auto writer = std::jthread([](std::shared_ptr<int>& ptr) {  
        for(int i = 0; i < 10; i++)  
            ptr = std::make_shared<int>(i);  
    }, std::ref(ptr));  
}
```

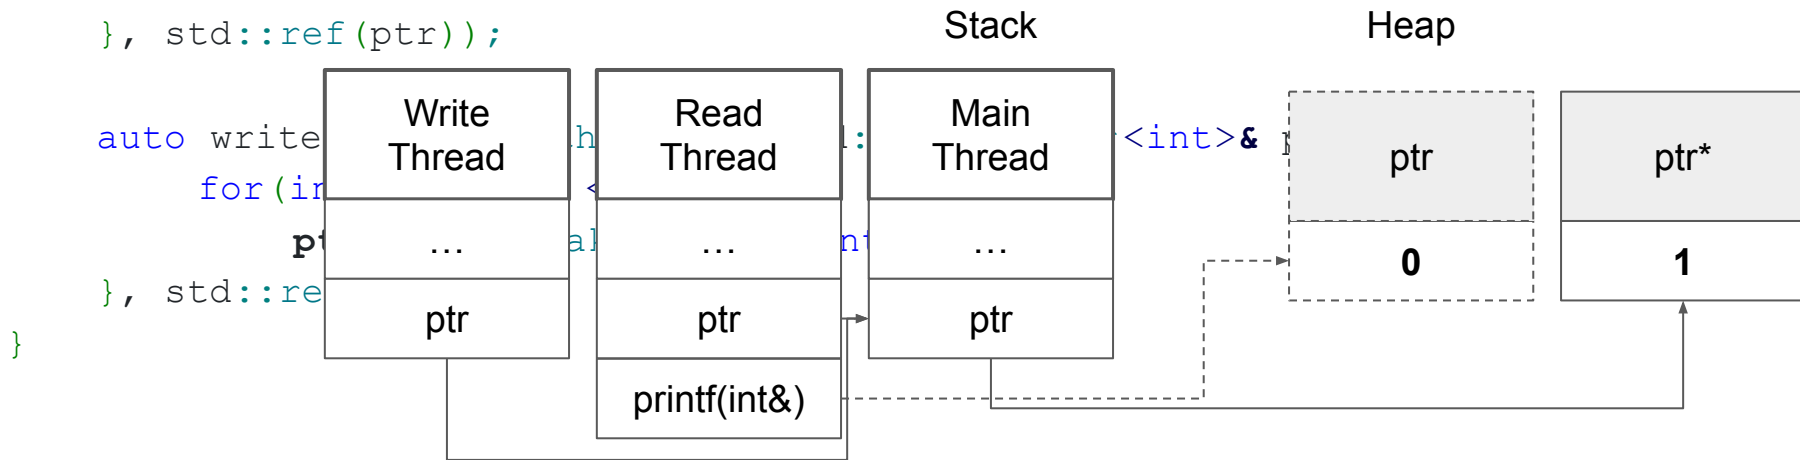


Problem 2. Data Race on Pointer (TSan)

```
int main() {
    std::shared_ptr<int> ptr;

    auto reader = std::jthread([](std::shared_ptr<int>& ptr) {
        while(ptr == nullptr)
            ;
        printf("%d\n", *ptr);
    }, std::ref(ptr));

    auto writer = std::jthread([](std::shared_ptr<int>& ptr) {
        for(int i = 0; i < 10; i++)
            ptr = std::make_shared<int>(i);
    }, std::ref(ptr));
}
```



Problem 2. Data Race on Pointer (TSan)

```
int main() {
    std::shared_ptr<int> ptr;

    auto reader = std::jthread(
        while(ptr == nullptr)
            ;
        printf("%d\n", *ptr);
    }, std::ref(ptr));

    auto writer = std::jthread(
        for(int i = 0; i < 100; i++)
            ptr = std::make_shared<int>(i);
    }, std::ref(ptr));
}
```

=====

WARNING: ThreadSanitizer: data race (pid=48675)

Write of size 8 at 0x7ffc4917a598 by thread T2:

(truncated)

#3 std::shared_ptr<int>::operator=(std::shared_ptr<int>&&)

(truncated)

Previous read of size 8 at 0x7ffc4917a598 by thread T1:

(truncated)

#1 bool std::operator==(int>(std::shared_ptr<int> const&, std::nullptr_t)

(truncated)

Problem 3. Find the bug

```
// Doubly Linked List
struct DLLNode {
    std::shared_ptr<DLLNode> prev;
    std::shared_ptr<DLLNode> next;
};

struct DLL {
    std::shared_ptr<DLLNode> head {};
    std::shared_ptr<DLLNode> tail {};

    void push_front (std::shared_ptr<DLLNode>);
    void push_back (std::shared_ptr<DLLNode>);

    std::shared_ptr<DLLNode> front ();
    std::shared_ptr<DLLNode> back ();
};
```

Problem 3. Circular Reference (ASan)

```
// Doubly Linked List
```

```
struct DLLNode {  
    std::shared_ptr<DLLNode> prev;  
    std::shared_ptr<DLLNode> next;  
};
```

```
struct DLL {  
    std::shared_ptr<DLLNode> head {};  
    std::shared_ptr<DLLNode> tail {};
```

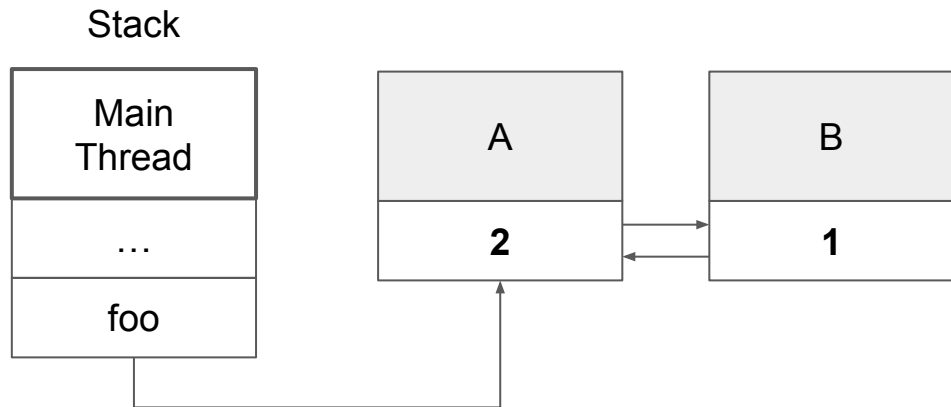
```
void push_front (std::shared_ptr<DLLNode>);
```

```
void push_back (std::shared_ptr<DLLNode>);
```

```
std::shared_ptr<DLLNode> front ();
```

```
std::shared_ptr<DLLNode> back ();
```

```
};
```



Problem 3. Circular Reference (ASan)

```
// Doubly Linked List
```

```
struct DLLNode {  
    std::shared_ptr<DLLNode> prev;  
    std::shared_ptr<DLLNode> next;  
};
```

```
struct DLL {  
    std::shared_ptr<DLLNode> head {};  
    std::shared_ptr<DLLNode> tail {};
```

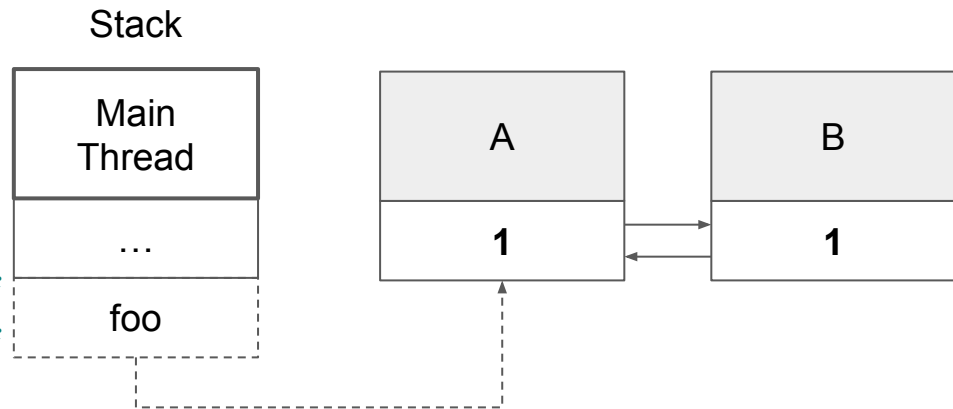
```
void push_front (std::shared_ptr<DLLNode>);
```

```
void push_back (std::shared_ptr<DLLNode>);
```

```
std::shared_ptr<DLLNode> front ();
```

```
std::shared_ptr<DLLNode> back ();
```

```
};
```



Problem 3. Circular Reference

```
// Doubly Linked List
struct DLLNode {
    std::shared_ptr<DLLNode> prev;
    std::shared_ptr<DLLNode> next;
};

struct DLL {
    std::shared_ptr<DLLNode> head {};
    std::shared_ptr<DLLNode> tail {};

    void push_front (std::shared_ptr<DLLNode> node);
    void push_back (std::shared_ptr<DLLNode> node);

    std::shared_ptr<DLLNode> front ();
    std::shared_ptr<DLLNode> back ();
};
```

```
=====
==1==ERROR: LeakSanitizer: detected memory leaks
```

```
Indirect leak of 48 byte(s) in 1 object(s) allocated from:
```

```
#0 0x5639046f60bd in operator new(unsigned long) /root/llvm-pro
#1 0x5639046f87c0 in std::__new_allocator<std::__Sp_counted_ptr<
#2 0x5639046f87c0 in std::allocator_traits<std::allocator<std::
#3 0x5639046f87c0 in std::__allocated_ptr<std::allocator<std::
#4 0x5639046f87c0 in std::__shared_count<(__gnu_cxx::__Lock_pol
#5 0x5639046f87c0 in std::__shared_ptr<DLLNode, (__gnu_cxx::__L
#6 0x5639046f87c0 in std::shared_ptr<DLLNode>::shared_ptr<std::
#7 0x5639046f87c0 in std::shared_ptr<std::enable_if<!is_array<D
#8 0x5639046f87c0 in main /app/example.cpp:28:11
#9 0x7fd5fa270082 in __libc_start_main (/lib/x86_64-linux-gnu/
```

```
Indirect leak of 48 byte(s) in 1 object(s) allocated from:
```

```
#0 0x5639046f60bd in operator new(unsigned long) /root/llvm-pro
#1 0x5639046f8710 in std::__new_allocator<std::__Sp_counted_ptr<
#2 0x5639046f8710 in std::allocator_traits<std::allocator<std::
#3 0x5639046f8710 in std::__allocated_ptr<std::allocator<std::
#4 0x5639046f8710 in std::__shared_count<(__gnu_cxx::__Lock_pol
#5 0x5639046f8710 in std::__shared_ptr<DLLNode, (__gnu_cxx::__L
#6 0x5639046f8710 in std::shared_ptr<DLLNode>::shared_ptr<std::
#7 0x5639046f8710 in std::shared_ptr<std::enable_if<!is_array<D
#8 0x5639046f8710 in main /app/example.cpp:27:11
#9 0x7fd5fa270082 in __libc_start_main (/lib/x86_64-linux-gnu/
```

```
SUMMARY: AddressSanitizer: 96 byte(s) leaked in 2 allocation(s).
```

Summary of Pitfalls

1. Data race on pointer
2. Data race on value
3. Memory leak due to circular reference

std::shared_ptr

Implementation

1. Start with sequential impl

```
template <typename T>
class SharedPtr {
    size_t m_count;
    T* m_ptr;
    // TODO: add additional fields

public:
    SharedPtr(T* ptr) : m_count(1), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        // TODO: synchronise this
        ++m_count;
    }

    ~SharedPtr() {
        // TODO: synchronise this
        if(--m_count == 0)
            delete m_ptr;
    }
};
```


2. Share the count

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    T* m_ptr;
    // TODO: add additional fields
```

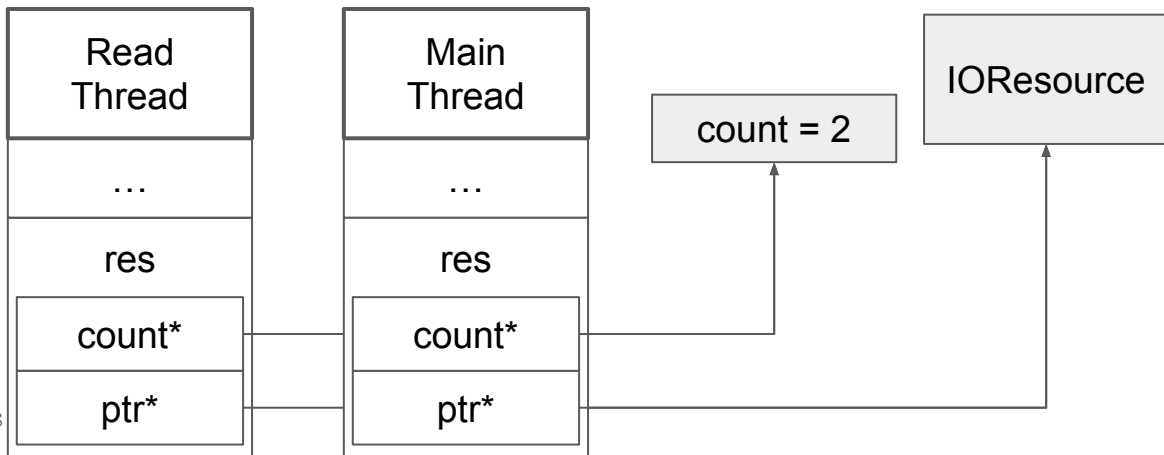
```
public:
```

```
    SharedPtr(T* ptr) : m_count(new int(1)), m_ptr(ptr) {}
```

```
    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        // TODO: synchronise this
        ++(*m_count);
    }
```

```
    ~SharedPtr() {
        // TODO: synchronise this
        if(--(*m_count) == 0)
            delete m_ptr;
```

```
};
```



2. ...vs copy (attempt #1)

```
template <typename T>
```

```
class SharedPtr {
```

```
    size_t* m_count;
```

```
    T* m_ptr;
```

```
    // TODO: add additional fields
```

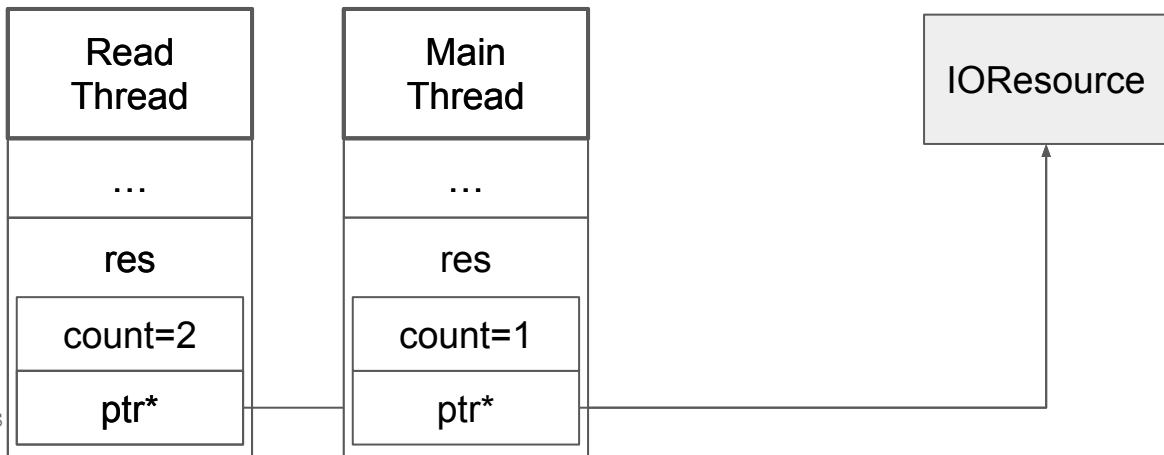
```
public:
```

```
    SharedPtr(T* ptr) : m_count(new int(1)), m_ptr(ptr) {}
```

```
    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {  
        // TODO: synchronise this  
        ++(*m_count);  
    }
```

```
    ~SharedPtr() {  
        // TODO: synchronise this  
        if(--(*m_count) == 0)  
            delete m_ptr;
```

```
};
```



3. Synchronize count (attempt #2)

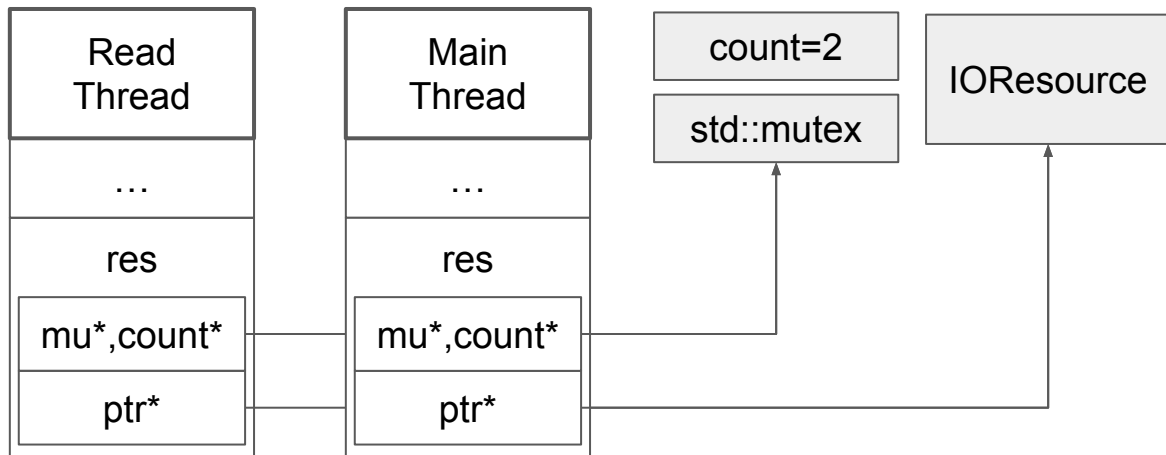
```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;
```

```
public:
```

```
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {}
```

```
    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex),
    m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }
```

```
    ~SharedPtr() {
        auto lk = std::unique_lock { *m_mutex };
        if(--(*m_count) == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    }
};
```



4. Safety?

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex),
    m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }

    ~SharedPtr() {
        auto lk = std::unique_lock { *m_mutex };
        if(--(*m_count) == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    }
};
```

4. Safety?

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex),
    m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }

    ~SharedPtr() {
        auto lk = std::unique_lock { *m_mutex };
        if(--(*m_count) == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    } // m_mutex.unlock()
};
```

Use after free!

4. Safety?

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }

    ~SharedPtr() {
        auto lk = std::unique_lock { *m_mutex };
        if(--(*m_count) == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    } // m_mutex.unlock()
};
```

WARNING: ThreadSanitizer: heap-use-after-free (pid=1)

Atomic read of size 1 at 0x7b0c00000030 by thread T9 (mutexes: write M0):

#0 <null> <null> (output.s+0x709fa)

#1 <null> <null> (output.s+0xd3add)

#2 <null> <null> (output.s+0xd387e)

#3 <null> <null> (libstdc++.so.6+0xe0a3b) (BuildId: 563574e8434b929f2596a1)

Previous write of size 8 at 0x7b0c00000030 by thread T9 (mutexes: write M0):

#0 <null> <null> (output.s+0xd2b6e)

#1 <null> <null> (output.s+0xd3ab0)

#2 <null> <null> (output.s+0xd387e)

#3 <null> <null> (libstdc++.so.6+0xe0a3b) (BuildId: 563574e8434b929f2596a1)

Mutex M0 (0x7b0c00000030) created at:

#0 <null> <null> (output.s+0x7085a)

#1 <null> <null> (output.s+0xd3762)

#2 <null> <null> (libc.so.6+0x24082) (BuildId: 1878e6b475720c7c51969e69ab2)

Thread T9 (tid=11, running) created by thread T8 at:

#0 <null> <null> (output.s+0x52b7d)

#1 <null> <null> (libstdc++.so.6+0xe0d3b) (BuildId: 563574e8434b929f2596a1)

#2 <null> <null> (output.s+0xd39ae)

#3 <null> <null> (libstdc++.so.6+0xe0a32) (BuildId: 563574e8434b929f2596a1)

SUMMARY: ThreadSanitizer: heap-use-after-free (/app/output.s+0x709fa)

=====

Use after free!

4. Unlock before Delete

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex), m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }

    ~SharedPtr() {
        size_t new_count = [this]() {
            auto lk = std::unique_lock { *m_mutex };
            return --(*m_count);
        }(); // m_mutex.unlock()
        if(new_count == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    }
};
```

Can we do better?

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {} // 2 alloc

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex), m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex }; // 1st cache miss + syscall
        ++(*m_count); // 2nd cache miss
    } // 2nd syscall

    ~SharedPtr() {
        size_t new_count = [this]() {
            auto lk = std::unique_lock { *m_mutex }; // 1st cache miss + syscall
            return --(*m_count); // 2nd cache miss
        }(); // 2nd syscall
        if(new_count == 0) {
            delete m_ptr; delete m_mutex; delete m_count; // whee!
        }
    }
};
```


Can we do better?

```
template <typename T>
class SharedPtr {
    size_t* m_count;
    std::mutex* m_mutex;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new size_t(1)), m_mutex(new std::mutex()), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_mutex(other.m_mutex), m_ptr(other.m_ptr) {
        auto lk = std::unique_lock { *m_mutex };
        ++(*m_count);
    }

    ~SharedPtr() {
        size_t new_count = [this]() {
            auto lk = std::unique_lock { *m_mutex };
            return --(*m_count);
        }();
        if(new_count == 0) {
            delete m_ptr; delete m_mutex; delete m_count;
        }
    }
};
```

When `m_count = 1`, only 1 shared ptr exists, no concurrency i.e. no copy constructor run concurrently in the middle of destructor. Hence it is safe to delete

Rewrite It In Atomic 🤪

```
template <typename T>
class SharedPtr {
    std::atomic<size_t>* m_count;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new std::atomic<size_t>(1)), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        m_count->fetch_add(1);
    }

    ~SharedPtr() {
        size_t old_count = m_count->fetch_sub(1);
        if(old_count == 1) {
            delete m_ptr; delete m_count;
        }
    }
};
```

**What is the weakest memory order we can use
for `fetch_add` and `fetch_sub`?**

How about atomics (how hard can it be)?

```
template <typename T>
class SharedPtr {
    std::atomic<size_t>* m_count;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new std::atomic<size_t>(1)), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        m_count->fetch_add(1, std::memory_order_relaxed);
    }

    ~SharedPtr() {
        size_t old_count = m_count->fetch_sub(1, std::memory_order_relaxed);
        if(old_count == 1) {
            delete m_ptr; delete m_count;
        }
    }
};
```

relaxed?

All threads will agree on the modification
order of each individual variable

How about atomics (how hard can it be)?

```
template <typename T>
class SharedPtr {
    std::atomic<size_t>* m_count;
    T* m_ptr;

public:
    SharedPtr(T* ptr) : m_count(new std::atomic<size_t>(1)), m_ptr(ptr) {}

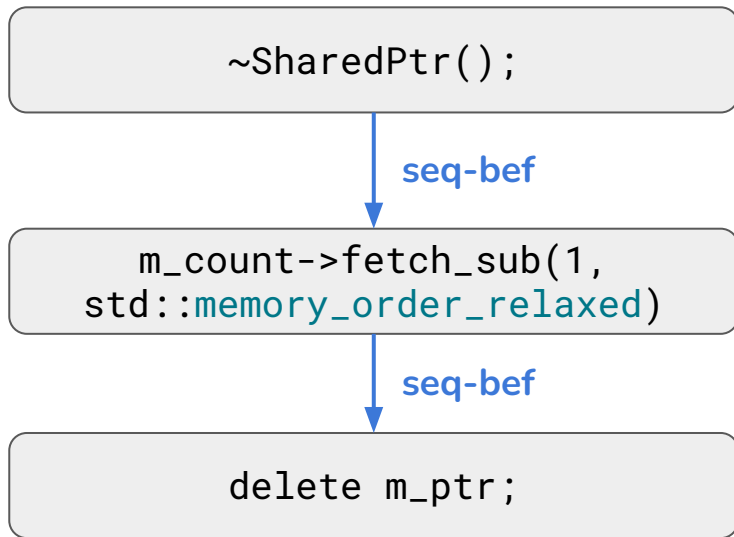
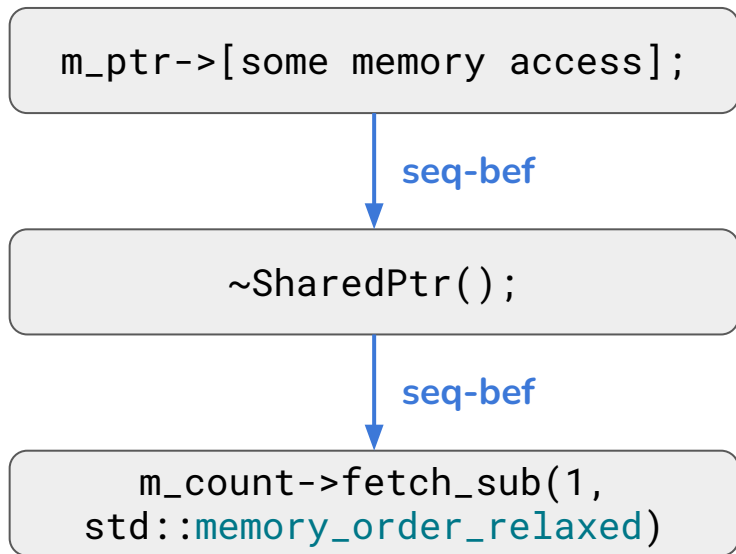
    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        m_count->fetch_add(1, std::memory_order_relaxed);
    }

    ~SharedPtr() {
        size_t old_count = m_count->fetch_sub(1, std::memory_order_relaxed);
        if(old_count == 1) {
            delete m_ptr; delete m_count;
        }
    }
};
```

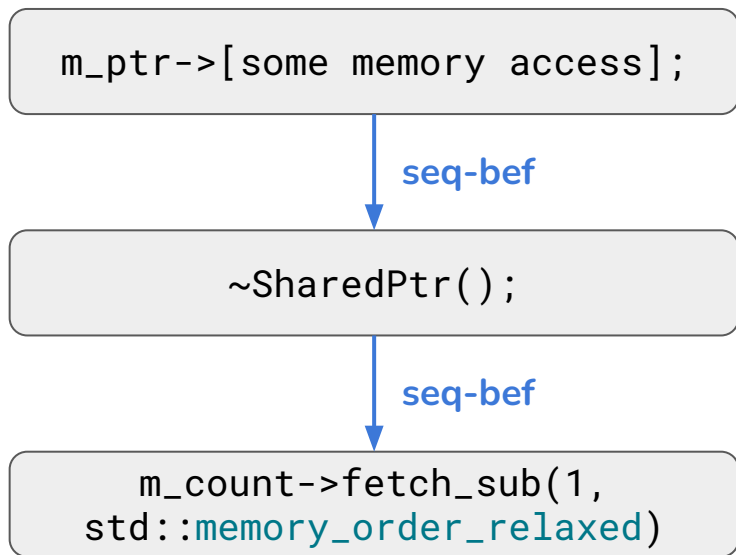
relaxed?

But what about m_count and m_ptr
between threads?

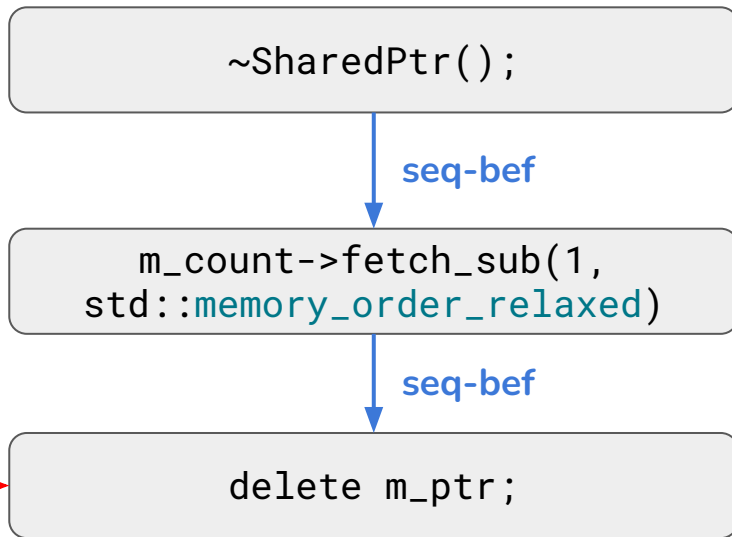
How about atomics (how hard can it be)?



How about atomics (how hard can it be)?



No synchronisation! `delete m_ptr;` not guaranteed to happen-after or happen-before any other memory access to `m_ptr` by the user, concurrent data race is possible



How about atomics (how hard can it be)?

```
template <typename T>
class SharedPtr {
    std::atomic<size_t>* m_count;
    T* m_ptr;

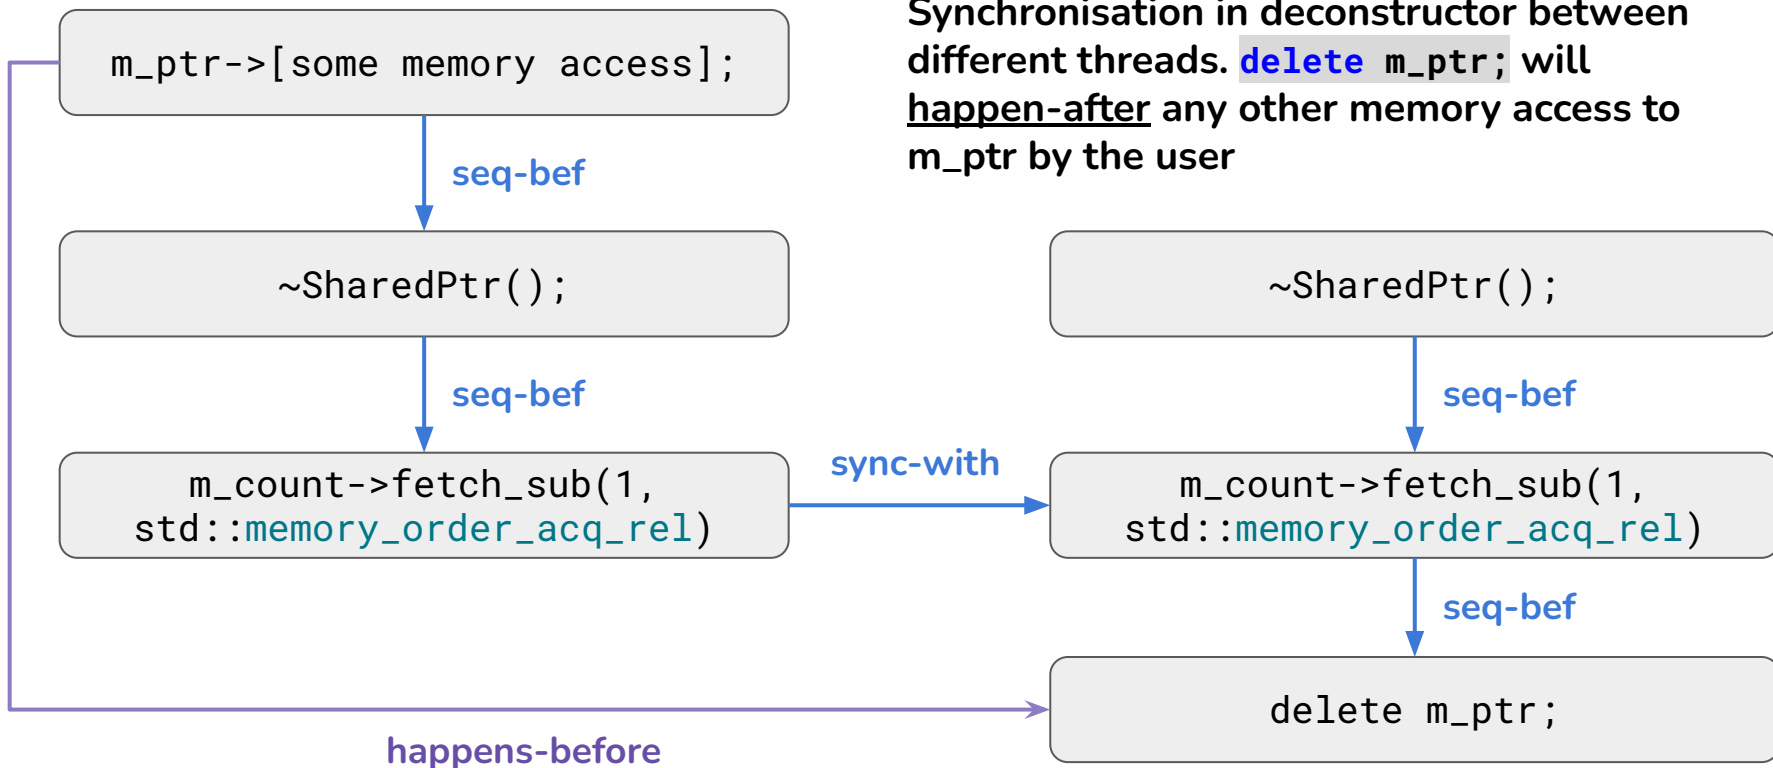
public:
    SharedPtr(T* ptr) : m_count(new std::atomic<size_t>(1)), m_ptr(ptr) {}

    SharedPtr(const SharedPtr& other) : m_count(other.m_count), m_ptr(other.m_ptr) {
        m_count->fetch_add(1, std::memory_order_relaxed);
    }

    ~SharedPtr() {
        size_t old_count = m_count->fetch_sub(1, std::memory_order_acq_rel);
        if(old_count == 1) {
            delete m_ptr; delete m_count;
        }
    }
};
```

Combined acquire release

How about atomics (how hard can it be)?



Have a great weekend ahead!

