


# Tutorial 05 - Go

Modified from Walter's and Sriram's

# Why Go?

# Message Passing as a First-Class Citizen


- **Shared memory  $\Rightarrow$  problem**



- So: no shared memory  $\Rightarrow$  no problem?
  - Hah. (also  $P \Rightarrow Q \nleftrightarrow !P \Rightarrow !Q$ )

## Golang

- However, some **highly scalable** and **correctness-focused** languages focus on **message passing**
  - Some *force* message passing
    - E.g., Elixir



## Elixir

# Today's lesson plan

Go through common mistakes when it comes to using Go

- go threads, wait group, defer, channels

Focusing on Go's concurrency, rather than the language

# Common Mistakes

# What's The Output?

```
func f() {  
    time.Sleep(time.Second)  
    fmt.Println("hello world")  
}
```

```
func main() {  
    for i := 0; i < 10; i++ {  
        go f();  
    }  
}
```

## Spawn, No Join

<https://fsmbolt.comp.nus.edu.sg/z/djPfWK>

# What's The Output?

```
func f() {  
    time.Sleep(time.Second)  
    fmt.Println("hello world")  
}
```

```
func main() {  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        go func() {  
            wg.Add(1)  
            f()  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
}
```

“New feature” : WaitGroup

Done() is safe,  
Add(int) is not

# What's The Output?

```
func f(i int) {  
    time.Sleep(time.Second)  
    fmt.Println(i)  
}
```

```
func main() {  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            f(i)  
            wg.Done()  
        }()  
    }  
    wg.Wait()  
}
```

Variable not captured



# What's The Output?

```
func f(i int) {  
    time.Sleep(time.Second)  
    fmt.Println(i)  
}  
  
func main() {  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func(i int) {  
            defer wg.Done()  
            f(i)  
        }(i)  
    }  
    wg.Wait()  
}
```

“New feature” : defer

# What's The Output?

```
func main() {  
    ch := make(chan int)  
  
    ch <- 1  
  
    fmt.Println(<-ch)  
}
```

<https://play.golang.com/p/XbZjZrMQdcN>

“New feature” : chan

Queue is by  
default unbuffered

How can we solve  
this?

# Solution 1 - Asynchronous Send

```
func main() {  
    ch := make(chan int)  
  
    go func() {  
        ch <- 1  
    }()  
  
    fmt.Println(<-ch)  
}
```

## Solution 2 - Buffered Channel

```
func main() {  
    ch := make(chan int, 1)  
    ch <- 1  
    fmt.Println(<-ch)  
}
```

How much buffer  
fixes deadlock?

## Solution 2 - Why Buffered Channel is always the solution

```
func main() {  
    ch := make(chan int, 1)  
  
    ch <- 1  
  
    ch <- 1  
  
    fmt.Println(<-ch)  
}
```

## Extra: Exploring Channels

- How many possible outputs does this program have, if any?
- What if it is buffered?
- 3 possible outputs
- Buffer does not make a difference since program does not depend on FIFO property

```
1  package main
2
3  import (
4      "fmt"
5  )
6
7  func producer(val int, c chan<- int) {
8      c <- val
9      c <- val
10 }
11
12
13 func main() {
14     c := make(chan int)
15
16     go producer(1, c)
17     go producer(7, c)
18
19     fmt.Println(<-c + <-c)
20 }
```

# What's The Output?

```
func f(i int) {  
    time.Sleep(time.Second)  
    fmt.Println(i)  
}  
  
func main() {  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func(i int) {  
            defer wg.Done()  
            f(i)  
        }(i)  
    }  
    wg.Wait()  
}
```

Var i is just copied around,  
not very useful

What if we want some true  
concurrent operation on  
some variable?

# Message Passing



# Why Not Shared Memory?

Memory Safety Issues:


- Use after free
- Double free
- Data races
- Stale atomic value
- ABA problem
- ...

# How to Solve Shared Memory Issues?

## CS3211 Part 1

1. Test all interleavings, reasoning

## CS3211 Part 2

2. Don't share memory 
3. Prevent unsafe access patterns (WW, WR, RW, ~~RR~~)

# How to Solve Shared Memory Issues?

## CS3211 Part 1

1. Test all interleavings



## CS3211 Part 2

2. Don't share memory **Go channels**
3. Prevent unsafe access patterns (WW, WR, RW, ~~RR~~) **Rust borrow checker**





# Concurrent Counter

...with channels

# Recall: Mutex Impl (Shared Mem)





# Recall: Mutex Impl






1. Lock

# Recall: Mutex Impl



1. Lock
2. Read




# Recall: Mutex Impl



1. Lock
2. Read
3. Modify





# Recall: Mutex Impl






1. Lock
2. Read
3. Modify
4. Write

# Recall: Mutex Impl




1. Lock
2. Read
3. Modify
4. Write
5. Unlock




# Message Passing (Unique Ownership)



Channel






# Message Passing



Channel

1. Read from ch




# Message Passing




Channel

1. Read from ch
2. Modify

# Message Passing





Channel



1. Read from ch
2. Modify
3. Push to ch




# Message Passing



Channel

1. **Read from ch**
2. Modify
3. Push to ch

# Message Passing





Channel


1. Read from ch
2. **Modify**
3. Push to ch



# Message Passing



Channel



1. Read from ch
2. Modify
3. **Push to ch**

# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
  
    for i := 0; i < 10; i++ {  
        go func() {  
            count := <-ch  
  
            count++  
  
            ch <- count  
  
        }()  
    }  
  
    ch <- 0  
  
    fmt.Println(<-ch)  
}
```

# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
  
    for i := 0; i < 10; i++ {  
        go func() {  
            count := <-ch  
            count++  
            ch <- count  
        }()  
    }  
  
    ch <- 0  
  
    fmt.Println(<-ch) // may read count < 10, should wait for workers  
}
```

We went through similar problem earlier on.

Master should wait for workers

# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```

# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```

Defer only runs after the go func() is done here!

1000th goroutine tries to write 1000 to channel


But main is stuck here!  
Not receiving

// A

// B

// C


// D



# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```


Master wait for worker  
to be done, worker wait  
for master take the  
count



# Code - Any problem?

```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    fmt.Println(<-ch)  
    wg.Wait()  
}
```

Master should wait for  
workers



# Code - Any problem?


```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```

// A

// B

// C


// D






# Solution - Break Cycle

```
func main() {  
    ch := make(chan int)  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            /* remove defer */  
            count := <-ch  
            count++  
            wg.Done()  
            ch <- count  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```




# Solution - Buffer

```
func main() {  
    ch := make(chan int, 1) // make buffered channel of size 1  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```



# Solution - More buffer. Will I read less than 1000?

```
func main() {  
    ch := make(chan int, 2) // make buffered channel of size 2  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```




No, because only 1 item passed around, more buffer or not makes no difference.

Can we do better?

# MPMC Queue

# Back to this. What's the “problem”? (performance)

```
func main() {  
    ch := make(chan int, 2) // make buffered channel of size 2  
    var wg sync.WaitGroup  
    for i := 0; i < 10; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            count := <-ch  
            count++  
            ch <- count  
            /* wg.Done() */  
        }()  
    }  
    ch <- 0  
    wg.Wait()  
    fmt.Println(<-ch)  
}
```



# Producer

```
func producer(done chan struct{}, q chan<- int) {  
    for {  
        select {  
            case q <- 1: // keeps incrementing...  
            case <-done: // until stopped (channel closed)  
                return  
        }  
    }  
}
```

# Consumer

```
func consumer(done chan struct{}, q chan int) {  
    sum := 0  
    for {  
        select {  
            case num := <-q:  
                sum += num  
            case <-done:  
                // TODO send local sum  
                return  
        }  
    }  
}
```

# Consumer

```
func consumer(done chan struct{}, q chan int, sumCh chan int) {  
    sum := 0  
    for {  
        select {  
        case num := <-q:  
            sum += num  
        case <-done:  
            sumCh <- sum  
            return  
        }  
    }  
}
```



# Consumer <-> Main

```
func consumer(  
    done chan struct{},  
    q chan int,  
    sumCh chan int) {  
    sum := 0  
    for {  
        select {  
        case num := <-q:  
            sum += num  
        case <-done:  
            sumCh <- sum  
            return  
        }  
    }  
}
```

```
func main() {  
    sumCh := make(chan int)  
  
    // spawn producers & consumers  
  
    sum := 0  
    for subSum := range sumCh {  
        sum += subSum  
    }  
    fmt.Println("Sum: ", sum)  
}
```

**Any problem?**

# Consumer <-> Main

```
func consumer(  
    done chan struct{},  
    q chan int,  
    sumCh chan int) {  
    sum := 0  
    for {  
        select {  
        case num := <-q:  
            sum += num  
        case <-done:  
            sumCh <- sum  
            return  
        }  
    }  
}
```

```
func main() {  
    sumCh := make(chan int)  
  
    // spawn producers & consumers  
    // close(done)  
    sum := 0  
    for subSum := range sumCh {  
        sum += subSum  
    }  
    fmt.Println("Sum: ", sum)  
}
```

**Close the channel 'done'!**  
**Never ends since channel 'sumCh' is open!**

# Solution 1 - Close > Read


```
func consumer(  
    done chan struct{},  
    q chan int,  
    sumCh chan int) {  
    sum := 0  
    for {  
        select {  
        case num := <-q:  
            sum += num  
        case <-done:  
            sumCh <- sum  
            return  
        }  
    }  
}
```

```
func main() {  
    sumCh := make(chan int)  
  
    // spawn producers & consumers  
    // close(done)  
    sum := 0  
    for i := 0; i < NCon; i++ {  
        sum += <-sumCh  
    }  
    close(sumCh)  
    fmt.Println("Sum: ", sum)  
}
```

Read exactly N times

**WARN: May panic if producer sends >1x**  
Idiomatic for writer to close  
But we must ensure only last writer closes

## Solution 2 - Fan-In vs Individual Queue



## Solution 2 - Fan-In vs Individual Queue

```
func consumer(  
    done chan struct{},  
    q chan int,  
    sumCh chan int) {  
    sum := 0  
    for {  
        select {  
        case num := <-q:  
            sum += num  
        case <-done:  
            sumCh <- sum  
            close(sumCh)  
            return  
        }  
    }  
}
```

```
func main() {  
    sumCh := make(chan int, NCon)  
  
    // spawn producers & consumers  
  
    sum := 0  
    for i := 0; i < NCon; i++ {  
        sum += <-sumCh[i]  
    }  
    fmt.Println("Sum: ", sum)  
}
```

# Bonus: Exploring Channels

- What is the final value of this program? (Assume that it correctly compiles)?
- There is data race

```
7 // Increment value in channel
8 func consumer(vals <-chan *int, results chan<- int) {
9     val := <-vals
10    *val++
11 }
12
13
14 func main() {
15     // Number of goroutines
16     const num_threads = 700_000
17
18     // Initialize channels and variables
19     v := 0
20     vals := make(chan *int, num_threads)
21     results := make(chan int)
22
23     // Create all consumers
24     for i := 0; i < num_threads; i++ {
25         go consumer(vals, results)
26     }
27
28     // Fill channel with values
29     for i := 0; i < num_threads; i++ {
30         vals <- &v
31     }
32     // ....
33     // ASSUME we wait correctly for all consumers to finish
34     fmt.Println("Final value: ", v)
35 }
```

# Takeaway

- “Undefined” Behavior
  - When passing args to goroutines, copy arguments (don’t share)
- “Catchable” Errors
  - Channels are synchronous by default. They can block
  - Double close / send after close on a channel panics

Notice it’s mostly deadlocks!

- Deadlocks are observable
  - eg long runtime, user QPS dropping, all threads asleep
- Deadlocks affect liveness, not safety
- This **DOES NOT MEAN** it’s always safe

See you next week!

