

# CS3211 Tutorial 7

Classic Concurrency Problems in C++ & Go

# H2O Problem


i.e. resource management

# 1. H2O Problem

- Hydrogen and oxygen atoms arrive at the factory.
- We **only** want to bond 2 H with 1 O.
- Atoms only start bonding when 2 H and 1 O arrives.
- Other atoms block.
- Only 1 bonding can occur at a time.


Models allocation of specific resources to a process.

Models distributed transactions.



# 1. H<sub>2</sub>O Problem

Incoming atoms




factory

completed

# 1. H<sub>2</sub>O Problem


Incoming atoms



factory


completed

# 1. H<sub>2</sub>O Problem



# 1. H<sub>2</sub>O Problem


Incoming atoms



factory


completed

# 1. H<sub>2</sub>O Problem






# 1. H2O Problem




# 1. H<sub>2</sub>O Problem

Incoming atoms



factory



# Solution 1: Using a Barrier

Any problem here?

```
struct WaterFactory {  
    std::barrier<> barrier;  
    WaterFactory() : barrier{3} {}  
  
    void oxygen(void (*bond)()) {  
        barrier.arrive_and_wait();  
        bond();  
    }  
  
    void hydrogen(void (*bond)()) {  
        barrier.arrive_and_wait();  
        bond();  
    }  
};
```

# Solution 1: Using a Barrier

```
auto oxygen_bond = []() {
    print("Bonding oxygen\n");
    std::this_thread::sleep_for(std::chrono::milliseconds{5});
    print("Done\n");
};

std::thread t1{[&]() { factory.oxygen(oxygen_bond); }};
std::thread t2{[&]() { factory.oxygen(oxygen_bond); }};
std::thread t3{[&]() { factory.oxygen(oxygen_bond); }};
// Bonding oxygen
// Bonding oxygen
// Bonding oxygen
// Done
// Done
// Done
```

Does not ensure correct types! Here we form ozone rather than water

```
struct WaterFactory {
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        barrier.arrive_and_wait();
        bond();
    }

    void hydrogen(void (*bond)()) {
        barrier.arrive_and_wait();
        bond();
    }
};
```

## Solution 2: Using Semaphores

Any problem here?

```
struct WaterFactory {  
    std::counting_semaphore oSem{1}, hSem{2};  
    WaterFactory() : barrier{3} {}  
  
    void oxygen(void (*bond)()) {  
        oSem.acquire(); -> 0  
        bond();  
        oSem.release();  
    }  
  
    void hydrogen(void (*bond)()) {  
        hSem.acquire();  
        bond();  
        hSem.release();  
    }  
};
```

# Solution 2: Using Semaphores

Any problem here? Yes!

Case 1: 1 or 2 H atoms call  
bond() -> we get H<sub>2</sub> or H

Case 2: 1 O atom call bond()  
-> We get O

```
struct WaterFactory {
    std::counting_semaphore oSem{1}, hSem{2};
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        oSem.acquire();
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        hSem.acquire();
        bond();
        hSem.release();
    }
};
```

# Solution 3: Using Barrier + Semaphores

Any problem here?

```
struct WaterFactory {
    std::counting_semaphore oSem{1}, hSem{2};
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        oSem.acquire();->0
        barrier.arrive_and_wait();
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        hSem.acquire();->1->0
        barrier.arrive_and_wait();
        bond();
        hSem.release();
    }
};
```

# Solution 3: Using Barrier + Semaphores

Any problem here? **NO!**

Case 1: there are less than 3 atoms -> atoms will block on barrier

Case 2: There are 3+ H atoms -> extra H will block on hSem

Case 3: There are 2+ O atoms -> extra O will block on oSem


```
struct WaterFactory {
    std::barrier<> oBarrier{1}, hSem{2};
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        Obarrier.arrive_and_wait();
        barrier.arrive_and_wait(); waitgroup
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        barrier.arrive_and_wait();
        bond();
        hSem.release();
    }
};
```



# Go Solution 1:



# Go Solution 1:

```
// Step 1: (Precommit)
//      Receive arrival requests from 2 hydrogen and 1 oxygen atoms
h1 := <-wfd.precomH
h2 := <-wfd.precomH
o := <-wfd.precom0

// Step 2: (Commit)
//      Tell the 3 atoms to start bonding
h1 <- struct{}{}
h2 <- struct{}{}
o <- struct{}{}

// Step 3: (Postcommit)
//      Wait until the 3 atoms have finished before looping
// We re-use the same communication channel as (Commit)
<-h1
<-h2
<-o
```

```
type WaterFactoryWithDaemon struct {
    // Channels for atoms to send their arrival requests
    precomH chan chan struct{}
    precom0 chan chan struct{}
}
```

```
func (wfd *WaterFactoryWithDaemon) hydrogen(bond func()) {
    commit := make(chan struct{}) // Step 1: Create private communication channel
    wfd.precomH <- commit          // Step 2: (Precommit)
    <-commit                        // Step 3: (Commit)
    bond()                        // Step 4: Bond
    commit <- struct{}{}          // Step 5: (Postcommit)
}
```

# Go Solution 1:

## Issues:

- daemon thread is a single point of failure
- potential performance bottleneck due to scheduling decisions

# Go Solution 1:

## Issues:

- daemon thread is a single point of failure
- potential performance bottleneck due to scheduling decisions

Recall: 1 oxygen per H<sub>2</sub>O, why not make oxygen the coordinator?

## Go Solution 2:

```
func (wf *WaterFactoryWithLeader) oxygen(bond func()) {  
    // Step 1: Become leader  
    <-wf.oxygenMutex // For fun, we can use a channel as a mutex  
  
    // Step 2: (Precommit)  
    //     Receive arrival requests from 2 hydrogen atoms  
    h1 := <-wf.precomH  
    h2 := <-wf.precomH  
  
    // Step 3: (Commit)  
    //     Tell the 2 hydrogen atoms to start bonding  
    h1 <- struct{}{}  
    h2 <- struct{}{}  
  
    // Step 4: Bond  
    bond()  
  
    // Step 5: (Postcommit)  
    //     Wait until the 2 hydrogen atoms to finish  
    // We re-use the same communication channel as (Commit)  
    <-h1  
    <-h2  
  
    // Step 6: Step down from being leader  
    wf.oxygenMutex <- struct{}{}  
}
```

```
type WaterFactoryWithLeader struct {  
    oxygenMutex chan struct{}  
    precomH      chan chan struct{}  
}
```

```
func (wf *WaterFactoryWithLeader) hydrogen(bond func()) {  
    commit := make(chan struct{}) // Step 1: Create private communication channel  
    wf.precomH <- commit          // Step 2: (Precommit)  
    <-commit                       // Step 3: (Commit)  
    bond()                       // Step 4: Bond  
    commit <- struct{}{}         // Step 5: (Postcommit)  
}
```


Can still have issues if oxygen holding the “mutex” dies but the probability is lower

# FIFO Semaphore

i.e. starvation-free


4 threads on sem of 2: 2- $\rightarrow$ 1- $\rightarrow$ 0(thread 3 and 4 blocked)- $\rightarrow$ 1

# FIFO Semaphore(=2) Problem Statement



**E.g. Oxygen atoms  
waiting to bond()**


# FIFO Semaphore(=2) Problem Statement



**Semaphores don't  
guarantee order!**




# FIFO Semaphore(=2) Problem Statement




**Even if is sole consumer,  
2 may be sleeping...**

# FIFO Semaphore(=2) Problem Statement




# FIFO Semaphore(=2) Problem Statement




... then another consumer comes

# FIFO Semaphore(=2) Problem Statement



# FIFO Semaphore(=2) Problem Statement



# FIFO using ticket queue

Demo 4




## Demo 4: ticket queue


```
struct FIFOSemaphore {  
    std::atomic<std::ptrdiff_t> next_ticket{1};  
    std::atomic<std::ptrdiff_t> now_serving;  
  
    FIFOSemaphore(std::ptrdiff_t initial_count) : now_serving{initial_count}  
    {}  
  
    void acquire() {  
        auto my_ticket = next_ticket.fetch_add(1);  
        while (now_serving.load() < my_ticket) {}  
    }  
  
    void release() {  
        now_serving.fetch_add(1);  
    }  
};
```




## Demo 4: ticket queue (initial\_count = 2)




## Demo 4: ticket queue (initial\_count = 2)



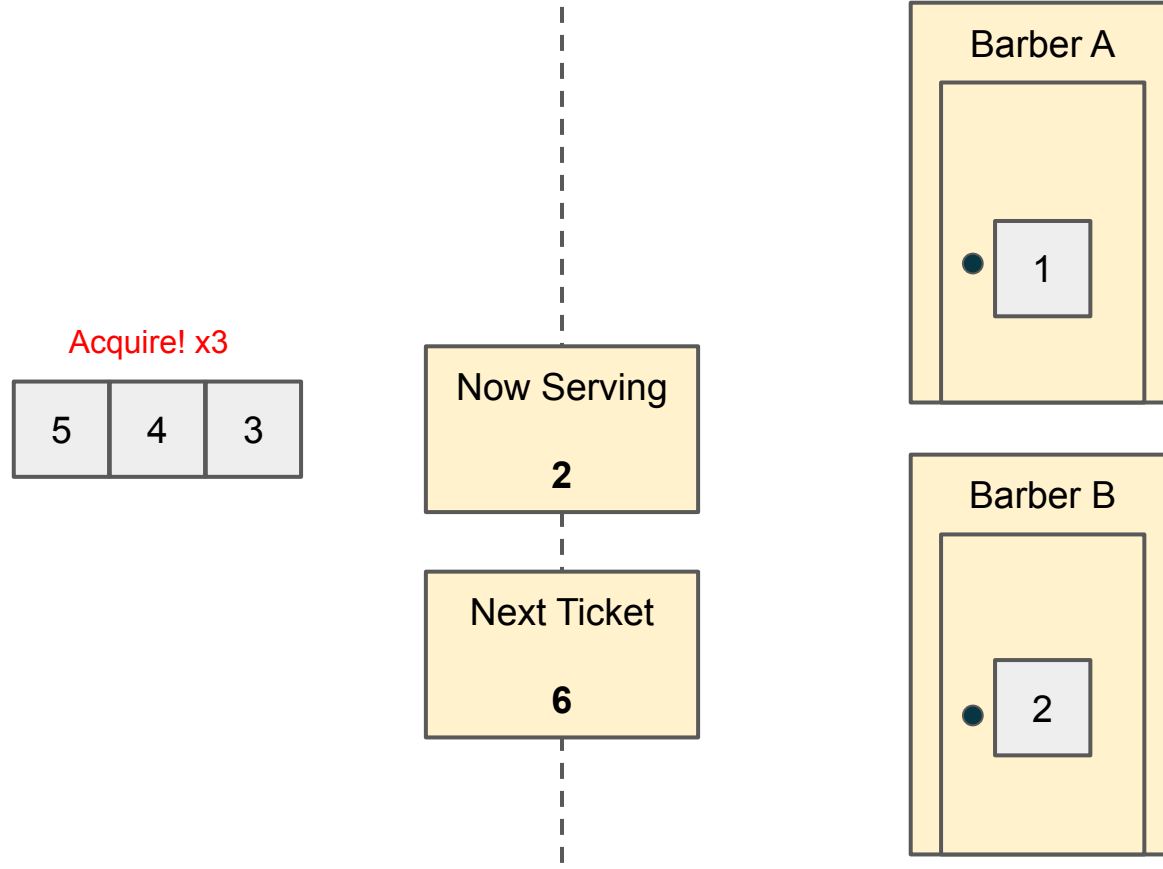
## Demo 4: ticket queue (initial\_count = 2)



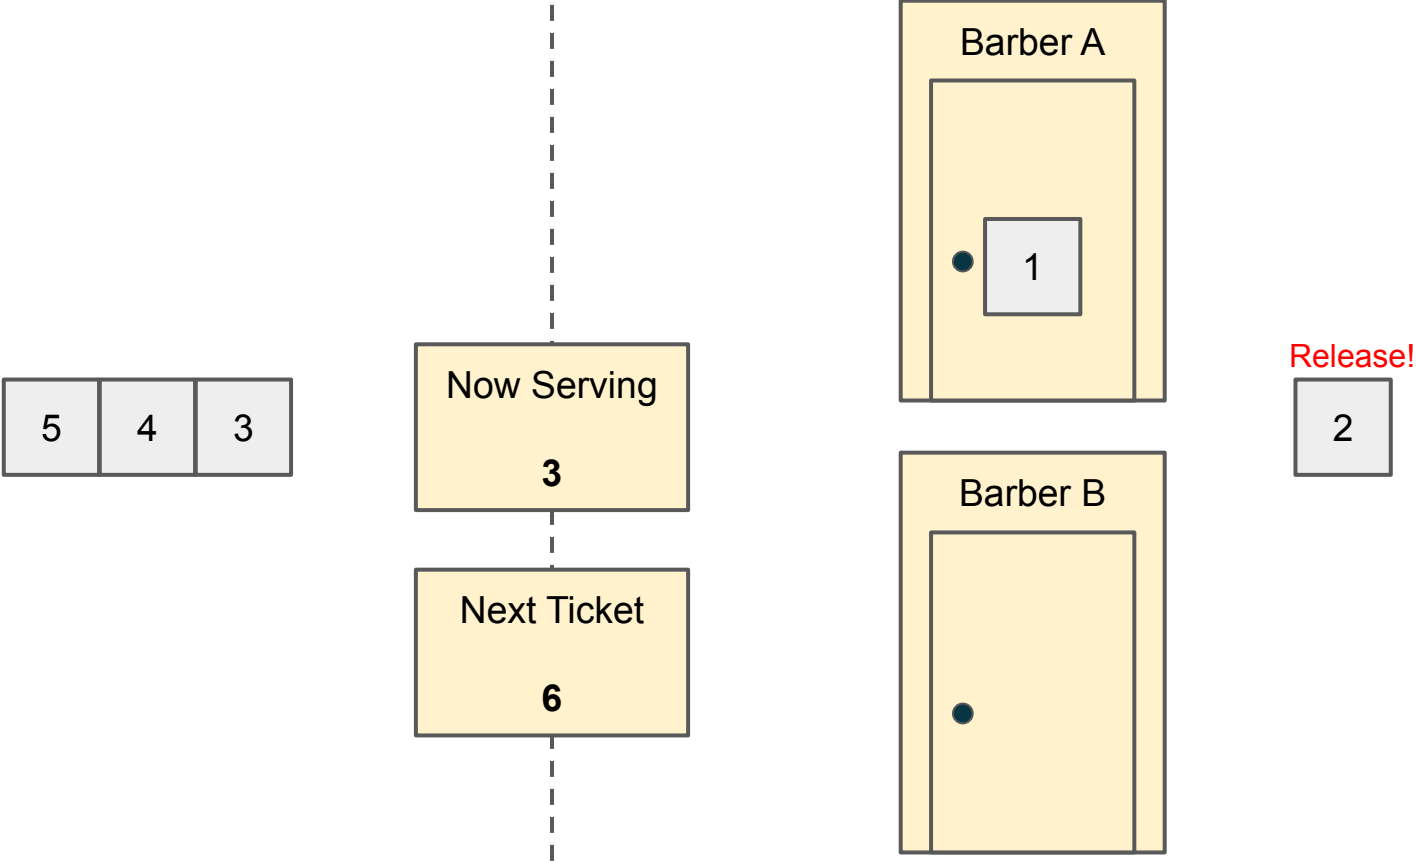
## Demo 4: ticket queue (initial\_count = 2)



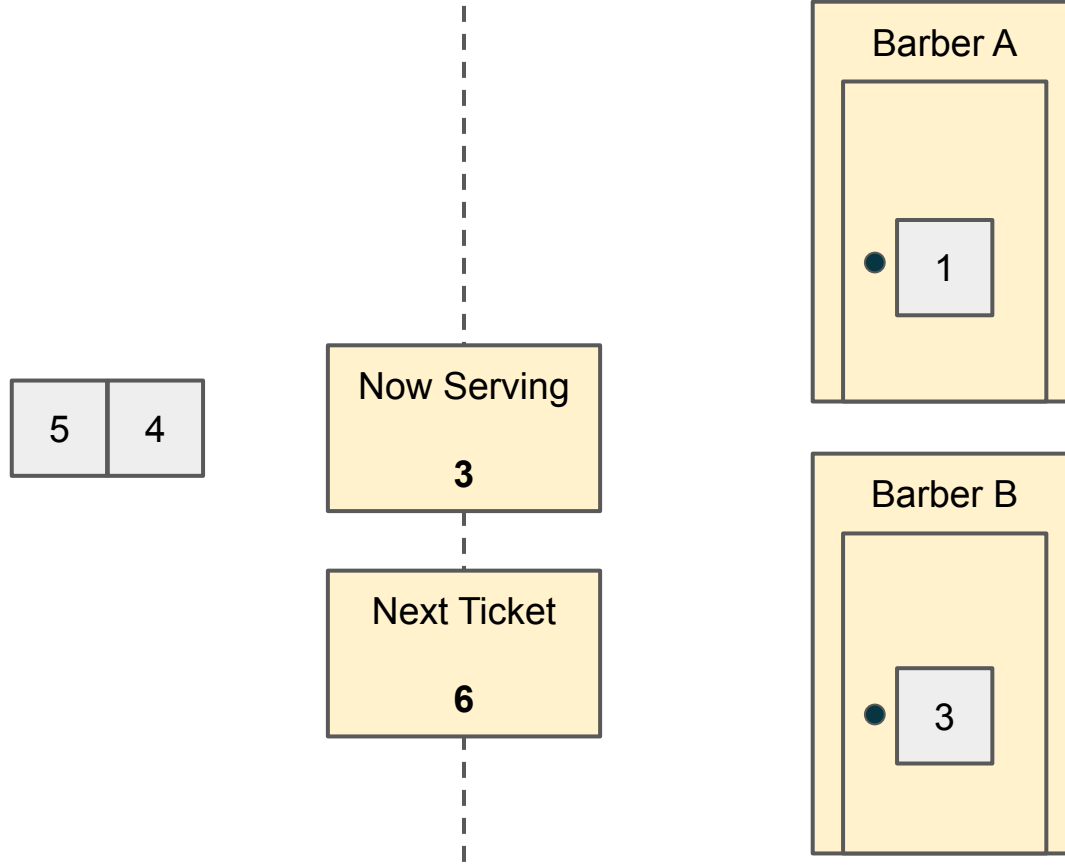
## Demo 4: ticket queue (initial\_count = 2)



# Demo 4: ticket queue (initial\_count = 2)



## Demo 4: ticket queue (initial\_count = 2)



# Task 1: Replace spinwait

Try it yourself!

Use

condition\_variable

or

atomic<T>::wait

```
struct FIFOSemaphore {
    std::atomic<std::ptrdiff_t> next_ticket{1};
    std::atomic<std::ptrdiff_t> now_serving;

    FIFOSemaphore(std::ptrdiff_t initial_count) : now_
    {}

    void acquire() {
        auto my_ticket = next_ticket.fetch_add(1);
        while(now_serving.load() < my_ticket) {}
    }

    void release() {
        now_serving.fetch_add(1);
    }
};
```



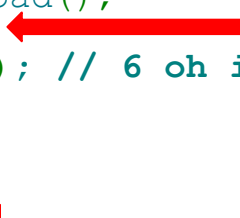
# Task 1: Replace spinwait

```
struct FIFOSemaphore {
    std::atomic<std::ptrdiff_t> next_ticket{1};
    std::atomic<std::ptrdiff_t> now_serving;

    FIFOSemaphore(std::ptrdiff_t initial_count) : now_serving{initial_count}
    {}

    void acquire() {
        auto my_ticket = next_ticket.fetch_add(1);
        auto old_now_serving = now_serving.load();
        if (old_now_serving < my_ticket)
            now_serving.wait(old_now_serving); // 6 oh its different unblock
        old == new
    }

    void release() {
        now_serving.fetch_add(1);
        now_serving.notify_one(); // 8
    }
};
```



Any problems?

# Task 1: Replace spinwait

Performs atomic waiting operations. Behaves as if it repeatedly performs the following steps:

- Compare the value representation of `this->load(order)` with that of old.
  - If those are equal, then blocks until `*this` is notified by `notify_one()` or `notify_all()`, or the thread is unblocked spuriously.
  - Otherwise, returns.





These functions are guaranteed to return only if value has changed, even if underlying implementation unblocks spuriously.

```
if (old now_serving < my_ticket) {  
    now_serving.wait(old_now_serving, my_ticket);  
}  
  
void release() {  
    now_serving.fetch_add(1);  
    now_serving.notify_one();  
};
```

Just because old changed doesn't mean the condition is true





# Task 1: Replace spinwait

`notify_one()` might not wake 3!



# Task 1: Replace spinwait

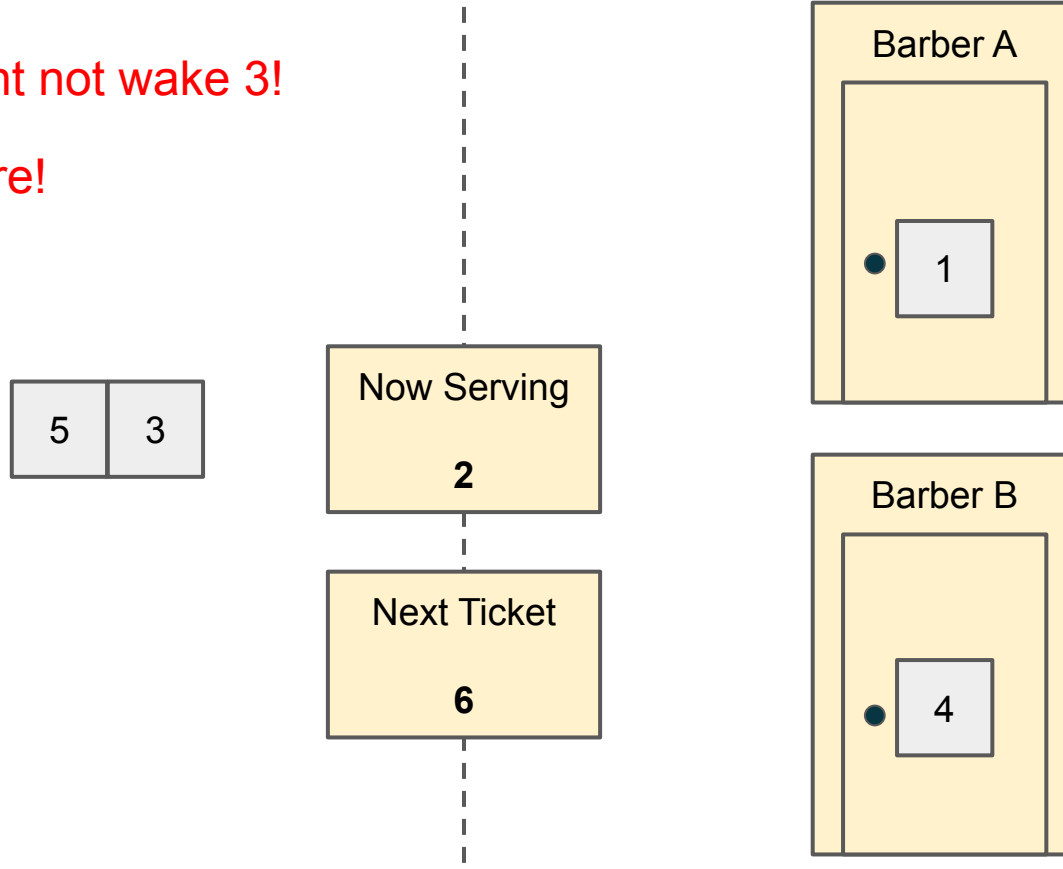
`notify_one()` might not wake 3!



# Task 1: Replace spinwait

`notify_one()` might not wake 3!

-> not FIFO anymore!



# Demo 5: Queue of semaphores

```
void acquire() {
    auto waiter = std::make_shared<Waiter>();
    {
        std::scoped_lock lock{mut};
        if (count > 0) {
            count--; // Positive count,
            return; // simply decrement without blocking
        }
        waiters.push(waiter); // Zero count, add to waiters
    }
    waiter->sem.acquire(); // and block on the semaphore
}
```

```
struct FIFOSemaphore5 {
    struct Waiter {
        std::binary_semaphore sem{0};
    };

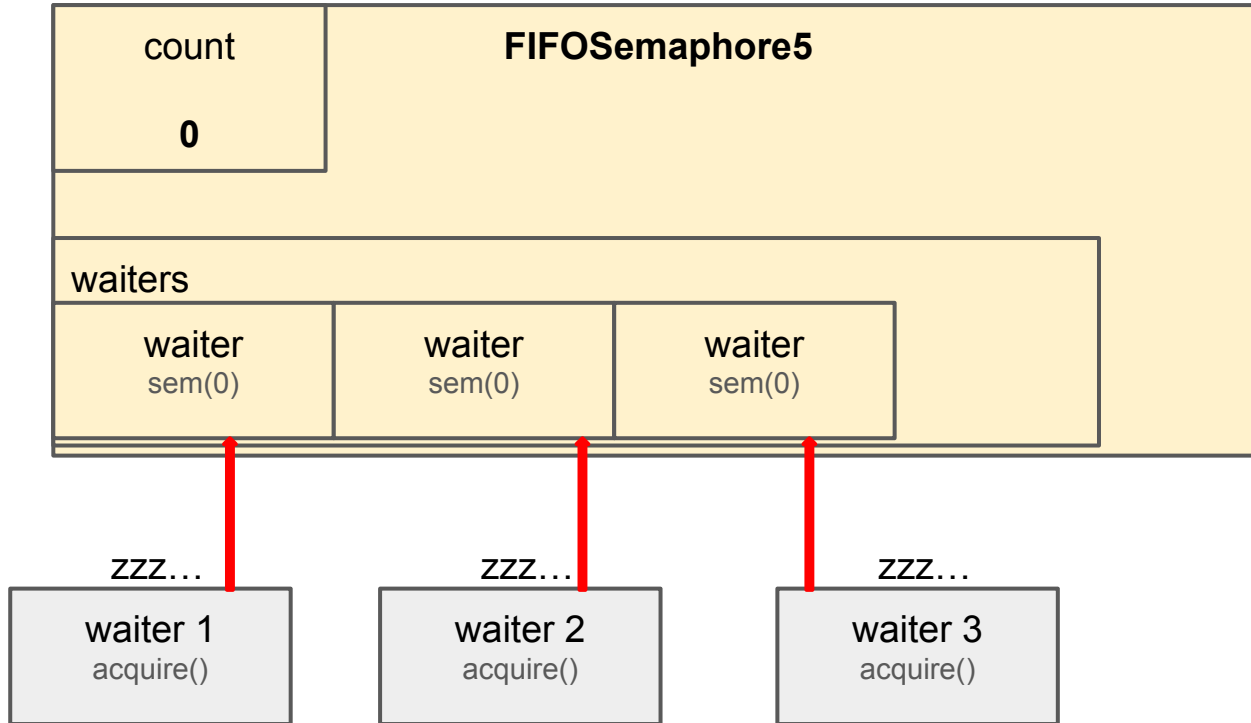
    std::mutex mut;
    std::queue<std::shared_ptr<Waiter>> waiters;
    std::ptrdiff_t count;

    FIFOSemaphore5(std::ptrdiff_t initial_count)
        : mut{}, waiters{}, count{initial_count} {}
}
```

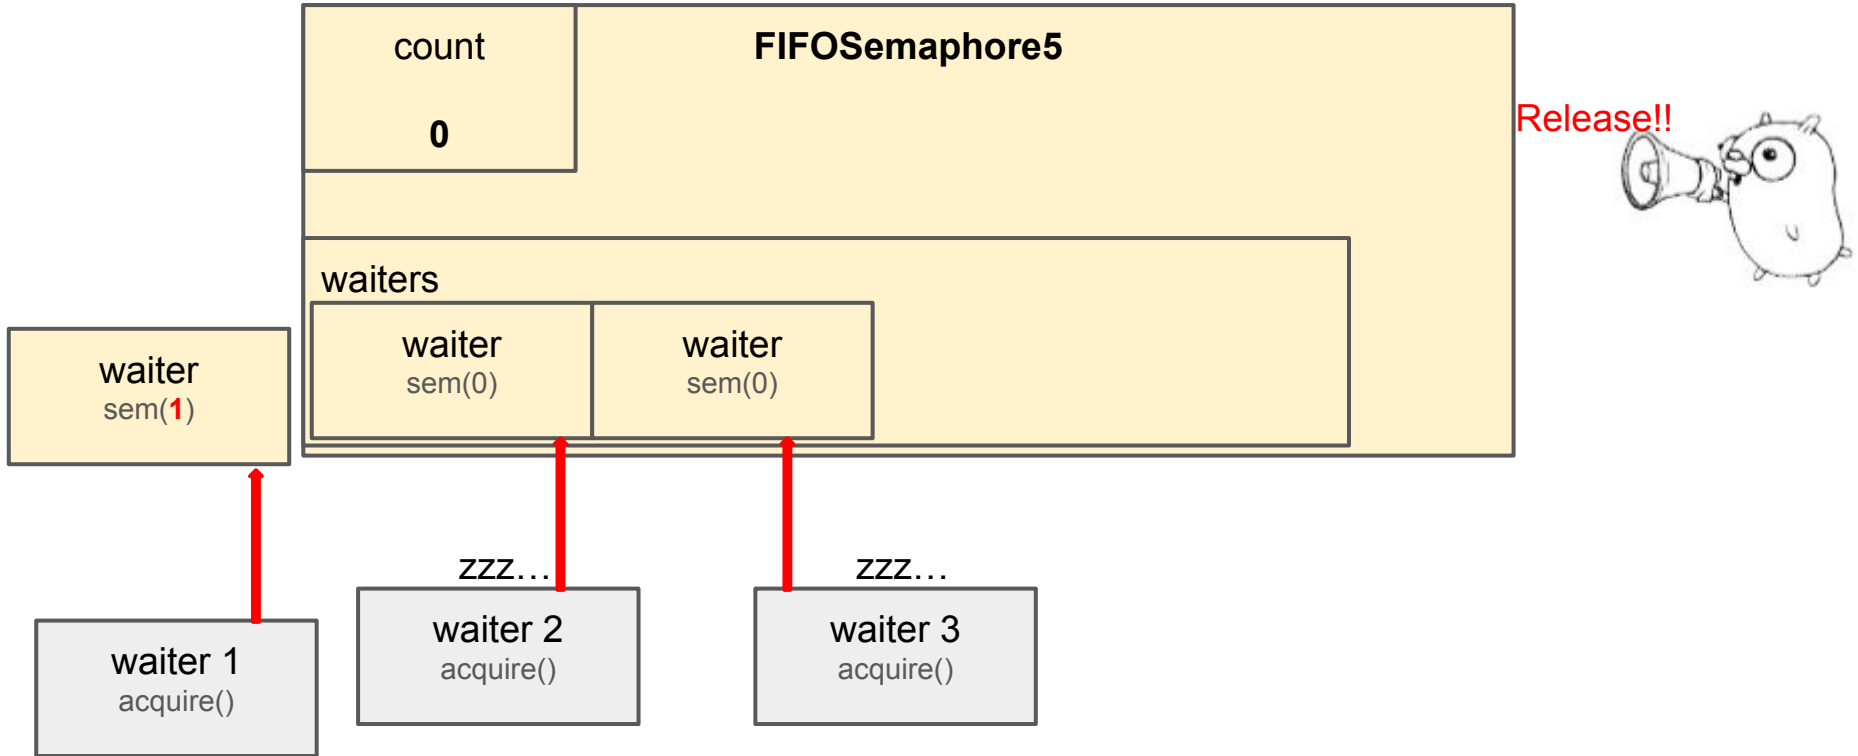
```
void release() {
    std::shared_ptr<Waiter> waiter;
    {
        std::scoped_lock lock{mut};
        if (waiters.empty()) {
            count++; // No waiters, simply increment count
            return;
        }

        waiter = waiters.front(); // Pop a waiter
        waiters.pop();
    }
    waiter->sem.release(); // and signal it
}
```

## Demo 5: Queue of semaphores

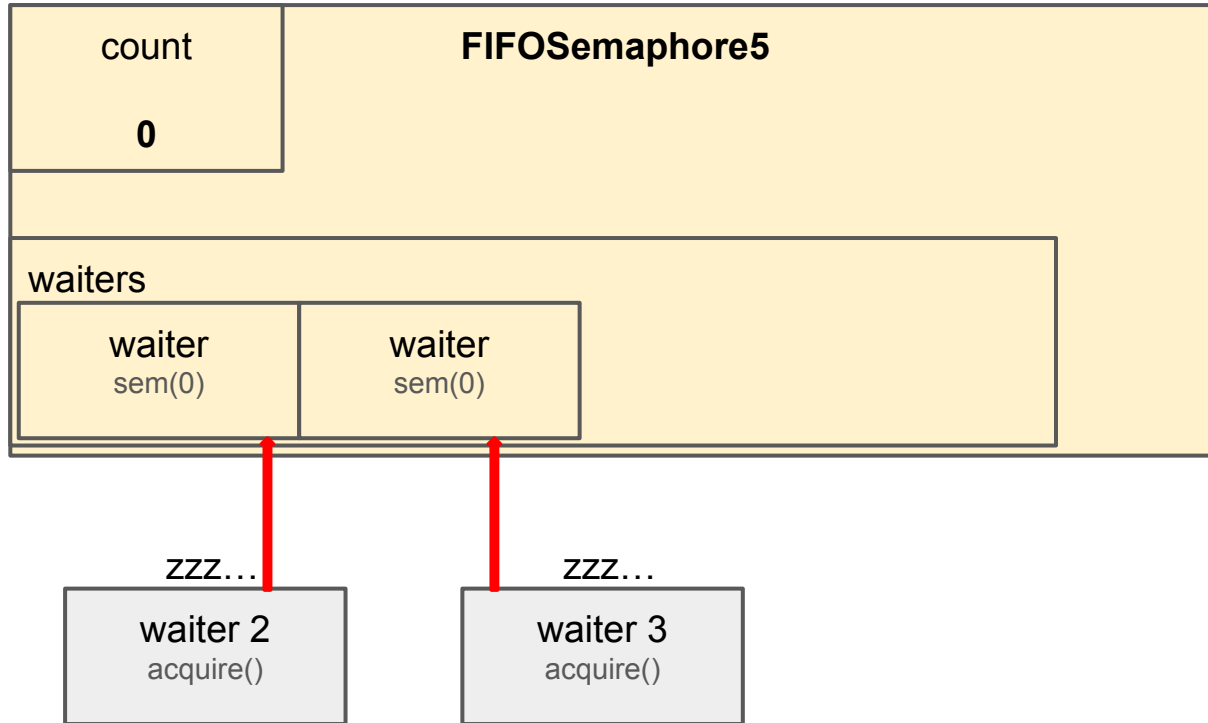


# Demo 5: Queue of semaphores






## Demo 5: Queue of semaphores




## Demo 5: Queue of semaphores



Release!!



## Demo 5: Queue of semaphores



## Demo 5: extras

Issues:

- 2 allocations each time waiter is added to queue

How about we do only 1 allocation?

## Demo 5: extras

```
void acquire() {
    Waiter waiter;

    {
        std::scoped_lock lock{mut};
        if (count > 0) {
            count--;
            return;
        }

        if (back == nullptr) {
            front = back = &waiter;
        } else {
            back->next = &waiter;
            back = &waiter;
        }
    }
    waiter.sem.acquire();
}
```

Since destruction of waiter is after all “external” accesses to it, no invalid memory access.  
tldr; no lifetime issues

```
struct FIFOSemaphore7 {
    struct Waiter {
        std::binary_semaphore sem{0};
        Waiter *next{nullptr};
    };

    std::mutex mut;
    Waiter *front;
    Waiter *back;
    std::ptrdiff_t count;

    FIFOSemaphore7(std::ptrdiff_t initial_count)
        : mut{}, front{nullptr}, back{nullptr}, count{initial_count} {}

    void release() {
        Waiter *head;
        {
            std::scoped_lock lock{mut};
            if (front == nullptr) {
                count++;
                return;
            }
            head = front;
            front = front->next;
            if (front == nullptr) {
                back = nullptr;
            }
        }
        head->sem.release();
    }
}
```

## Demo 6: Using a buffered channel

- Sending on a full/unbuffered channel -> block
- Suppose `initial_count = N`
  - Make a buffered channel of size N.
  - If  $M > N$  goroutines arrive at the same time,  $(M - N)$  goroutines block.
  - release -> send to channel
  - acquire -> recv from channel

Basically a semaphore!

## Demo 6: Using a buffered channel

```
type hchan struct {
    qcount    uint           // total data in the queue
    dataqsiz  uint           // size of the circular queue
    buf       unsafe.Pointer // points to an array of dataqsiz elements
    elemsize  uint16
    closed    uint32
    elemtype  *_type // element type
    sendx     uint     // send index
    recvx     uint     // receive index
    recvq     waitq    // list of recv waiters
    sendq     waitq    // list of send waiters
}

type waitq struct {
    first *sudog
    last  *sudog
}
```

```
type sudog struct {
    // The following fields are protected by the hchan.lock of the
    // channel this sudog is blocking on. shrinkstack depends on
    // this for sudogs involved in channel ops.

    g *g

    next *sudog
    prev *sudog
}
```

current impl is fifo, but not guaranteed by the standard.

## Demo 7: Using a daemon goroutine

Bad idea to rely on implementation-defined behaviour.

How about we have a goroutine manage others!



## Demo 7: Using a daemon goroutine

```
type Semaphore2 struct {  
    acquireCh chan chan struct{}  
    releaseCh chan struct{}  
}
```

```
func (s *Semaphore2) Acquire() {  
    ch := make(chan struct{})  
    // Send daemon a channel that can be used to unblock us  
    s.acquireCh <- ch  
    // Block until daemon decides to unblock us  
    <-ch  
}
```

```
func (s *Semaphore2) Release() {  
    s.releaseCh <- struct{}{}  
}
```

```
func NewSemaphore2(initial_count int) *Semaphore2 {  
    sem := new(Semaphore2)  
    sem.acquireCh = make(chan chan struct{}, 100)  
    sem.releaseCh = make(chan struct{}, 100)  
  
    go func() {  
        count := initial_count  
        // The FIFO queue that stores the channels used to unblock waiters  
        waiters := NewChanQueue()
```

## Demo 7: Using a daemon goroutine

```
for {  
    select {  
        case <-sem.releaseCh:  
            // ...  
  
        case ch := <-sem.acquireCh:  
            // ...  
    }  
}
```

Either process a release or acquire request

## Demo 7: Using a daemon goroutine

```
case ch := <-sem.acquireCh: // Decrement or add a waiter
    if count > 0 {
        count--
        ch <- struct{}{} // Since count is +ve, don't block the waiter
    } else {
        waiters.PushBack(ch) // Add waiter to back of the wait queue
    }
}
```

Either process a release or acquire request

## Demo 7: Using a daemon goroutine

```
case <-sem.releaseCh: // Increment or unblock a waiter
    if waiters.Len() > 0 {
        ch := waiters.Pop()
        ch <- struct{}{} // Unblocks the oldest waiter
    } else {
        count++
    }
```

Either process a release or acquire request

# Demo 7: Using a daemon goroutine

Why a buffer of 100?

```
func NewSemaphore2(initial_count int) *Semaphore2 {  
    sem := new(Semaphore2)  
    sem.acquireCh = make(chan chan struct{}, 100)  
    sem.releaseCh = make(chan struct{}, 100)  
  
    go func() {  
        count := initial_count  
        // The FIFO queue that stores the channels used to unblock waiters  
        waiters := NewChanQueue()
```

# Demo 7: Using a daemon goroutine

Why a buffer of 100?

If goroutine blocks on

```
func (s *Semaphore2) Acquire() {  
    ch := make(chan struct{})  
    // Send daemon a channel that can be used to  
    unblock us  
    s.acquireCh <- ch  
    // Block until daemon decides to unblock us  
    <-ch  
}
```

for a long time (wakeups not in FIFO), no guarantees of FIFO.

**BUT**, if send succeeds (buffered in channel), we get FIFO behavior.

```
type hchan struct {  
    qcount    uint           // total data in the queue  
    dataqsiz  uint           // size of the circular queue  
    buf       unsafe.Pointer // points to an array of dataqsiz elements  
    elemsize  uint16  
    closed    uint32  
    elemtype  *_type // element type  
    sendx     uint       // send index  
    recvx     uint       // receive index  
    recvq     waitq      // list of recv waiters  
    sendq     waitq      // list of send waiters  
  
    // lock protects all fields in hchan, as well as several  
    // fields in sudogs blocked on this channel.  
    //  
    // Do not change another G's status while holding this lock  
    // (in particular, do not ready a G), as this can deadlock  
    // with stack shrinking.  
    lock mutex  
}
```

Pls scan for attendance



See you next week!